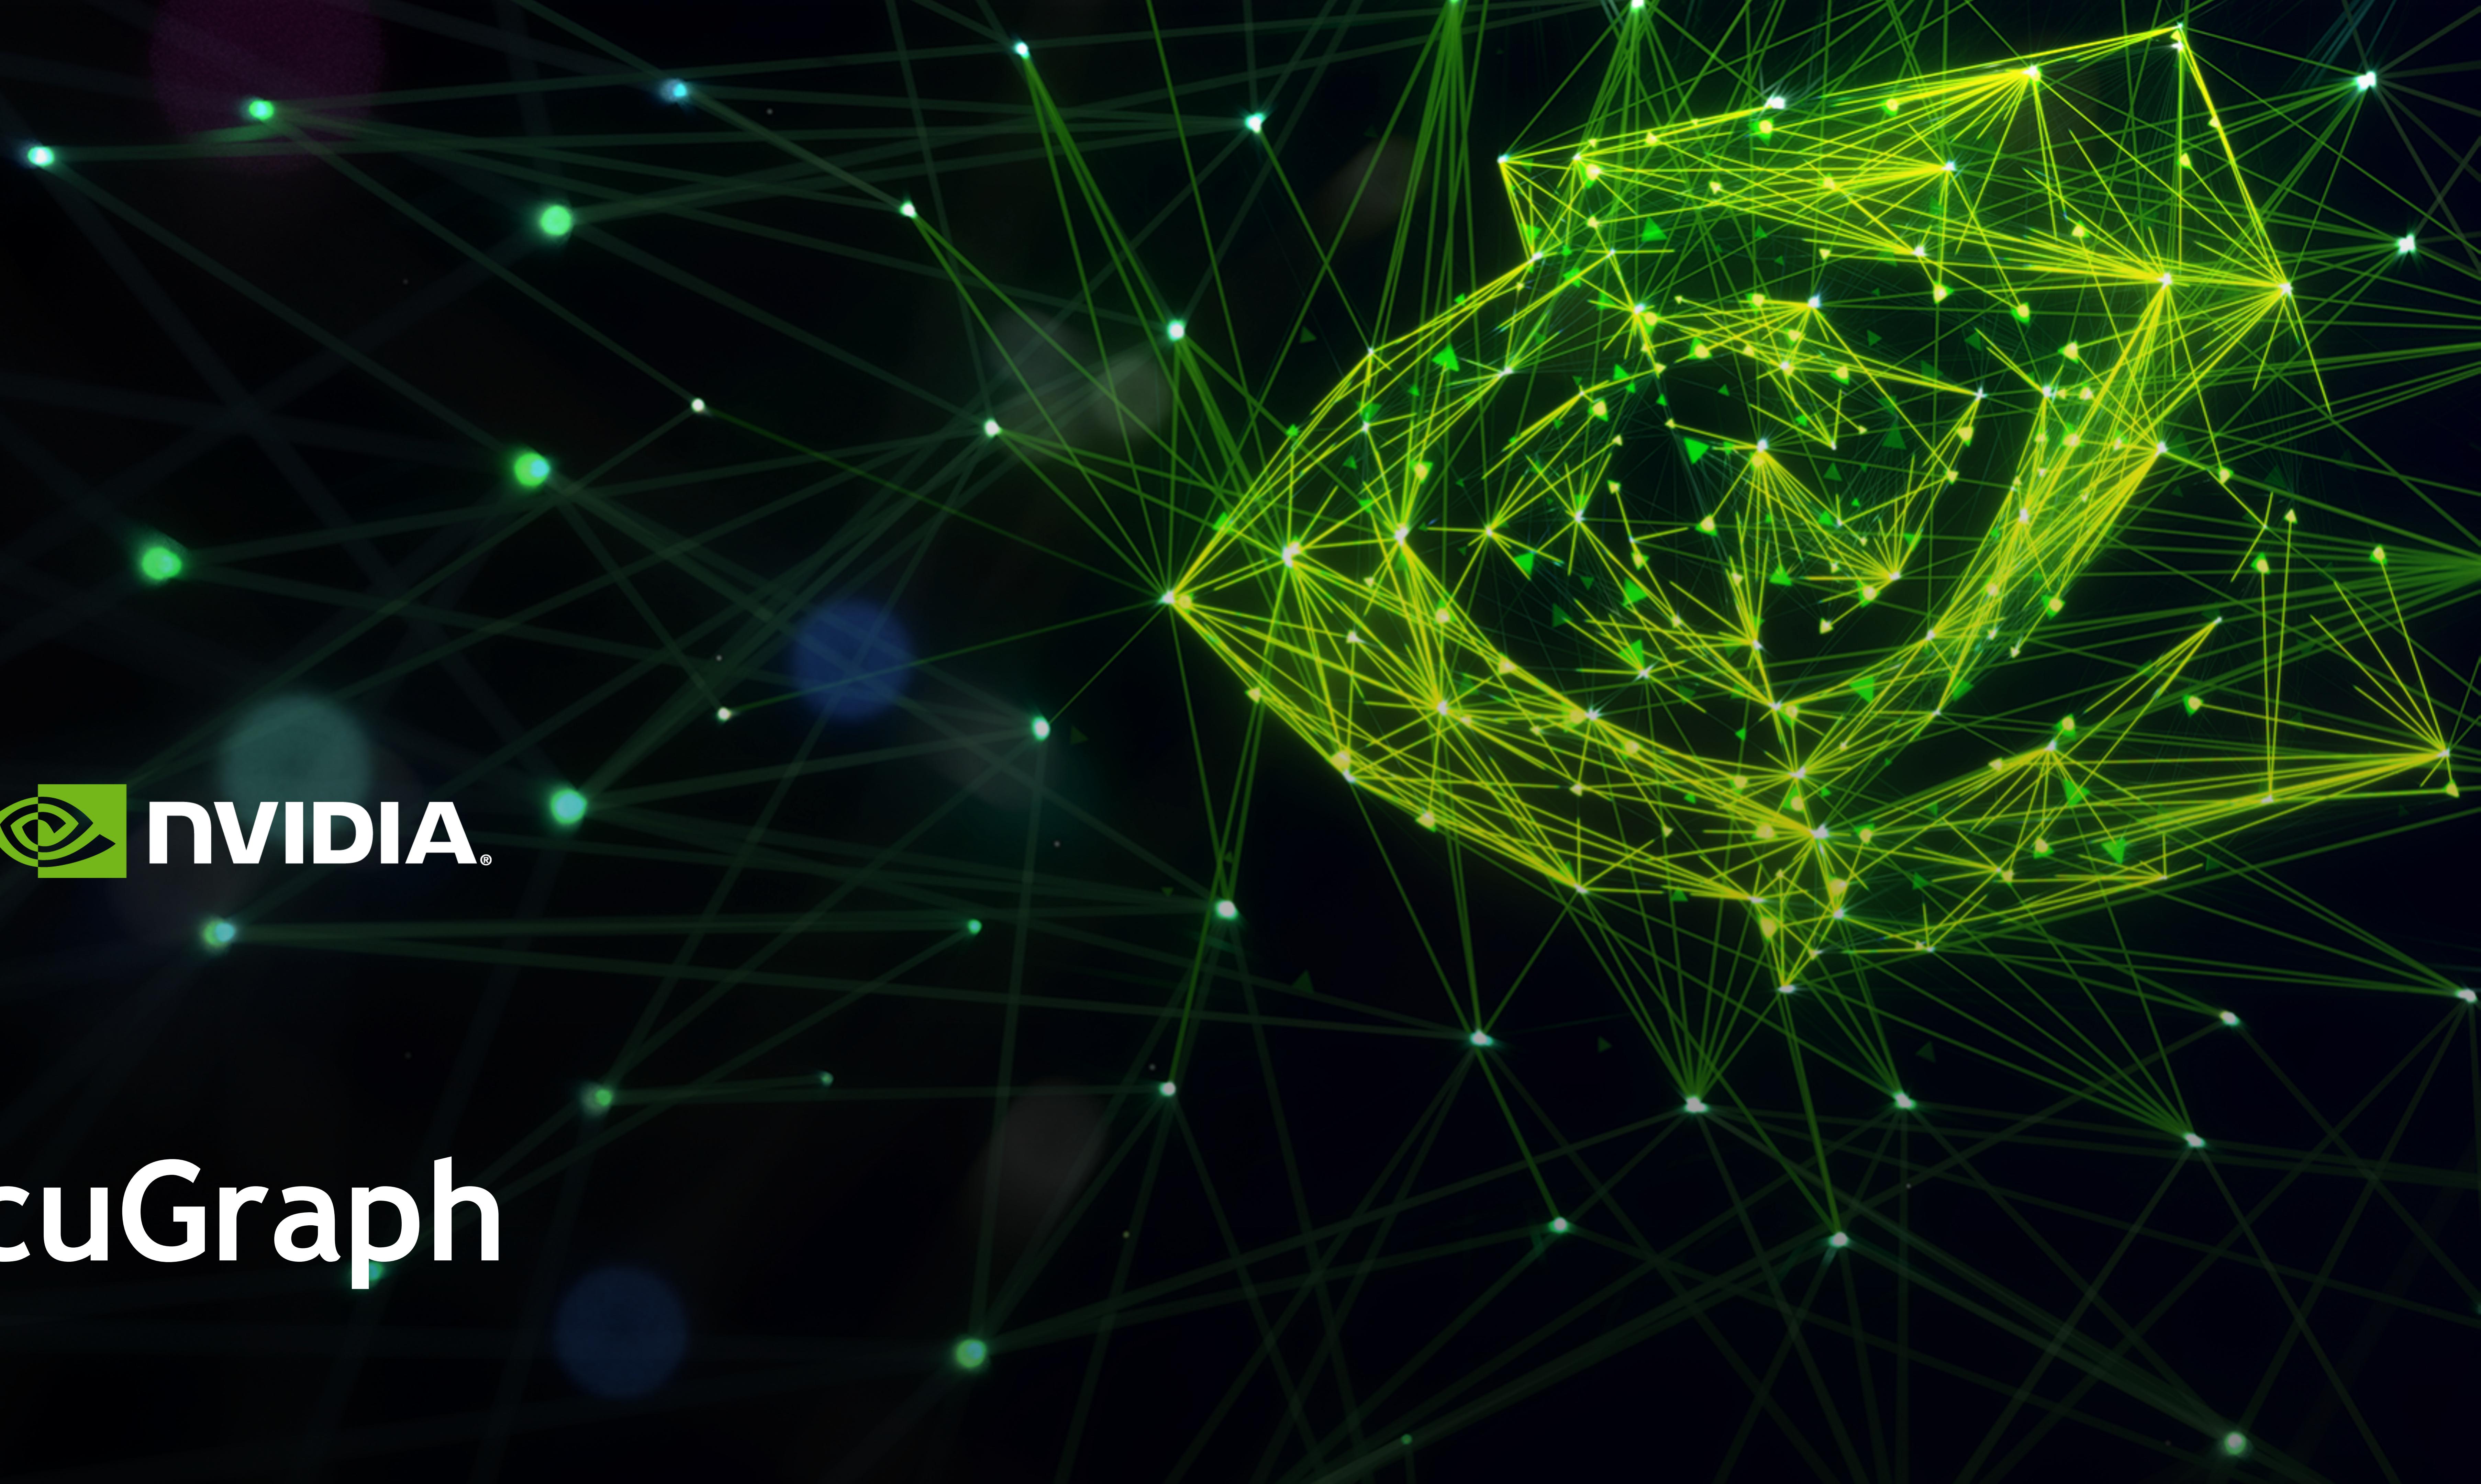




nVIDIA[®]

cuGraph



ACCELERATED GRAPH ANALYTICS

- cuGraph library is a collection of GPU accelerated graph algorithms
 - Over three dozen algorithms, more planned
 - Scaling from small single GPUs graphs to massive multi-GPU clusters
- users familiar with NetworkX will quickly recognize the NetworkX-like API provided in cuGraph, with the goal to allow existing code to be ported with minimal effort into RAPIDS.
 - Support cuDF Dataframes
 - CuPy and Numpy sparse matrixes
 - NetworkX graph objects
 - more ..
- Python API
- C API
- C++ API

```
import cudf
import cugraph

# read data into a cuDF DataFrame using read_csv
gdf = cudf.read_csv("graph_data.csv", names=["src", "dst"], dtype=["int32", "int32"])

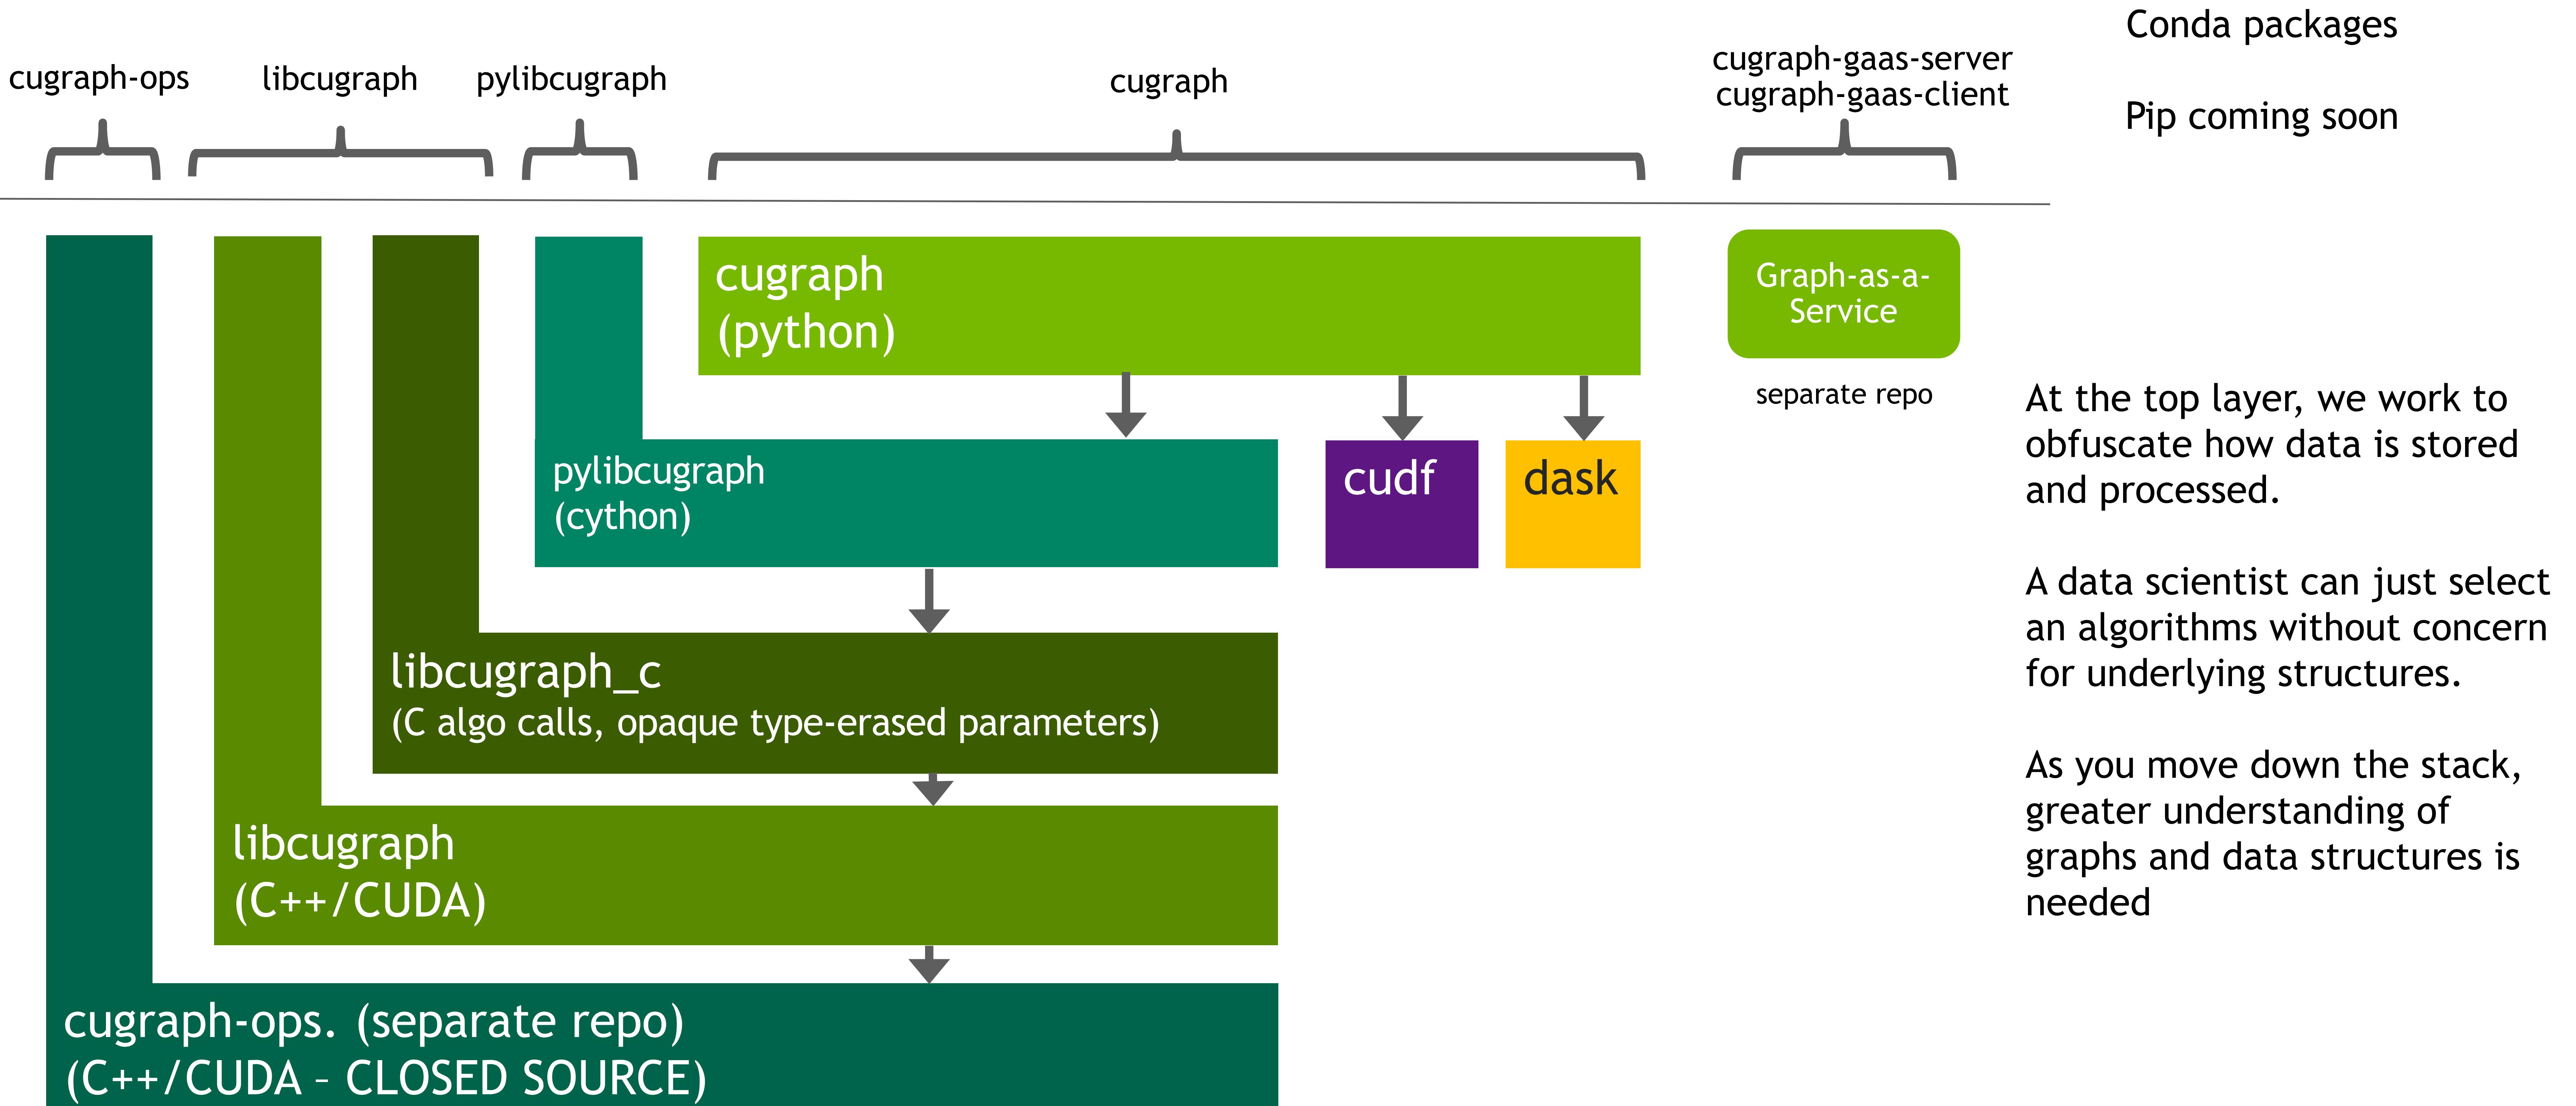
# We now have data as edge pairs
# create a Graph using the source (src) and destination (dst) vertex pairs
G = cugraph.Graph()
G.from_cudf_edgelist(gdf, source='src', destination='dst')

# Let's now get the PageRank score of each vertex by calling cugraph.pagerank
df_page = cugraph.pagerank(G)

# Let's look at the top 10 PageRank Score
df_page.sort_values('pagerank', ascending=False).head(10)
```

THE CUGRAPH STACK

many layers of cugraph

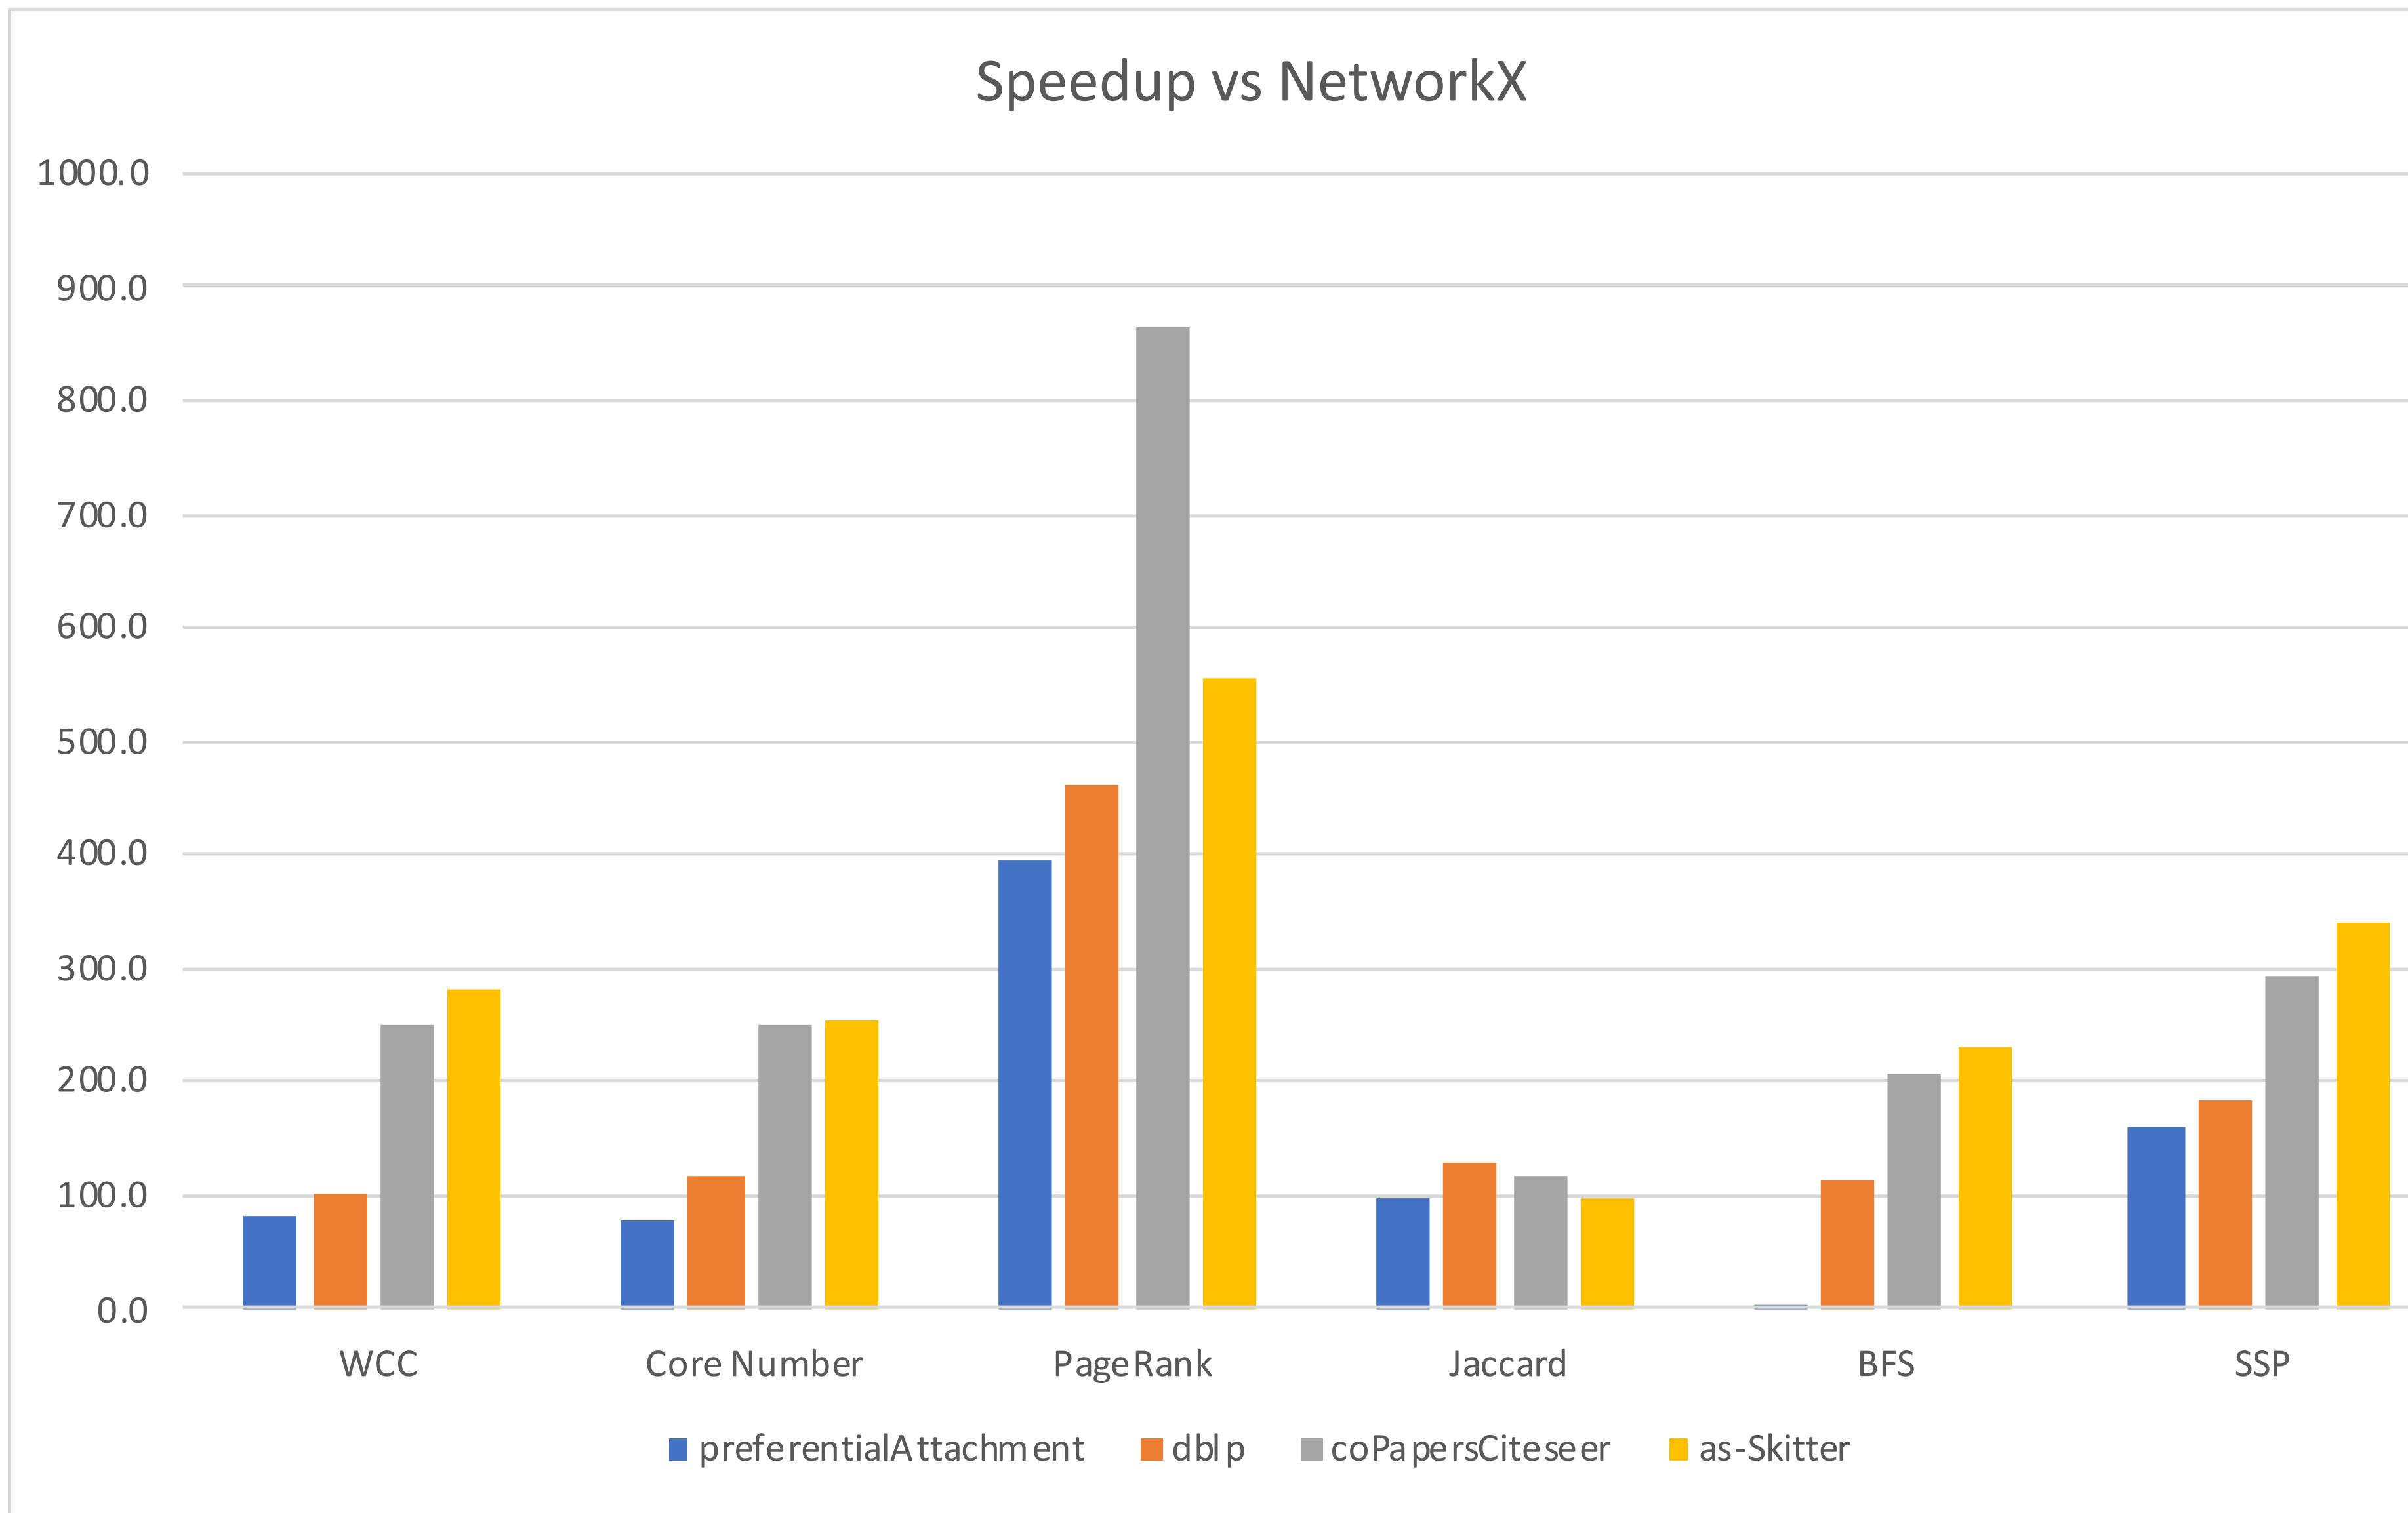


AVAILABLE ALGORITHMS

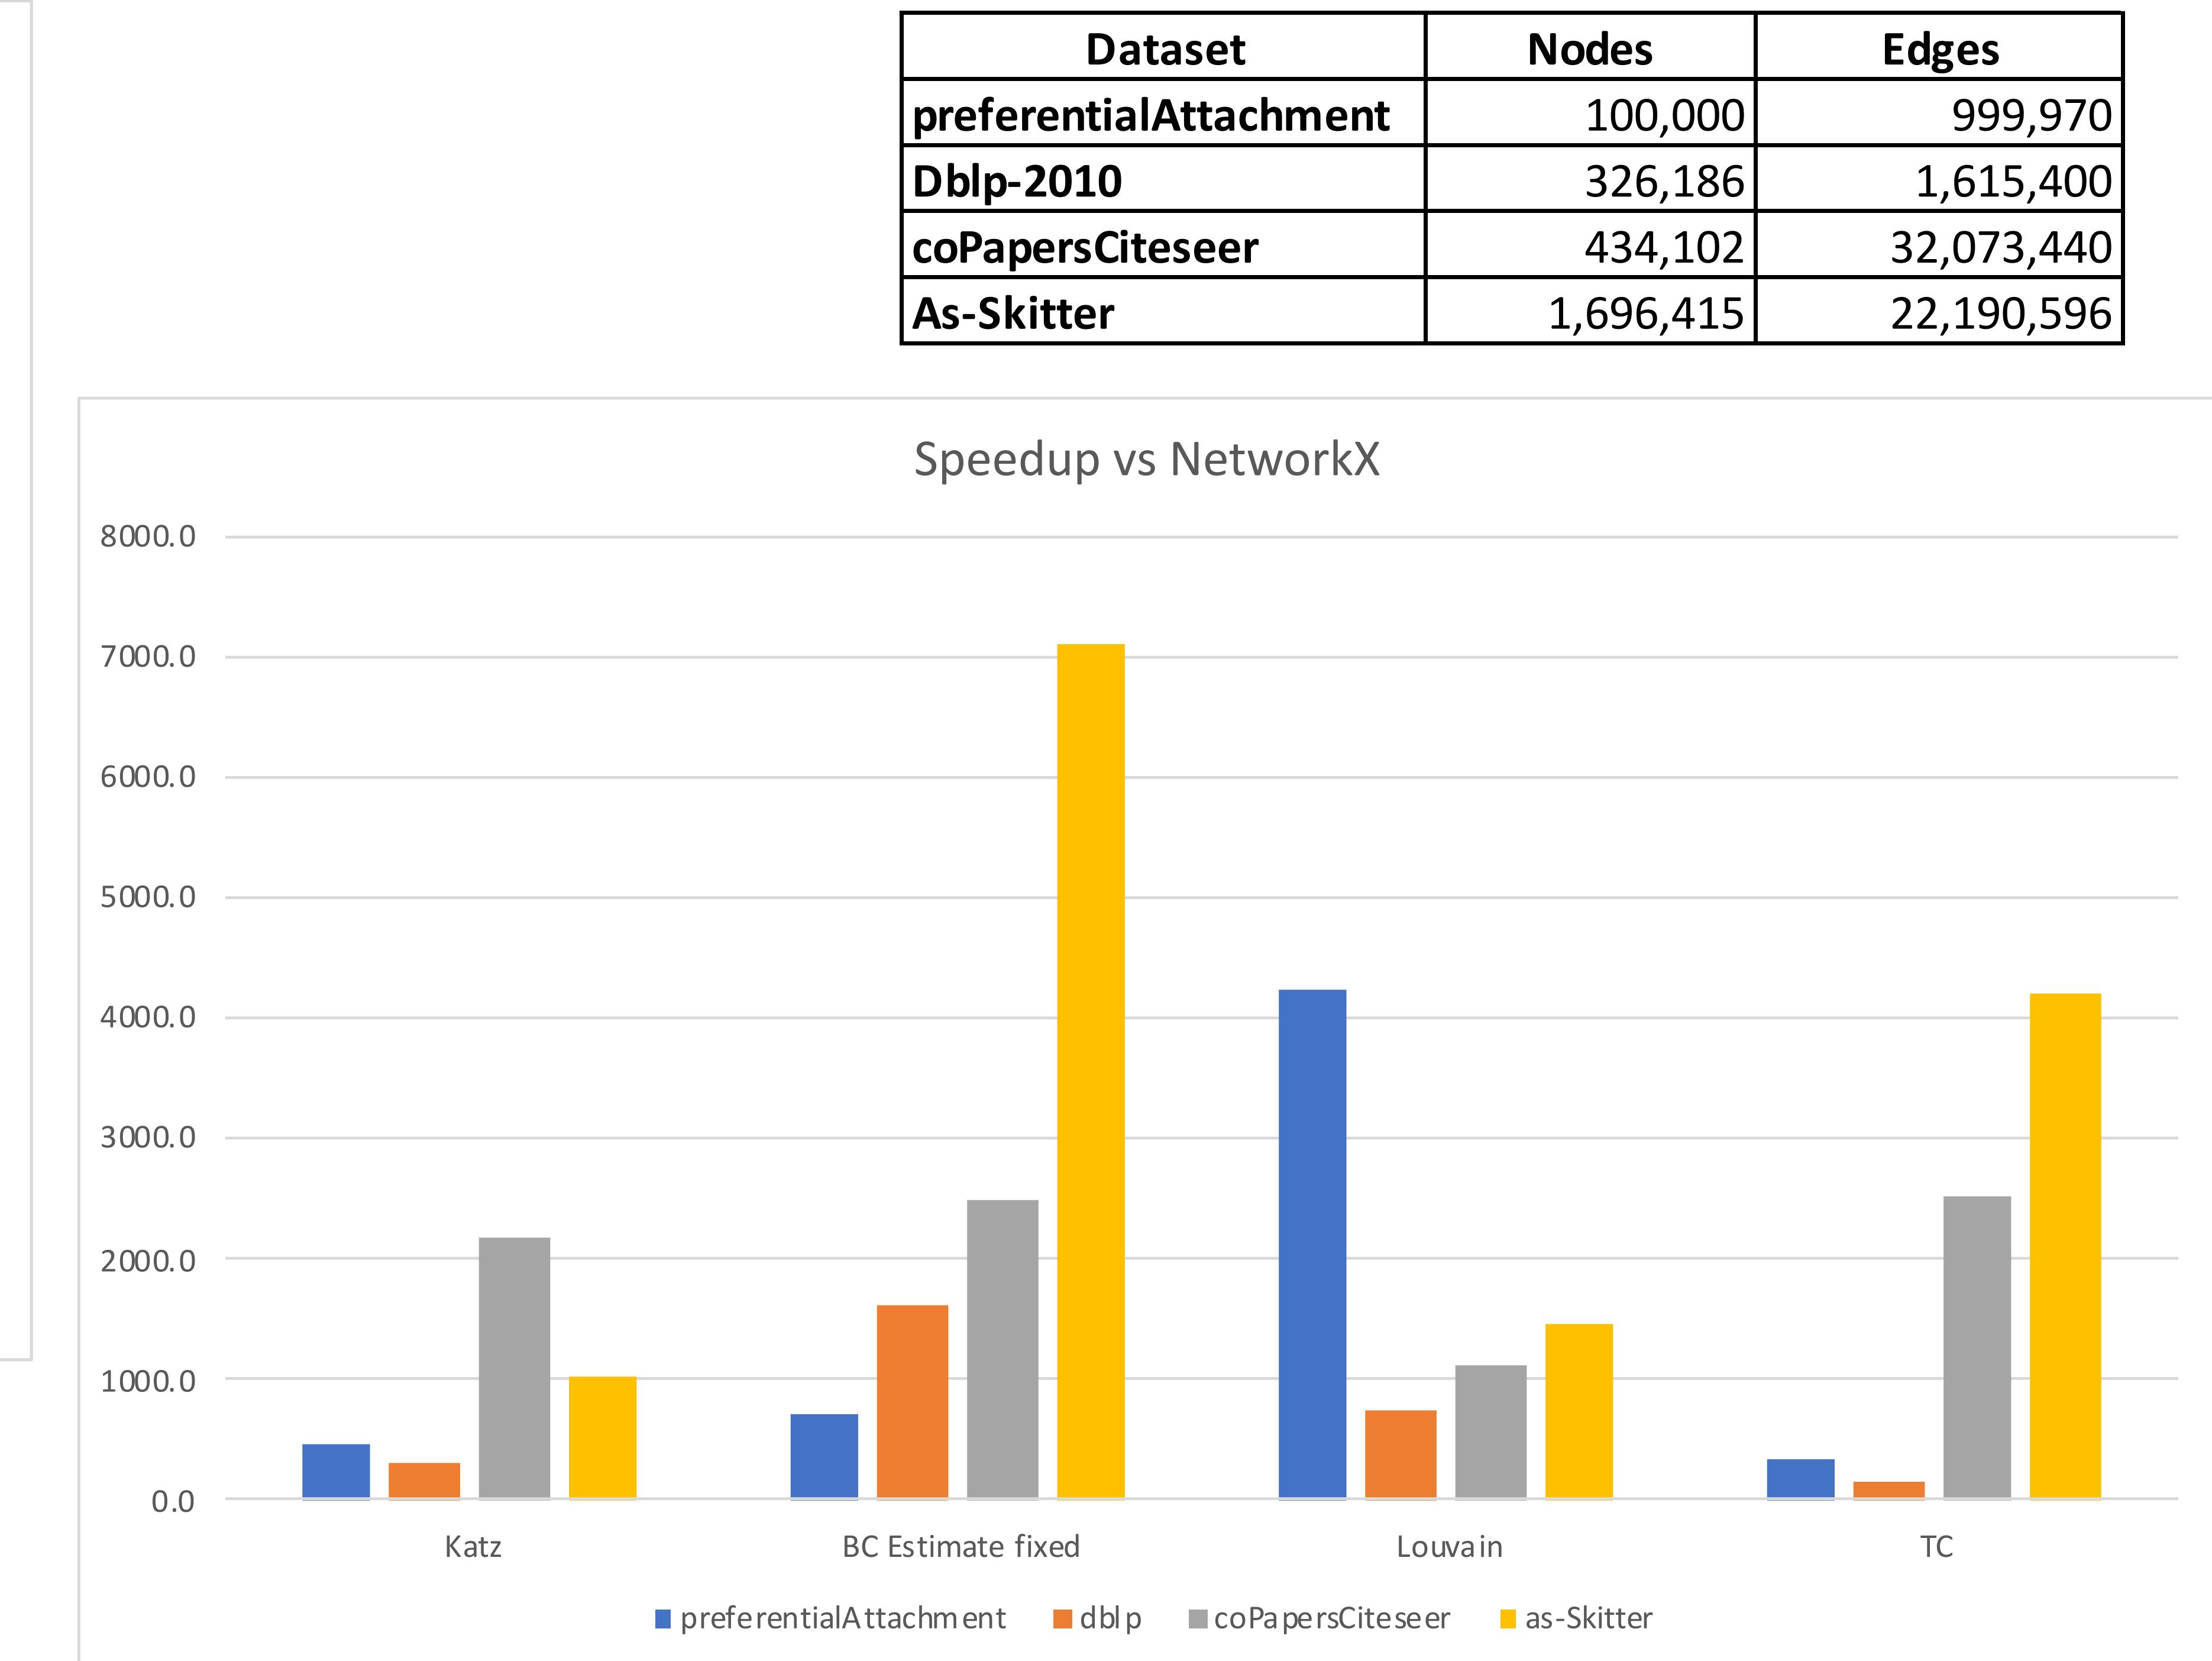
Class	Algorithms	MNMG	Class	Algorithms	MNMG
Centrality	Katz	Yes	Link Analysis	PageRank	Yes
	Betweenness Centrality			Personal PageRank	Yes
	Edge Betweenness Centrality			HITS	Yes
	Degree Centrality	Yes	Link Prediction	Jaccard Similarity	
	Eigenvector Centrality	Yes		Weighted Jaccard Similarity	
				Overlap Similarity	
Community	Leiden			Sorensen	
	Louvain	Yes	Traversal	Breadth First Search (BFS)	Yes
	Ensemble Clustering for Graphs			Single Source Shortest Path (SSSP)	Yes
	Spectral-Clustering - Balanced Cut		Sampling	Random Walks (Uniform and Biased)	
	Spectral-Clustering - Modularity			EgoNet	
	Subgraph Extraction			Node2Vec	
	Triangle Counting	Yes		Neighborhood	Yes
Components	K-Truss		Other	Minimum Spanning Tree	
				Maximum Spanning Tree	
	Weakly Connected Components	Yes		Renumbering	Yes
Core	Strongly Connected Components			Symmetrize	Yes
			Structure		
	K-Core				
Layout	Core Number				
	Force Atlas 2				
Linear Assignment					
	Hungarian				

PERFORMANCE SPEEDUP

Single GPU Comparison Against NetworkX



Two charts since there is a change in Y axis size
Betweenness Centrality is with K = 100

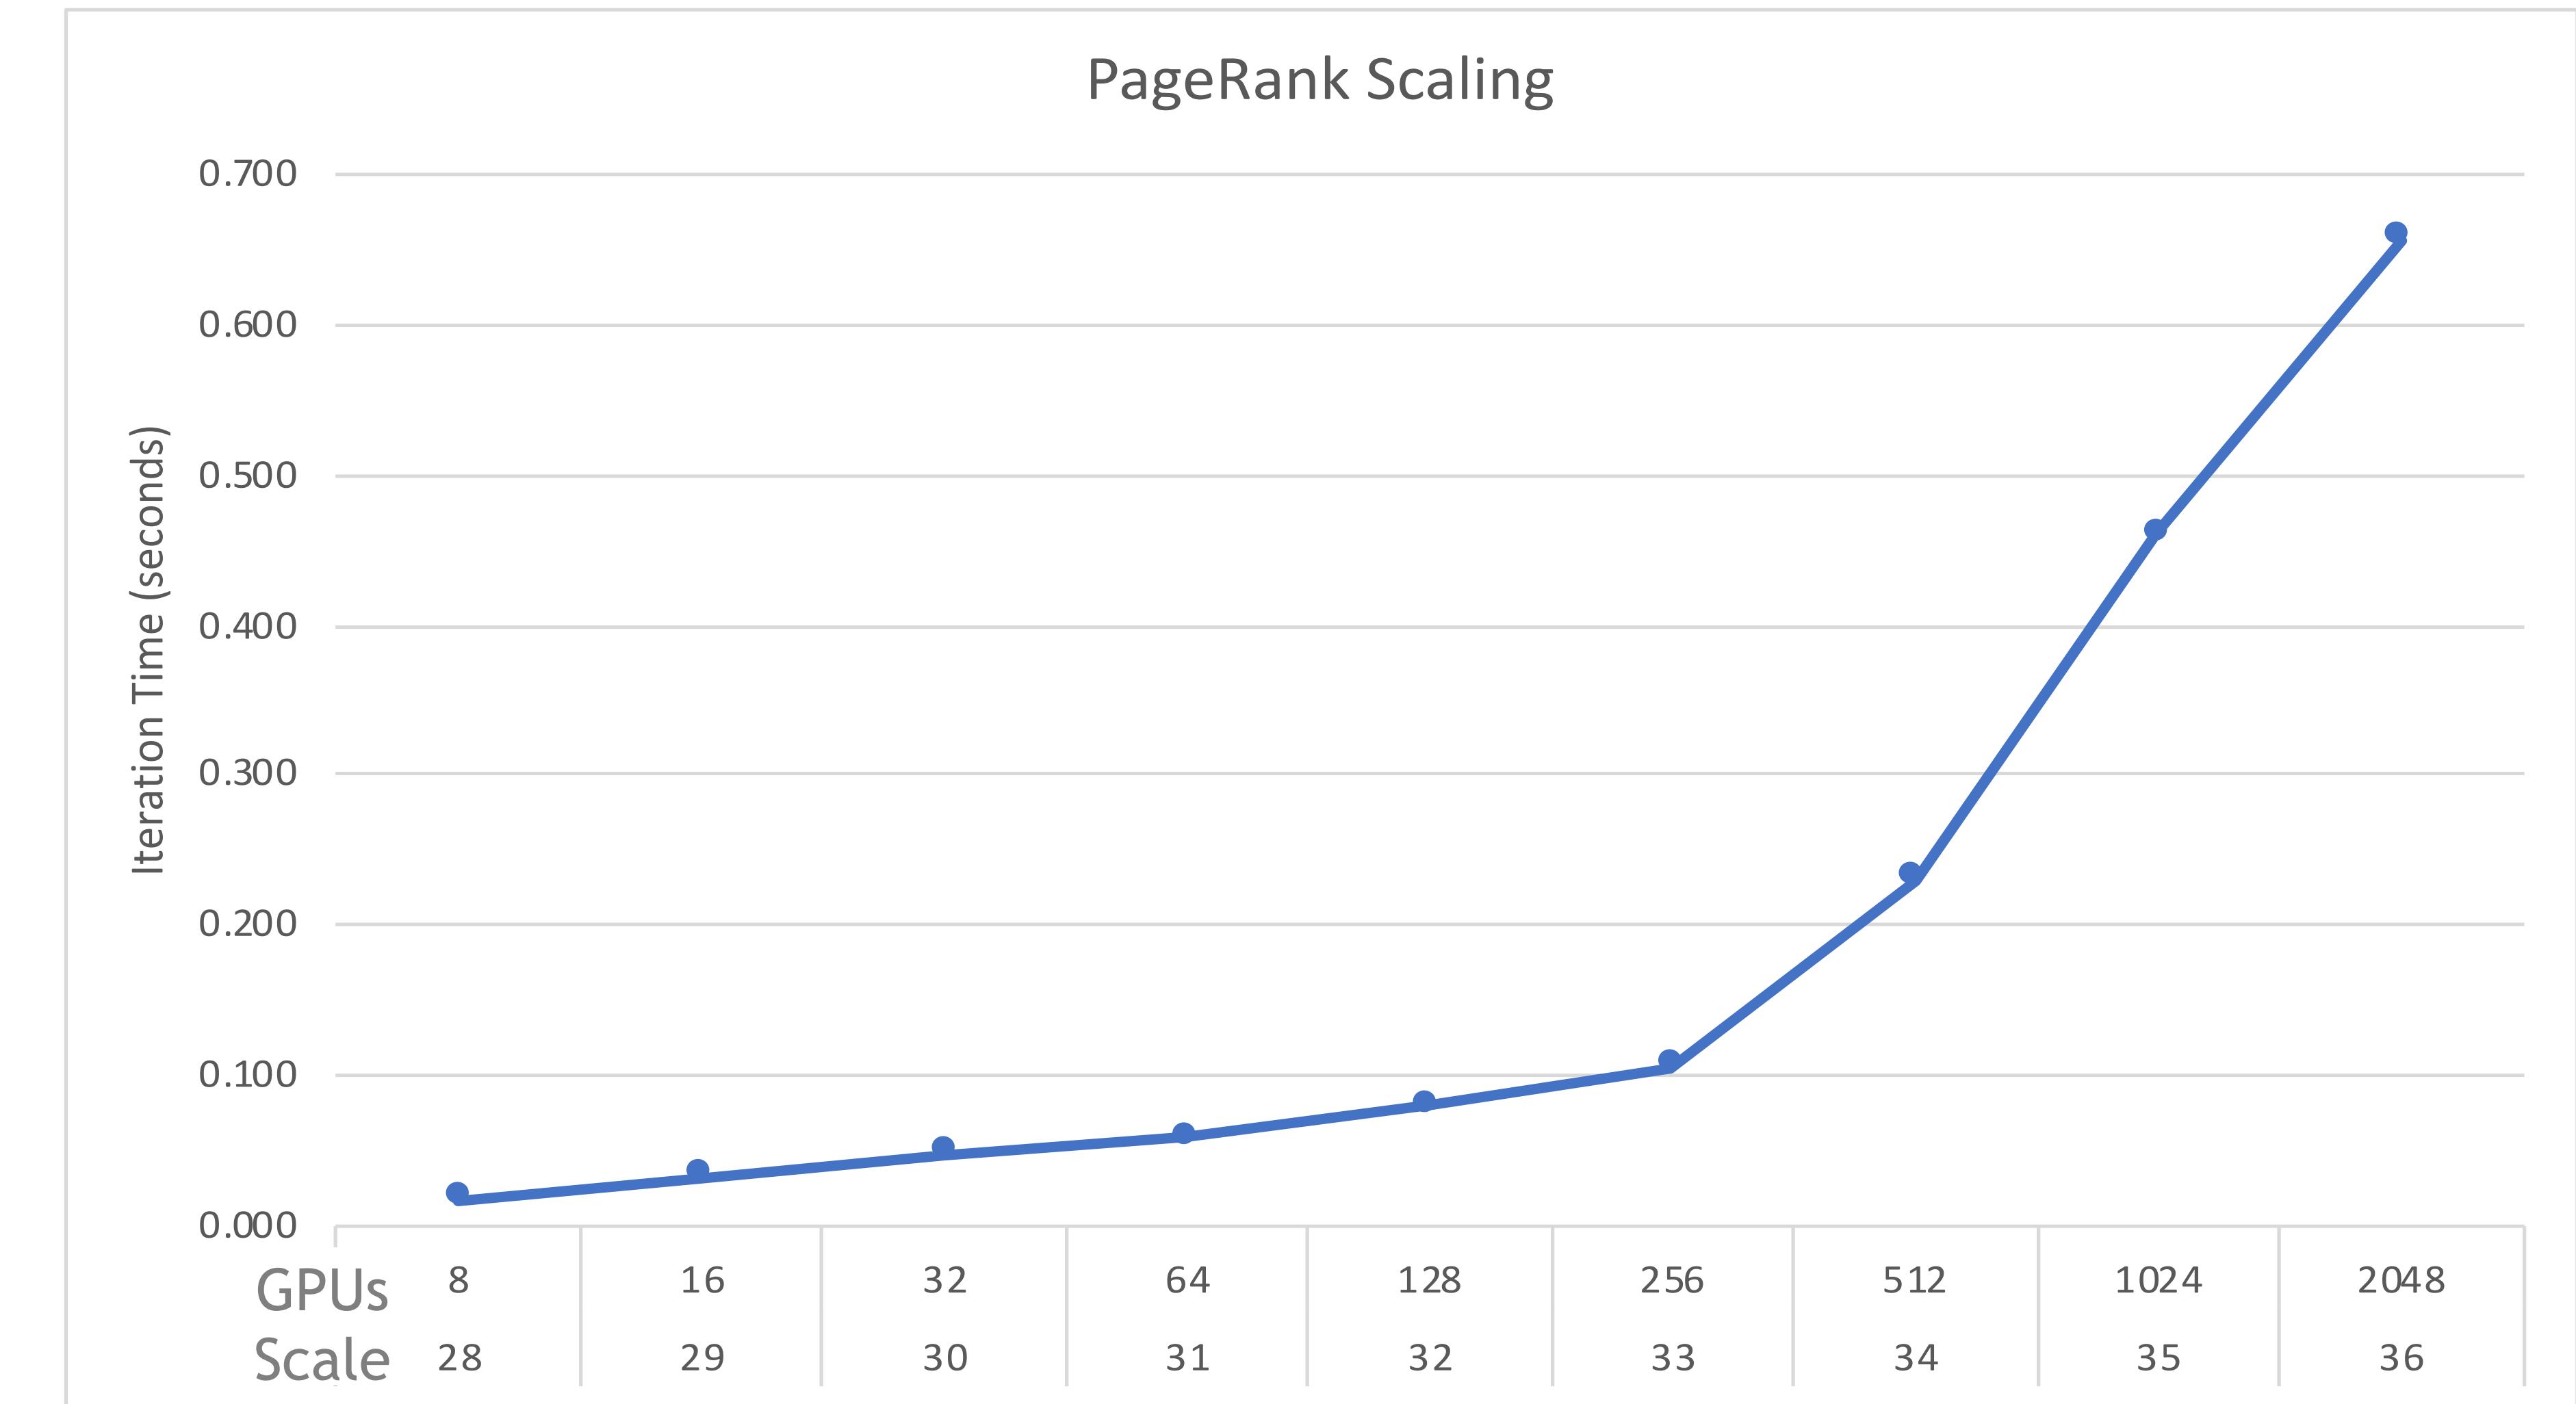


OVER 1 TRILLION EDGES !!!

cuGraph Scaling to Massive Scale

cuGraph scales smoothly from small graphs on 1 GPU to massive graphs with trillions of edges on 2,048 GPUs

- Being able to get answers in a matter of seconds regardless of graph scale if important
- **PageRank:** Scale 36 (1.1 trillion directed edges) in 19.3 seconds (0.66 seconds per iteration, 2,048 GPUs)
- **Louvain:** Scale 35 (0.55 trillion undirected edges or 1.1 trillion directed edges) in 336 seconds (1024 GPUs)
- Large scale testing just started
 - More analytics will be supported at 1000+ GPU scale
 - Continuous optimization for both memory usage and performance scalability



RMAT Data Generator
using Graph500 Specs
Edge Factor of 16

Scale	Number of Vertices	Number of Edges	COO Data Size (GB) in GPU
28	268,435,456	4,294,967,296	80
29	536,870,912	8,589,934,592	160
30	1,073,741,824	17,179,869,184	320
31	2,147,483,648	34,359,738,368	640
32	4,294,967,296	68,719,476,736	1,280
33	8,589,934,592	137,438,953,472	2,560
34	17,179,869,184	274,877,906,944	5,120
35	34,359,738,368	549,755,813,888	10,240
36	68,719,476,736	1,099,511,627,776	20,480
37	137,438,953,472	2,199,023,255,552	40,960

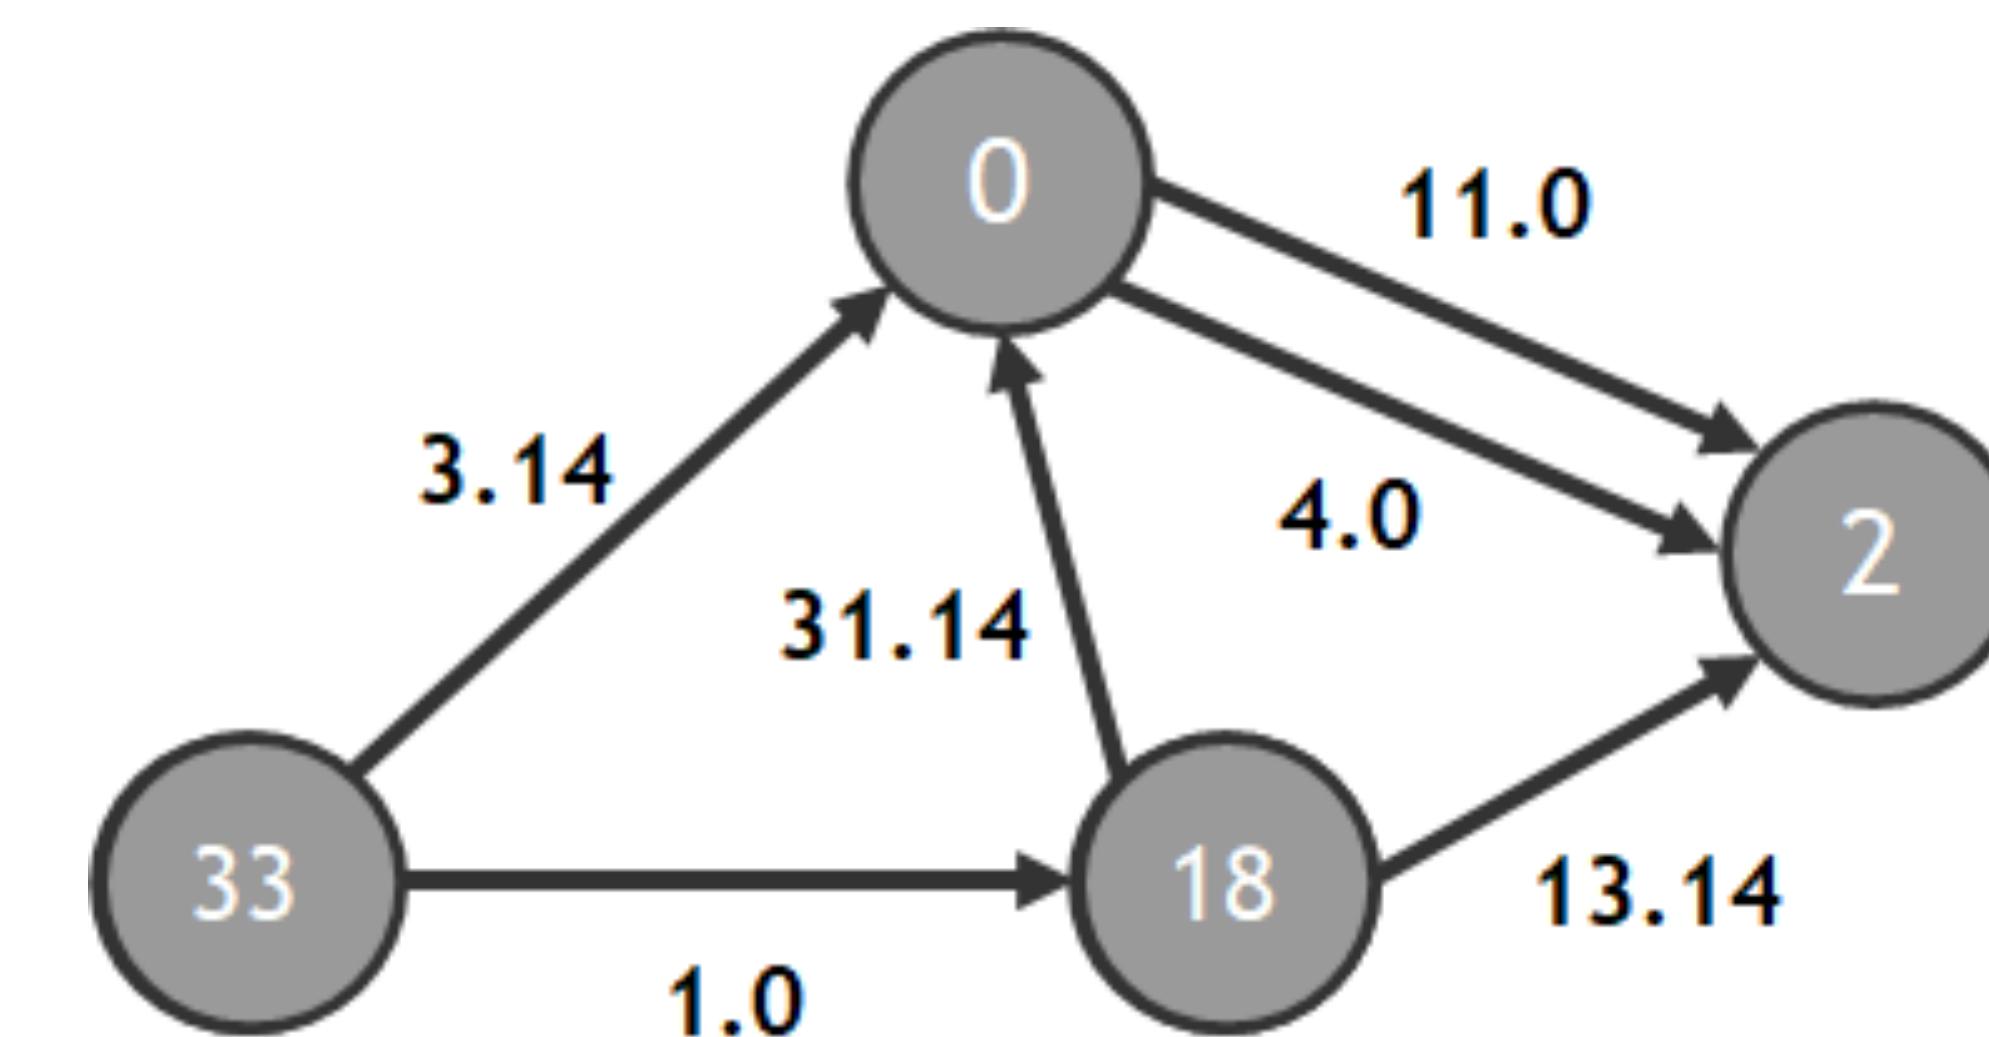
Property Graphs

PROPERTY GRAPH

Added in release 21.12

- Prior to 21.12, cugraph users were limited to creating graphs from source and destination vertices, and a single (optional) numeric weight value for each edge.

source	destination	weight
33	0	3.14
33	18	1.0
18	0	31.14
0	2	4.0
0	2	11.0
18	2	13.14



- ...but real data can contain numerous attributes of various types on both vertices and edges
 - Being able to store, access, and process attributes is important
- Property Graphs offer this ability, and are critical for supporting Heterogenous graph in GNNs
- Note: this is not proposal to support a graph database or even a graph query language*

PROPERTY GRAPH

Added in release 21.12

- Property Graphs offer the ability to load graph data containing multiple, heterogeneous attributes on items that represent vertices and edges.
- Attributes can then be used:
 - as a means to filter or select certain edges and/or vertices for further analysis using graph analytics
 - as weight values for graph algorithms that consider edge weights
 - as data that can be added to graph algorithm results for use by client applications (e.g. GNN training)

	Cust_ID	Zip	Card_num
customers	18	78757	23451
	33	78750	12345

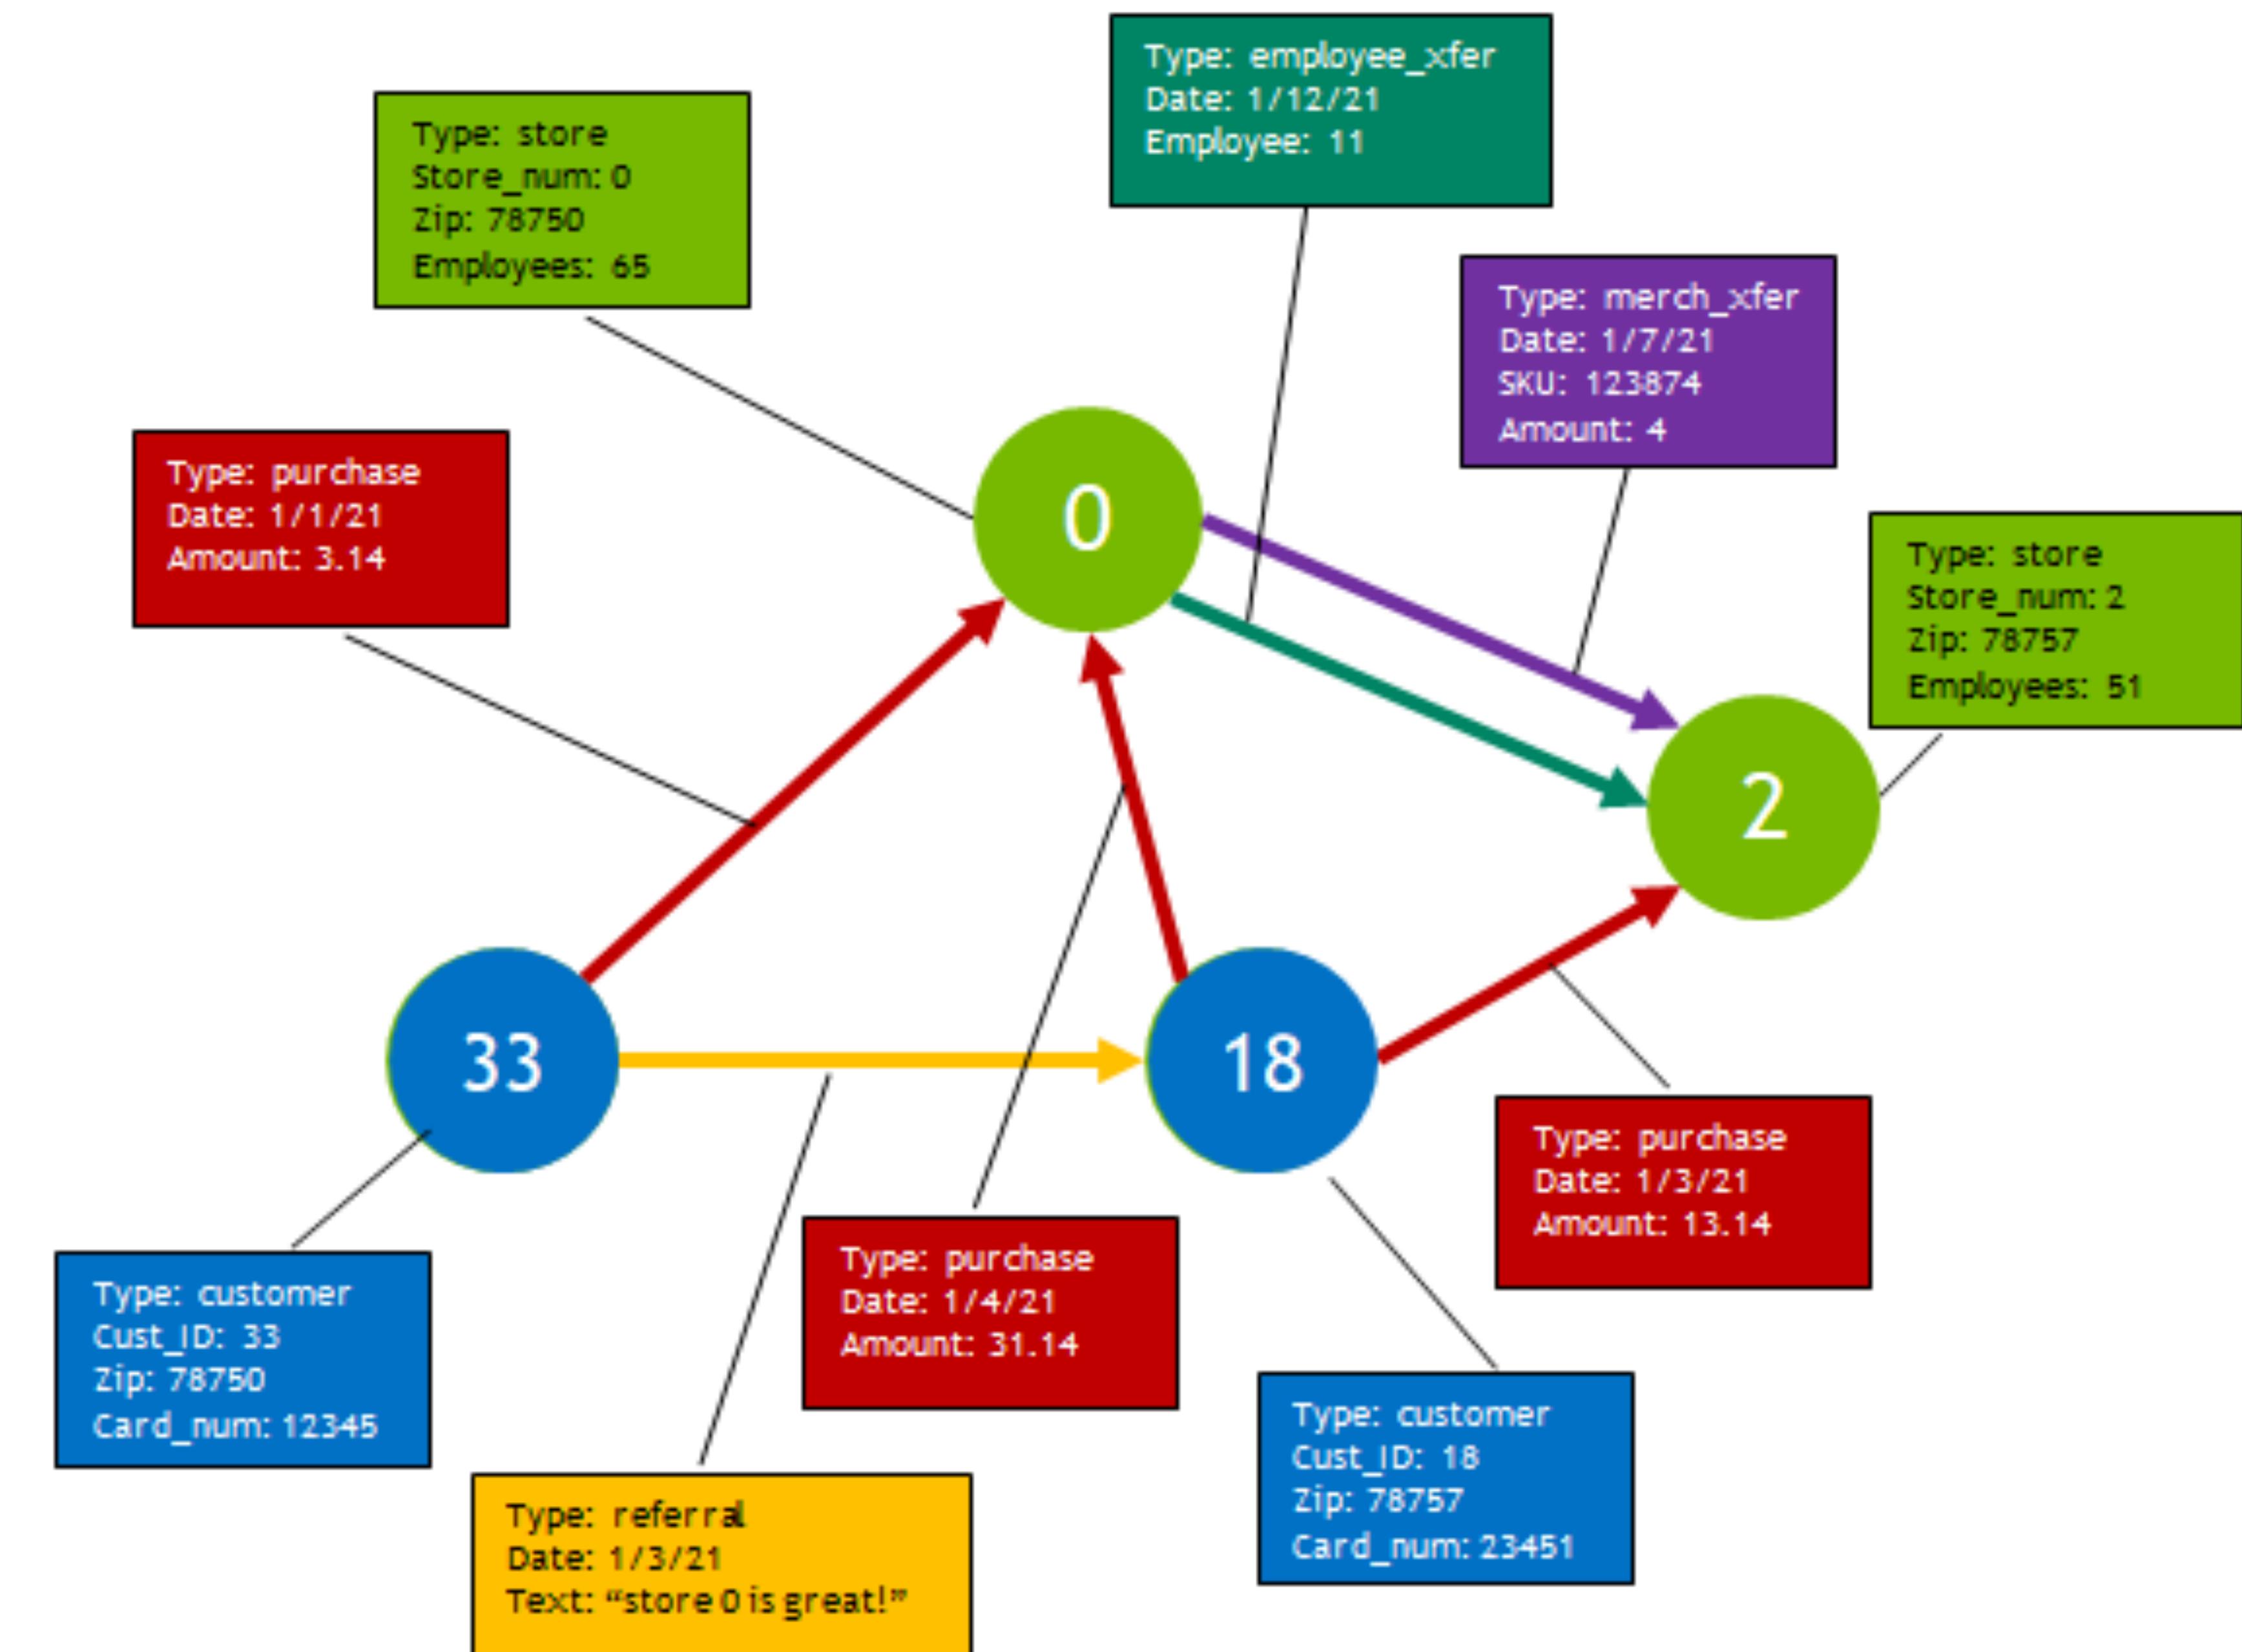
	Store_num	Zip	Employees
stores	0	78750	65
	2	78757	51

	Cust_ID	Store_num	Date	Amount
purchases	33	0	1/1/21	3.14
	18	2	1/3/21	13.14
	18	0	1/4/21	31.14

	Cust_ID_1	Cust_ID_2	Date	Text
referrals	33	18	1/3/21	"store 0 is great!"

	From	To	Date	SKU	Amount
merch xfers	0	2	1/7/21	123874	4

	From	To	Date	Employee
employee xfers	0	2	1/12/21	11



PROPERTY GRAPH

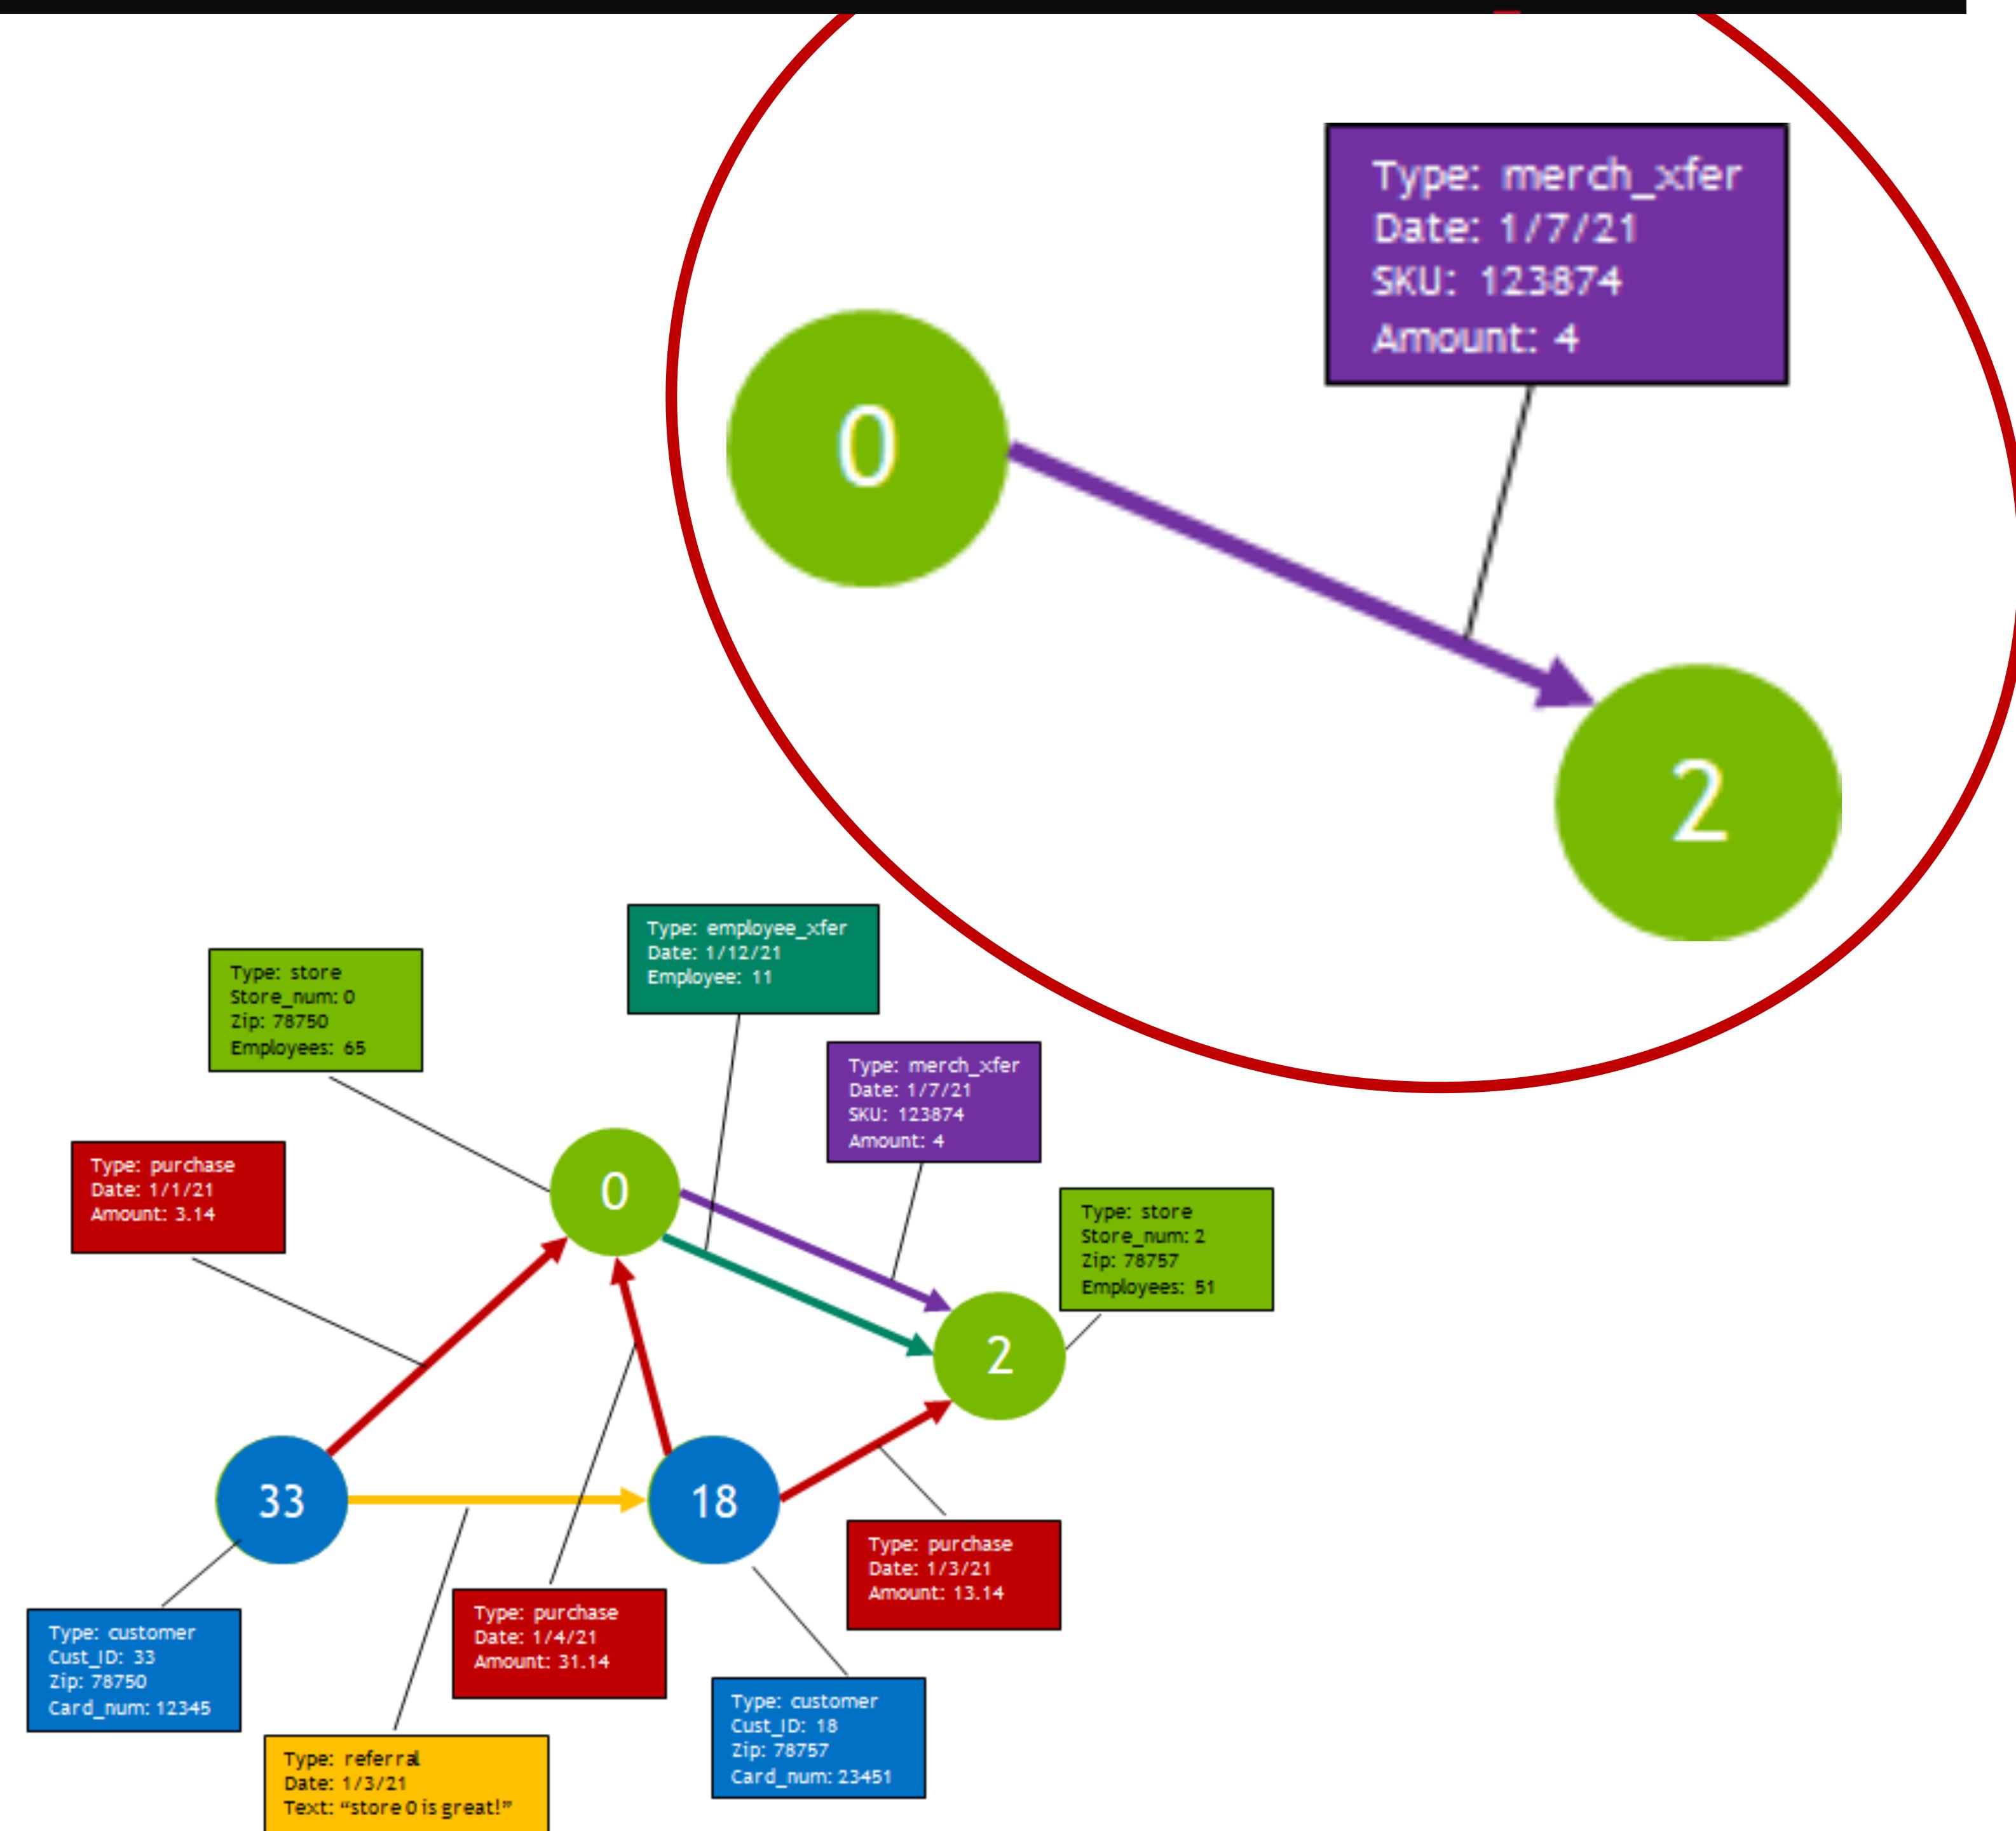
- Example: use a Property Graph to load various datasets as edges and vertices with attributes, use the Property Graph API to extract different graphs based on attributes to run analysis.

```
pG = cugraph.experimental.PropertyGraph()

pG.add_vertex_data(customers_df,
                    type_name="customers",
                    vertex_col_name="Cust_ID")
...
pG.add_edge_data(purchases_df,
                  type_name="purchases",
                  vertex_col_names=("Cust_ID", "Store_num"))
...
selection = pG.select_vertices(f"{pG.type_col_name}=='stores'")
selection += pG.select_edges(f"{pG.type_col_name}=='merch_xfers'")
G = pG.extract_subgraph(selection=selection, edge_weight_property="Amount")

print(G.view_edge_list())
```

```
(rapids) root@17e55ac97b56:/Projects/cugraph/python/cugraph# python pgdemo4.py
      weights   src   dst
0        4.0     0     2
```



PROPERTY GRAPH

- Example: use a Property Graph to load the Zachary Karate Club dataset, use Louvain to find the two primary partitions, use Pagerank to find the top 3 influential vertices in each partition.

```
import cudf
import cugraph
from cugraph.experimental import PropertyGraph

# Read edgelist data into a DataFrame, load into PropertyGraph as edge data.
df = cudf.read_csv("karate_directed.csv",
                    delimiter=" ",
                    dtype=["int32", "int32", "float32"],
                    header=None)

pG = PropertyGraph()
pG.add_edge_data(df, vertex_col_names=("0", "1"))

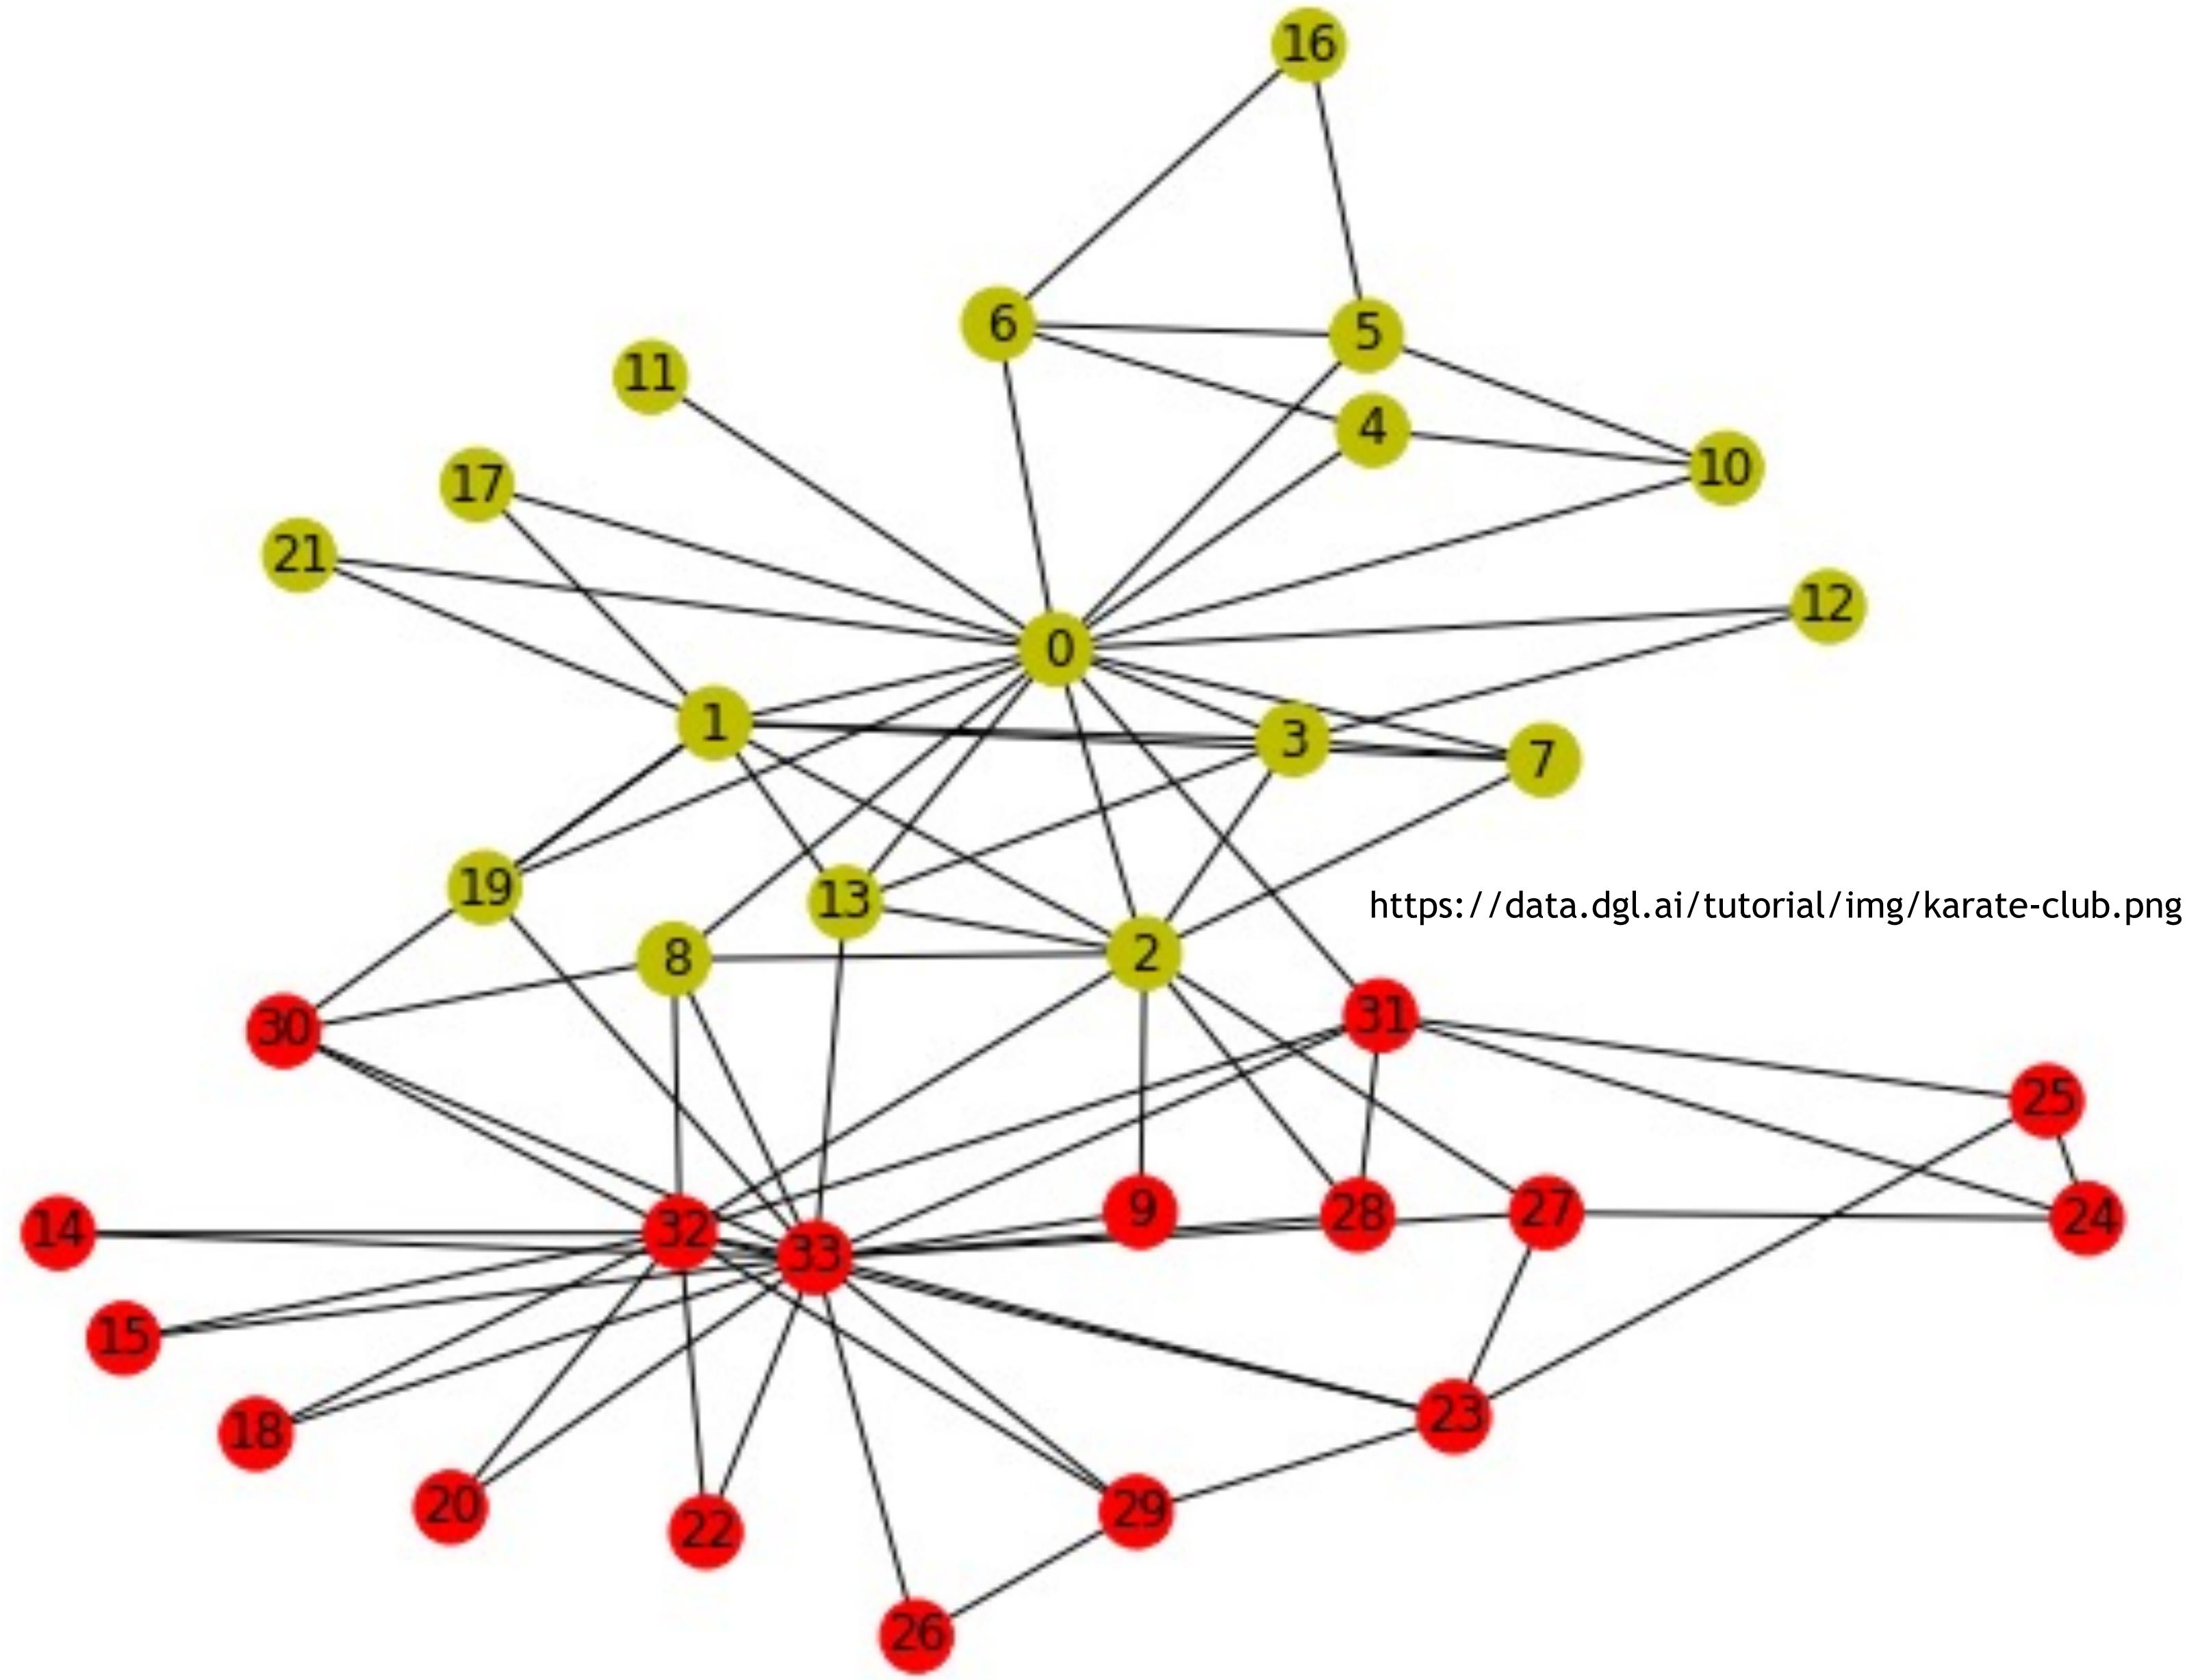
# Run Louvain to get the partition number for each vertex.
# Set resolution accordingly to identify two primary partitions.
(partition_info, _) = cugraph.louvain(pG.extract_subgraph(), resolution=0.6)

# Add the partition numbers back to the PropertyGraph as vertex properties.
pG.add_vertex_data(partition_info, vertex_col_name="vertex")

# Use the partition properties to extract a Graph for each partition.
G0 = pG.extract_subgraph(selection=pG.select_vertices("partition == 0"))
G1 = pG.extract_subgraph(selection=pG.select_vertices("partition == 1"))

# Run pagerank on each graph, print results.
pageranks0 = cugraph.pagerank(G0)
pageranks1 = cugraph.pagerank(G1)
print(pageranks0.sort_values(by="pagerank", ascending=False).head(3))
print(pageranks1.sort_values(by="pagerank", ascending=False).head(3))
```

```
(rapids) /cugraph/python/cugraph# python pgdemo3.py
pagerank  vertex
0  0.192222      0
12  0.108455     1
6   0.084853     2
pagerank  vertex
1  0.191775    33
2  0.150855    32
3  0.072723    31
```

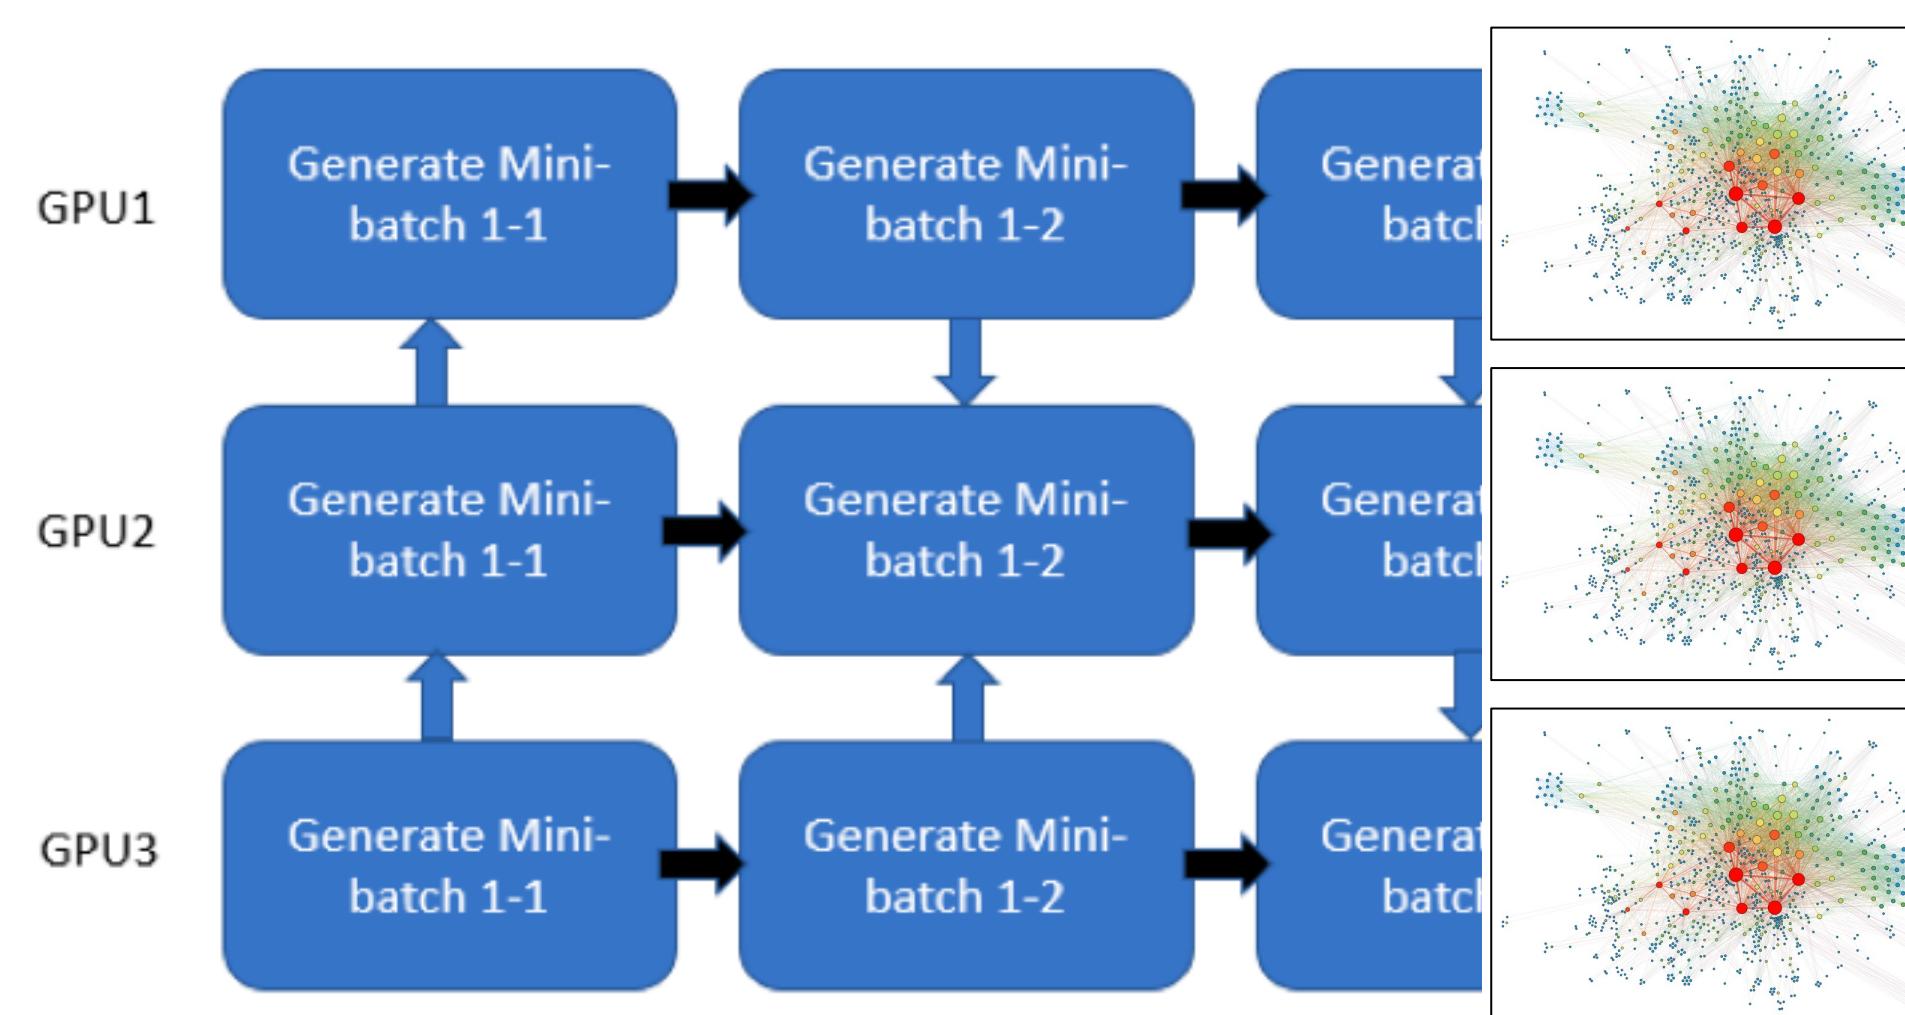




Graph-as-a-Service - GaaS

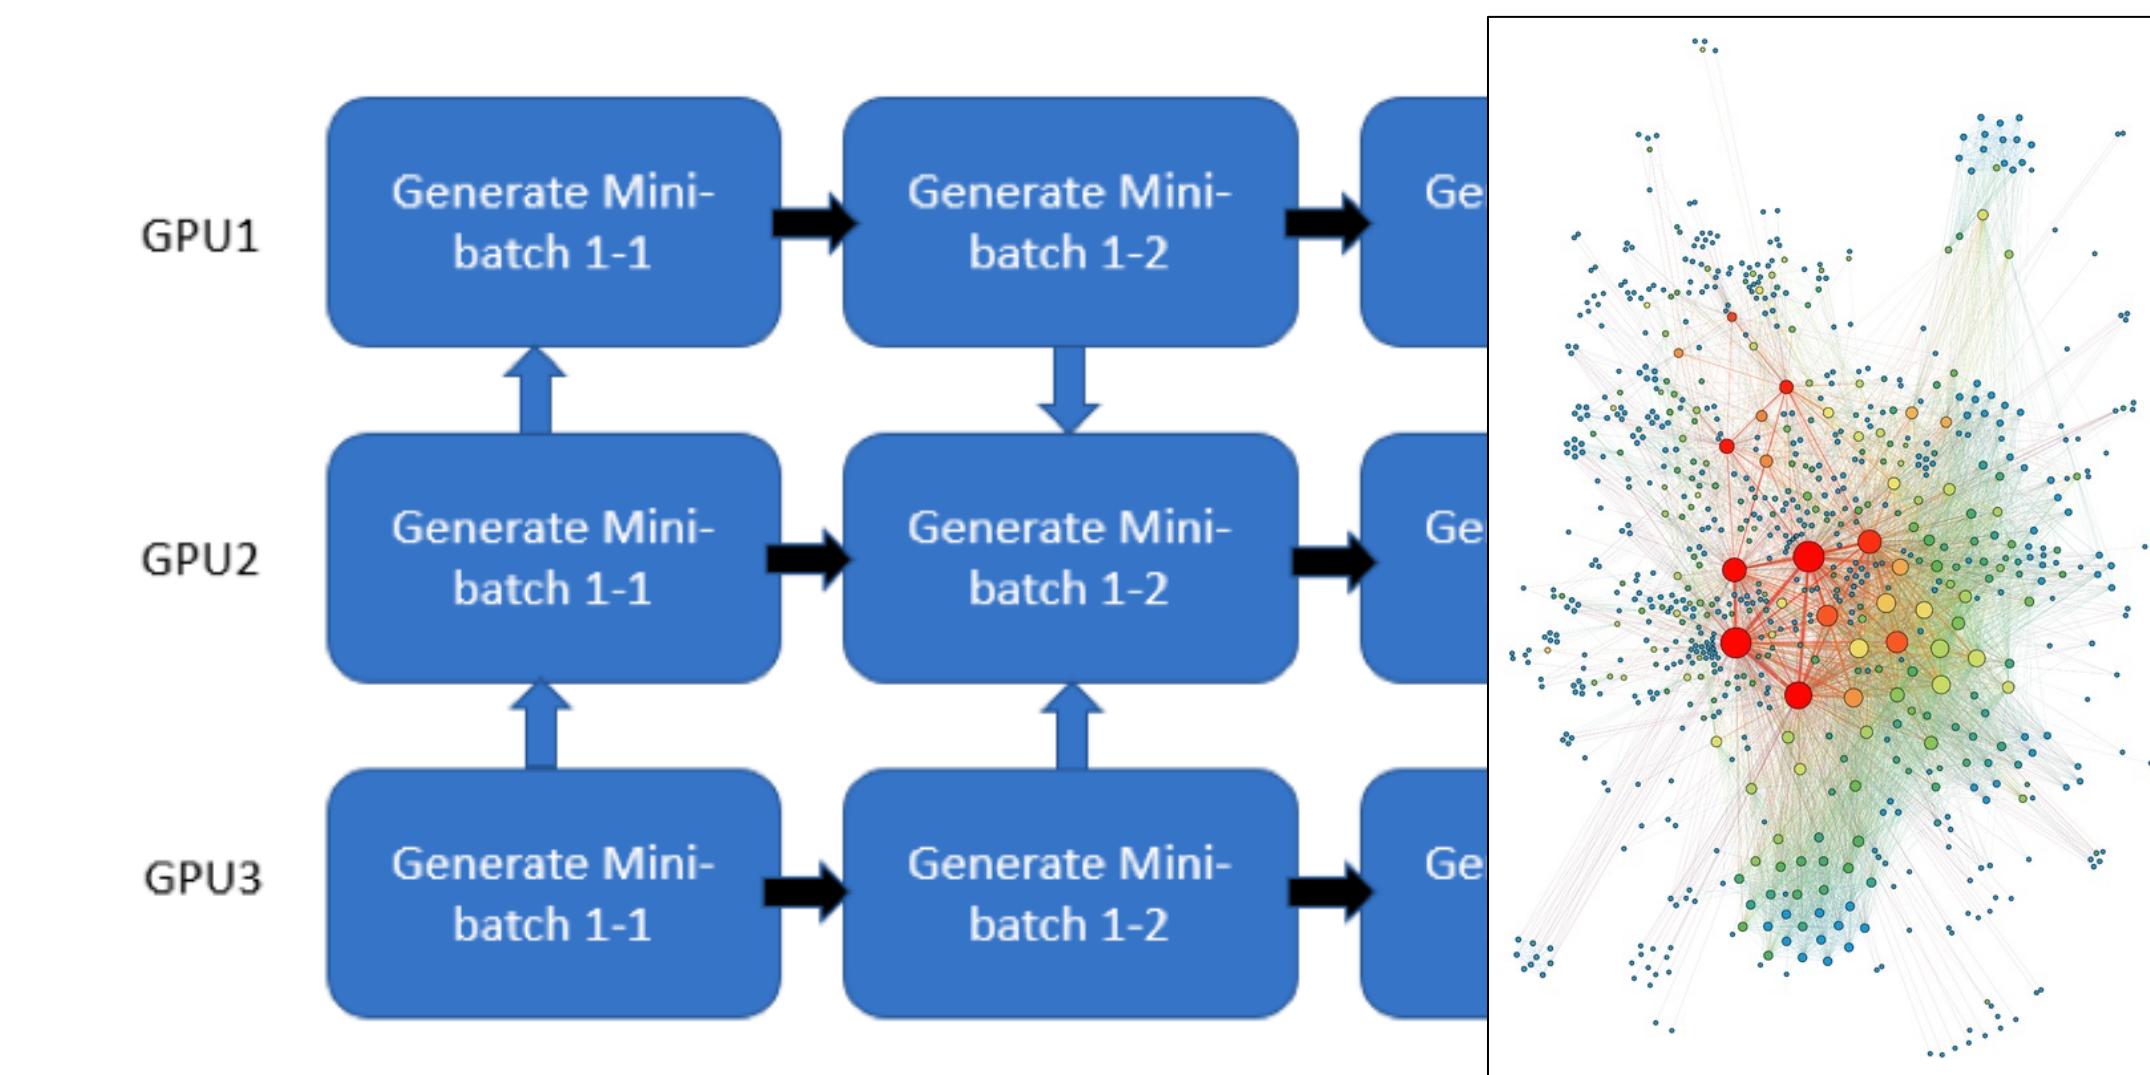
GAAS - GRAPH-AS-A-SERVICE

Background in a GNN Environment

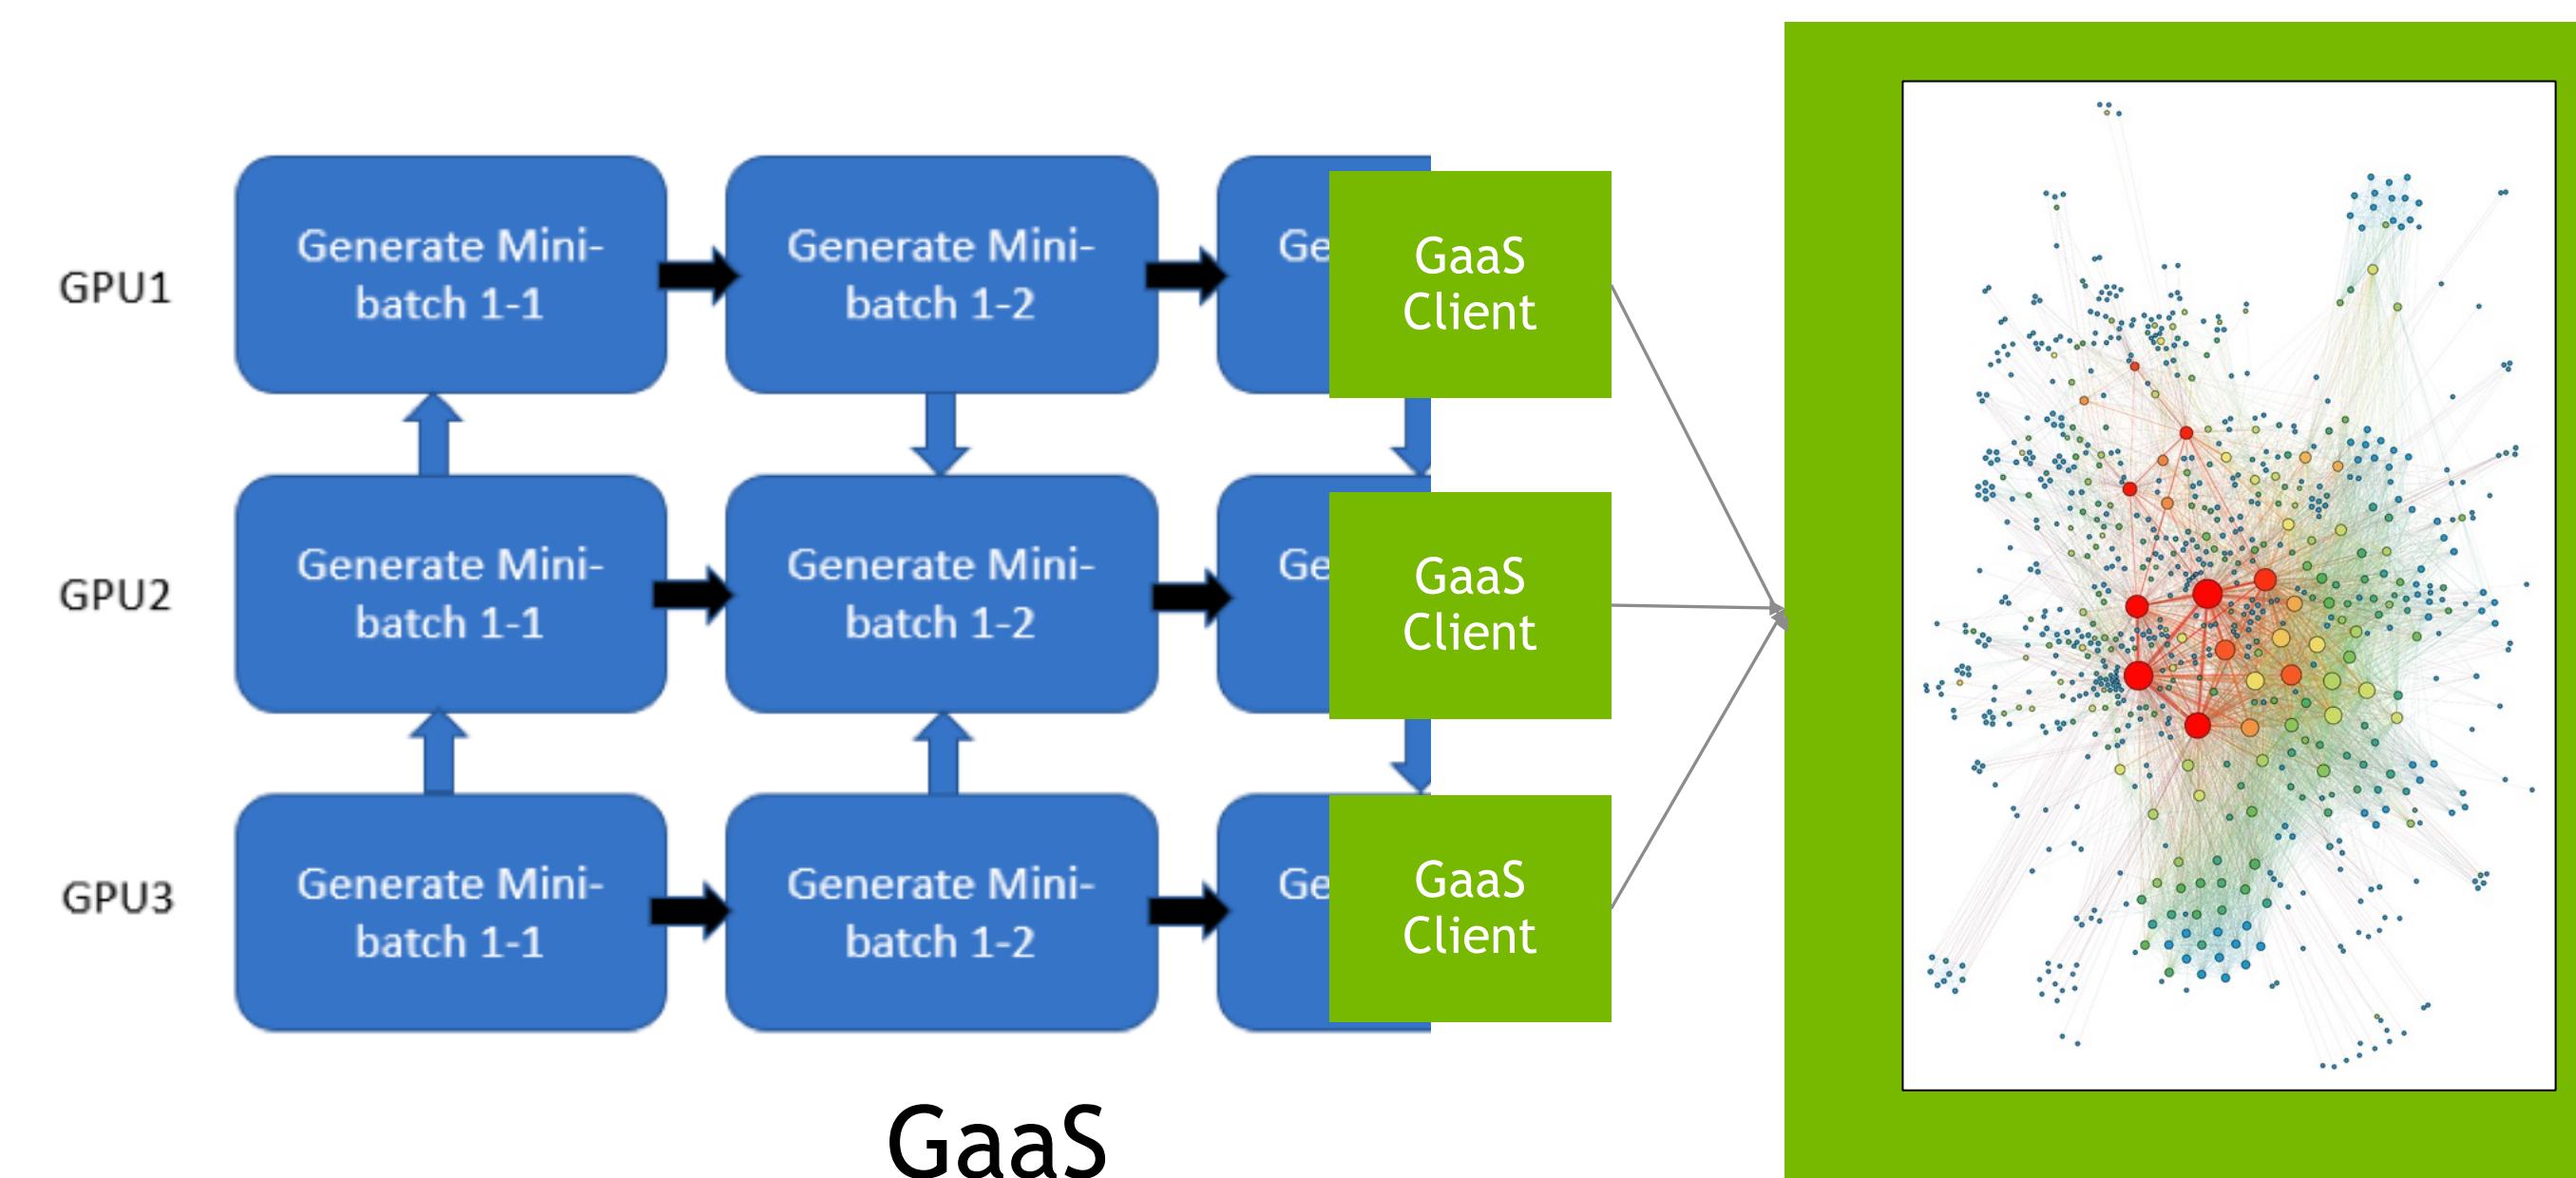


Data parallel
No issues with small graphs

... BUT ...
What happens when the graph
gets large?



Data parallel no longer works
Resource contention



Addresses issues by separating
Graph from application

GaaS

GAAS - GRAPH-AS-A-SERVICE

Background and goals

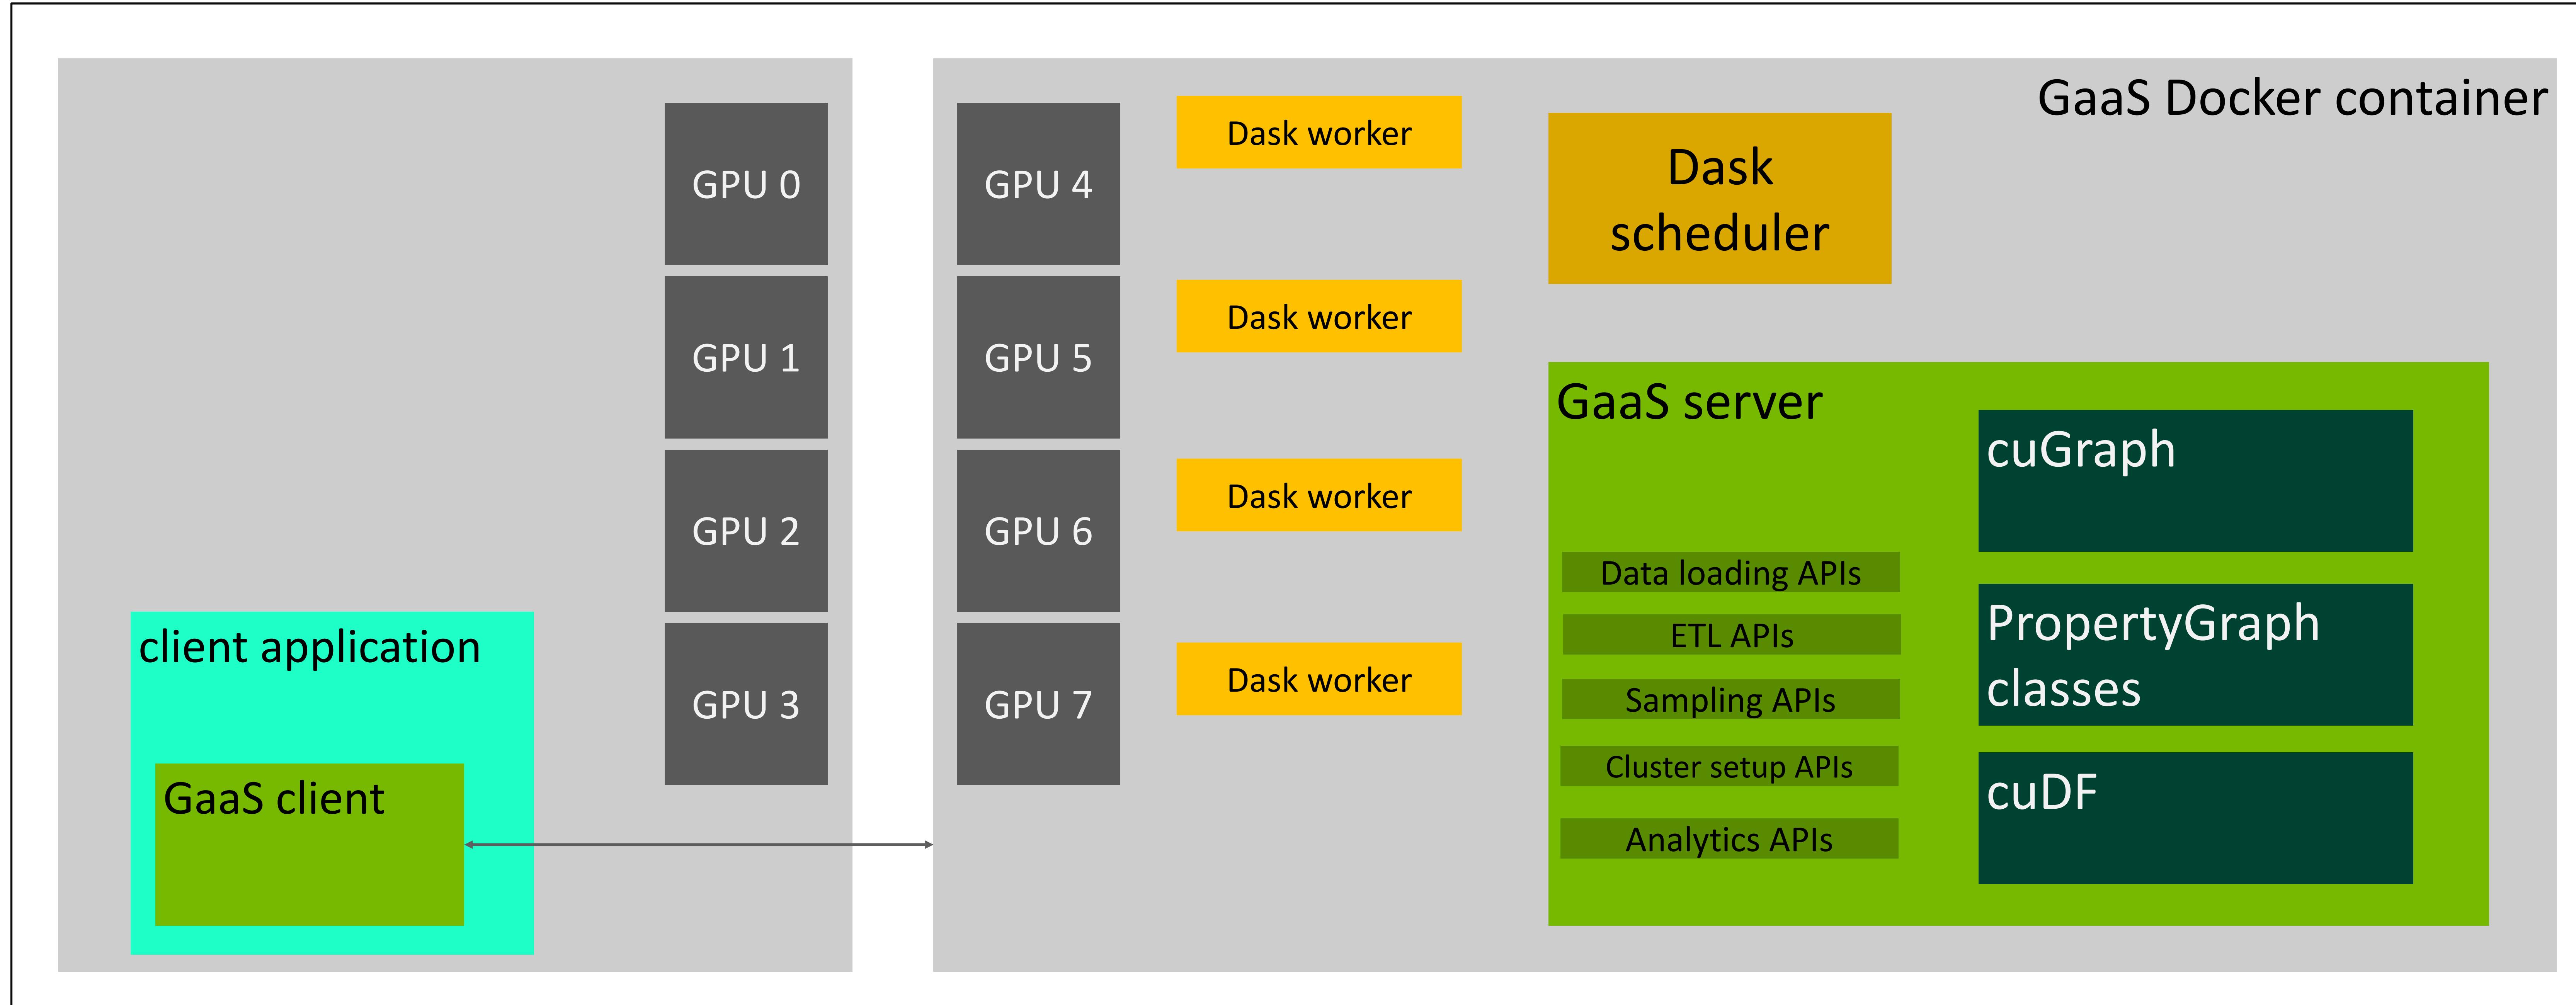
- Separate large graph management and analytic code from application code
 - The application, like GNN code, should be isolated from the details of cuGraph graph management, dedicated multi-node/multi-GPU setup, feature storage and retrieval, etc.
 - Scaling from single GPU (SG), to multi-GPU (MG), to multi-node/multi-GPU (MNMG) should not require changes to the graph integration code
- Support multiple concurrent clients/processes/threads sharing one or more graphs
 - No need for each client to have access to graph data files, perform ETL, etc.
 - Simplify concurrent programming - synchronization, batch processing details, etc. are implemented server-side
 - The GNN user should be able to easily partition and isolate hardware resources between graph analytics and training
 - The application/GNN code should be able to prefetch graph samples while training without resource contention
- Simplify MNMG deployments
 - Docker images contain all required, compatible packages
 - Clients use the same GaaS APIs for both SG and MG

NOTE: GaaS is still a work-in-progress and changes frequently as we add functionality and improve performance.



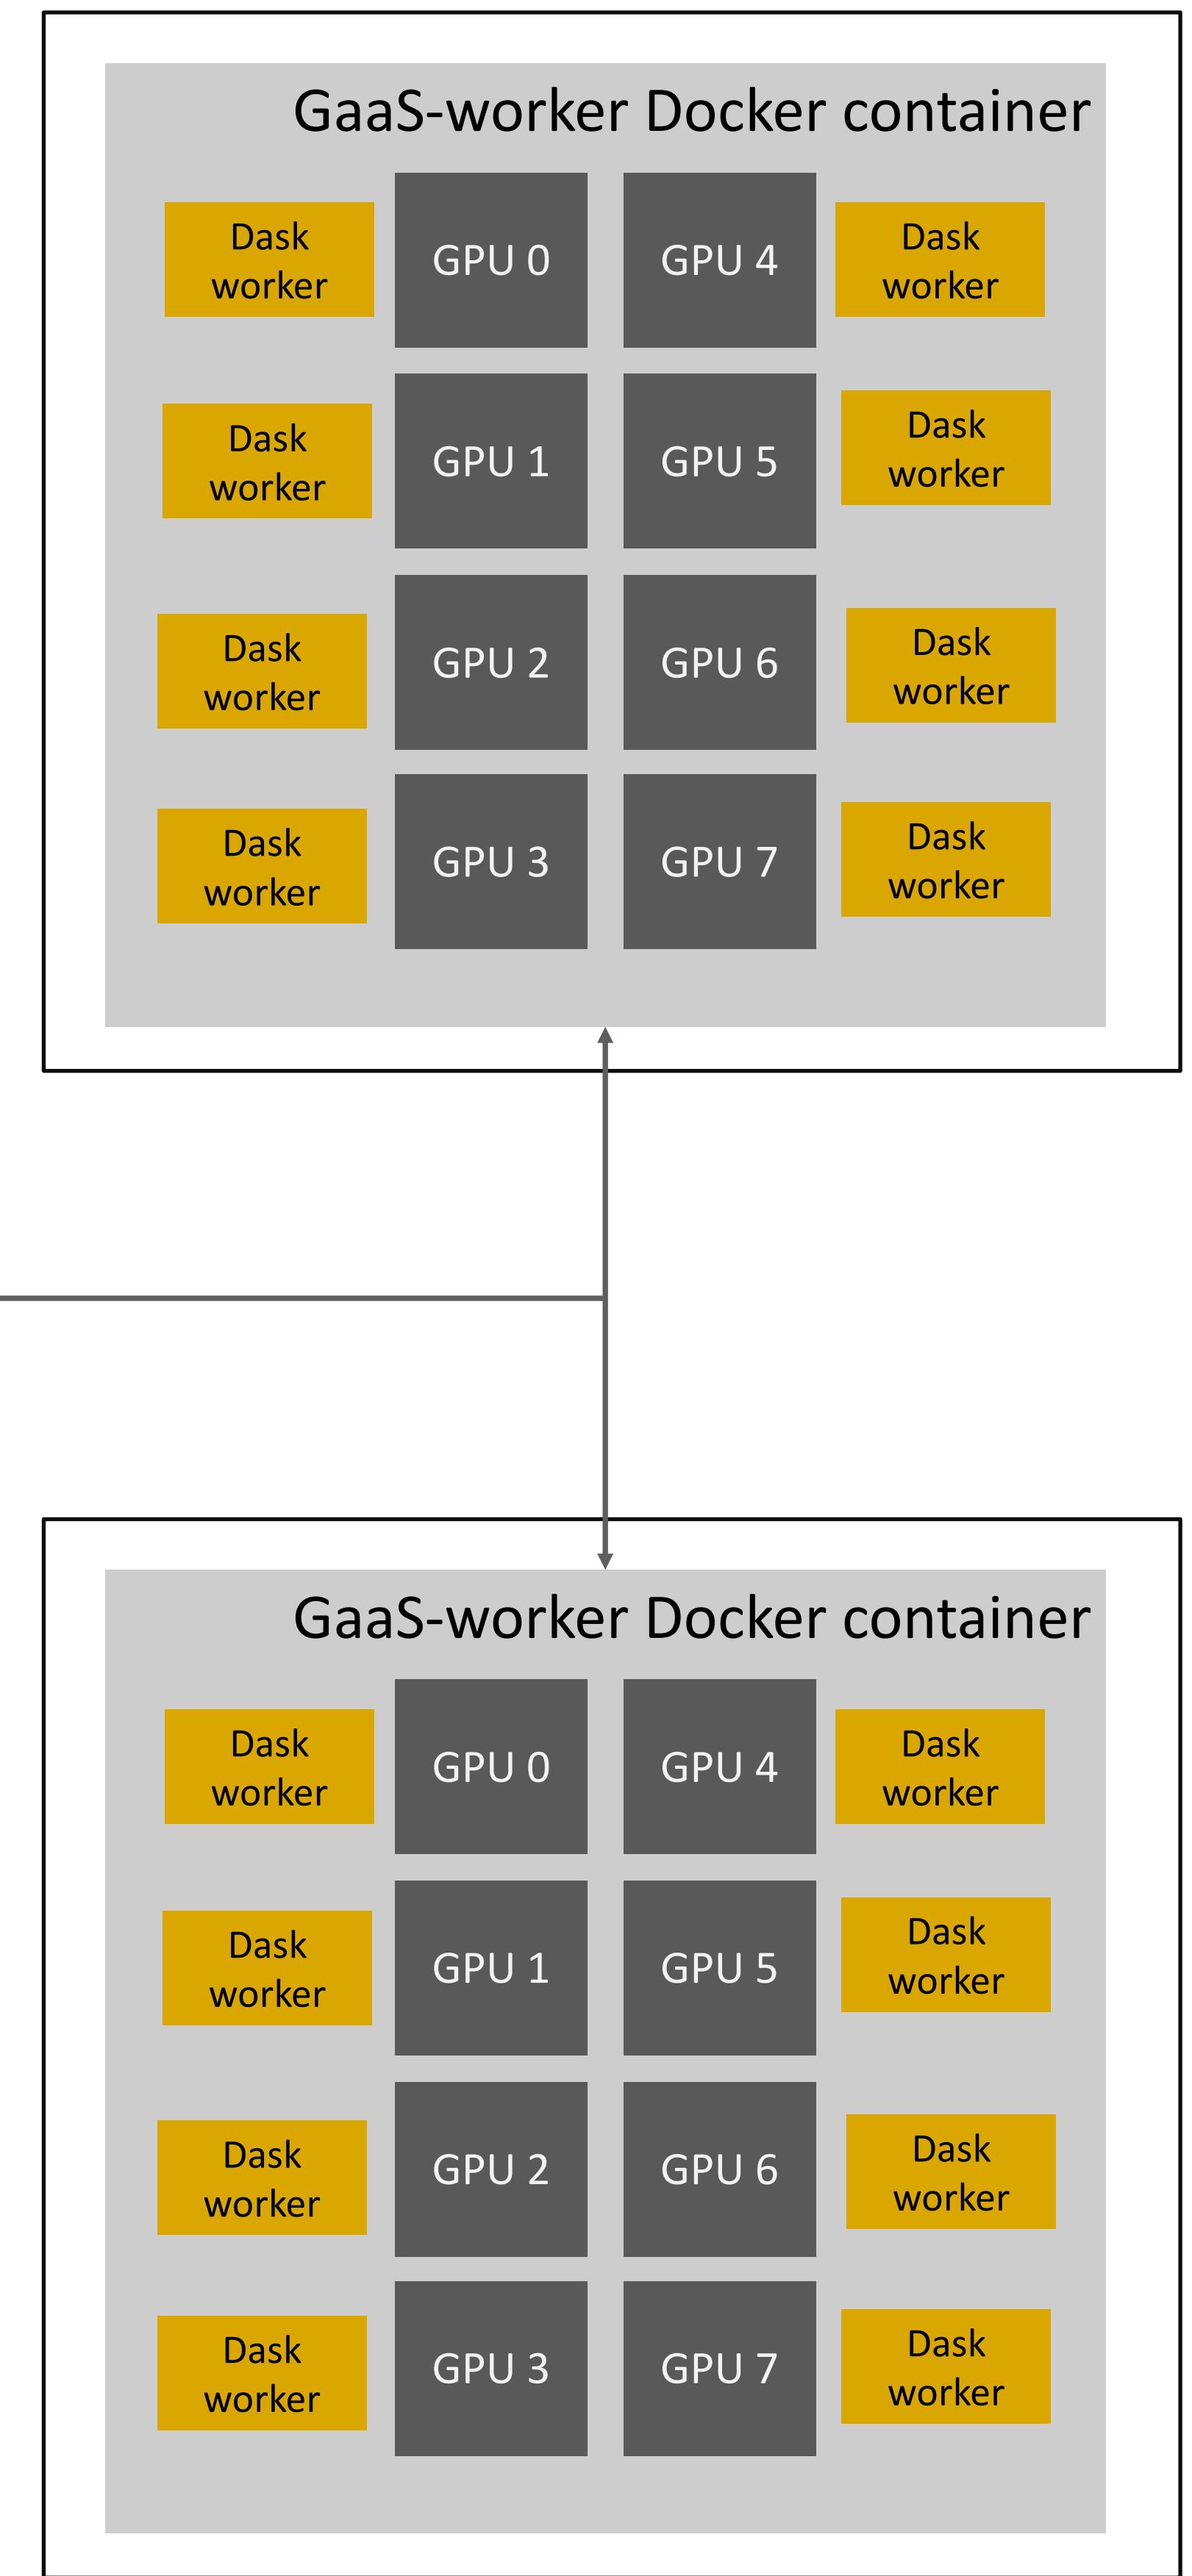
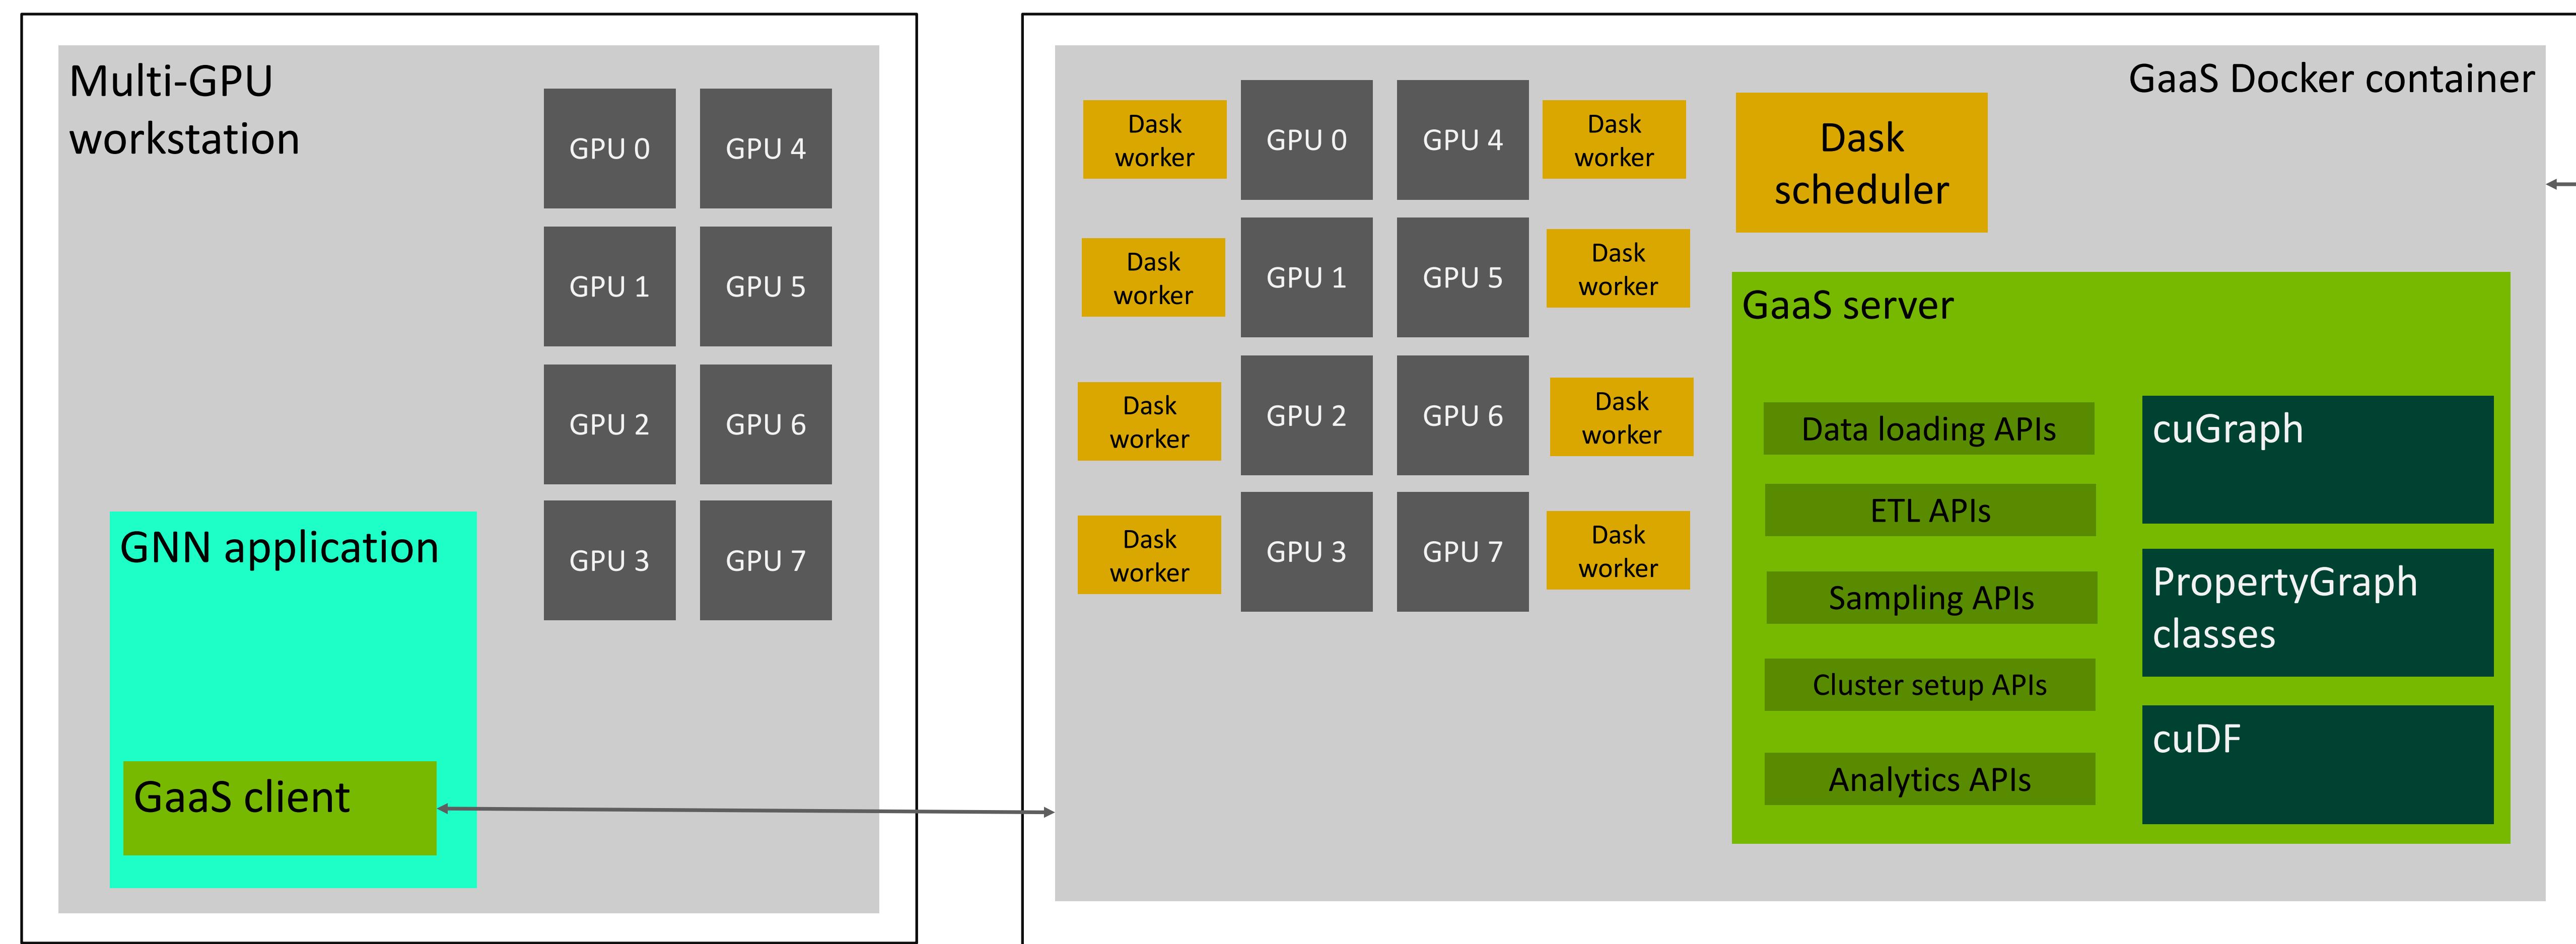
GAAS - GRAPH-AS-A-SERVICE

multi-GPU deployment



GAAS - GRAPH-AS-A-SERVICE

multi-node multi-GPU deployment

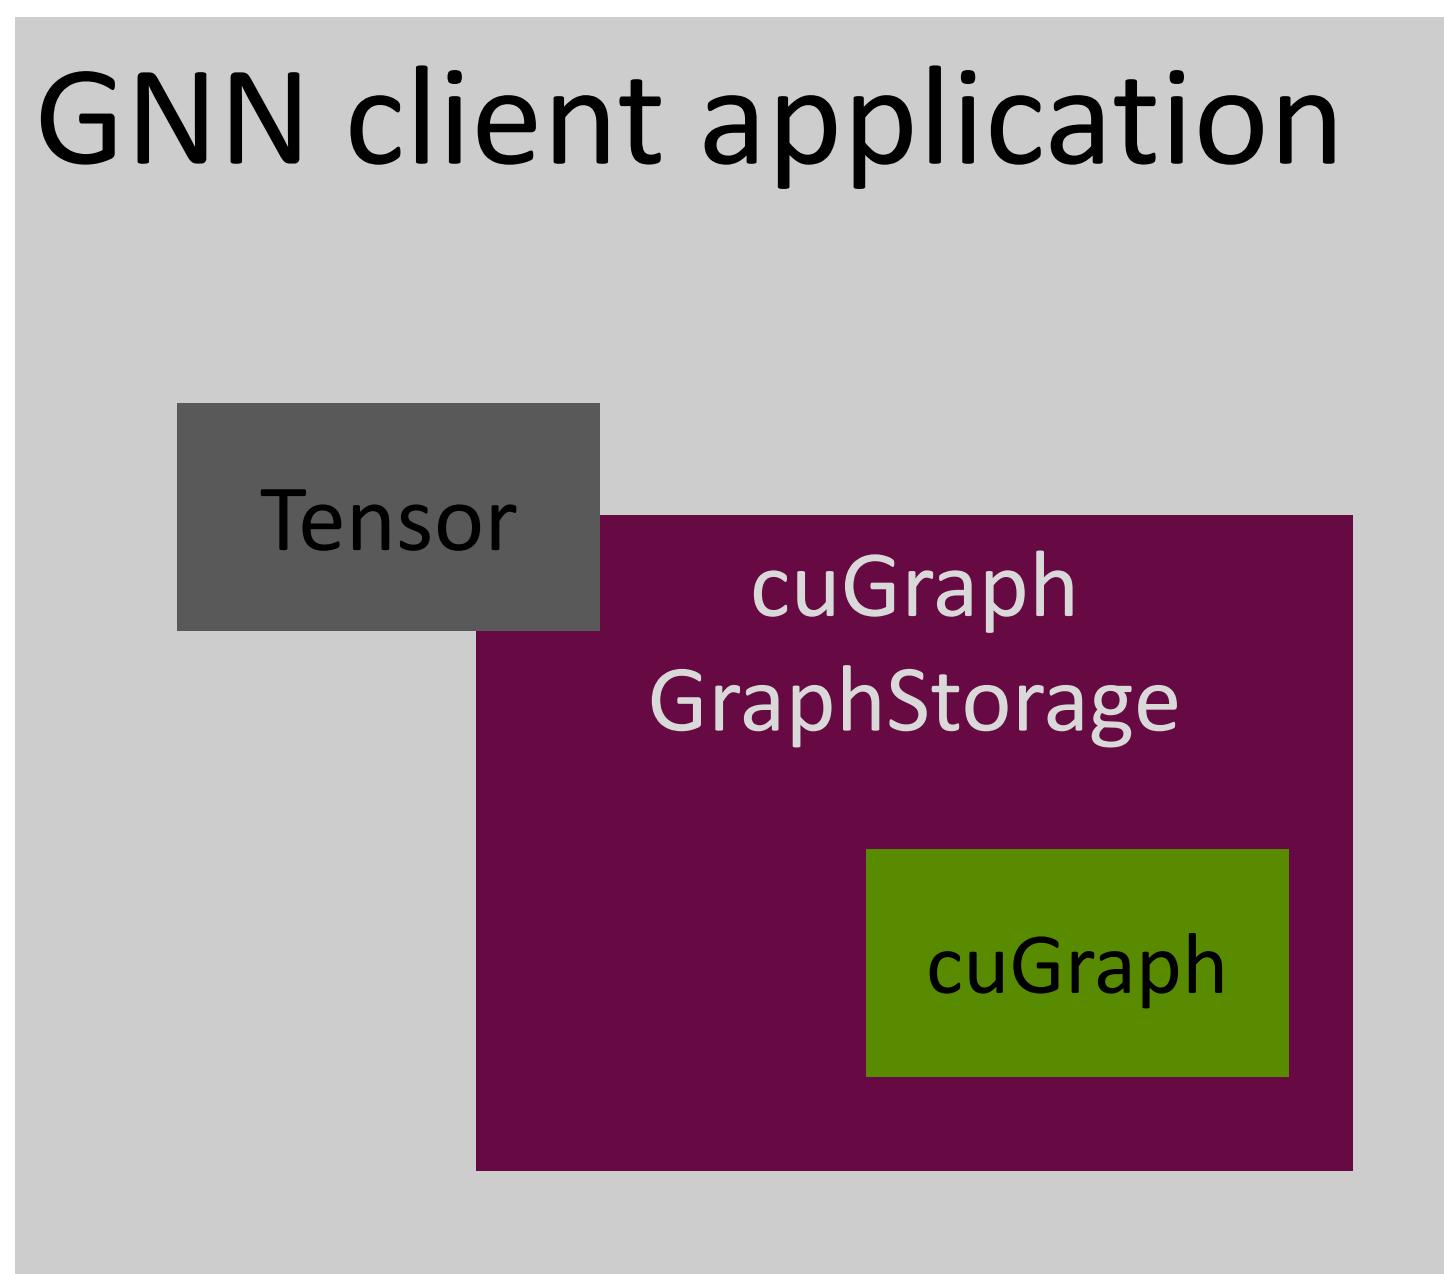


GAAS - DESIGN PATTERNS

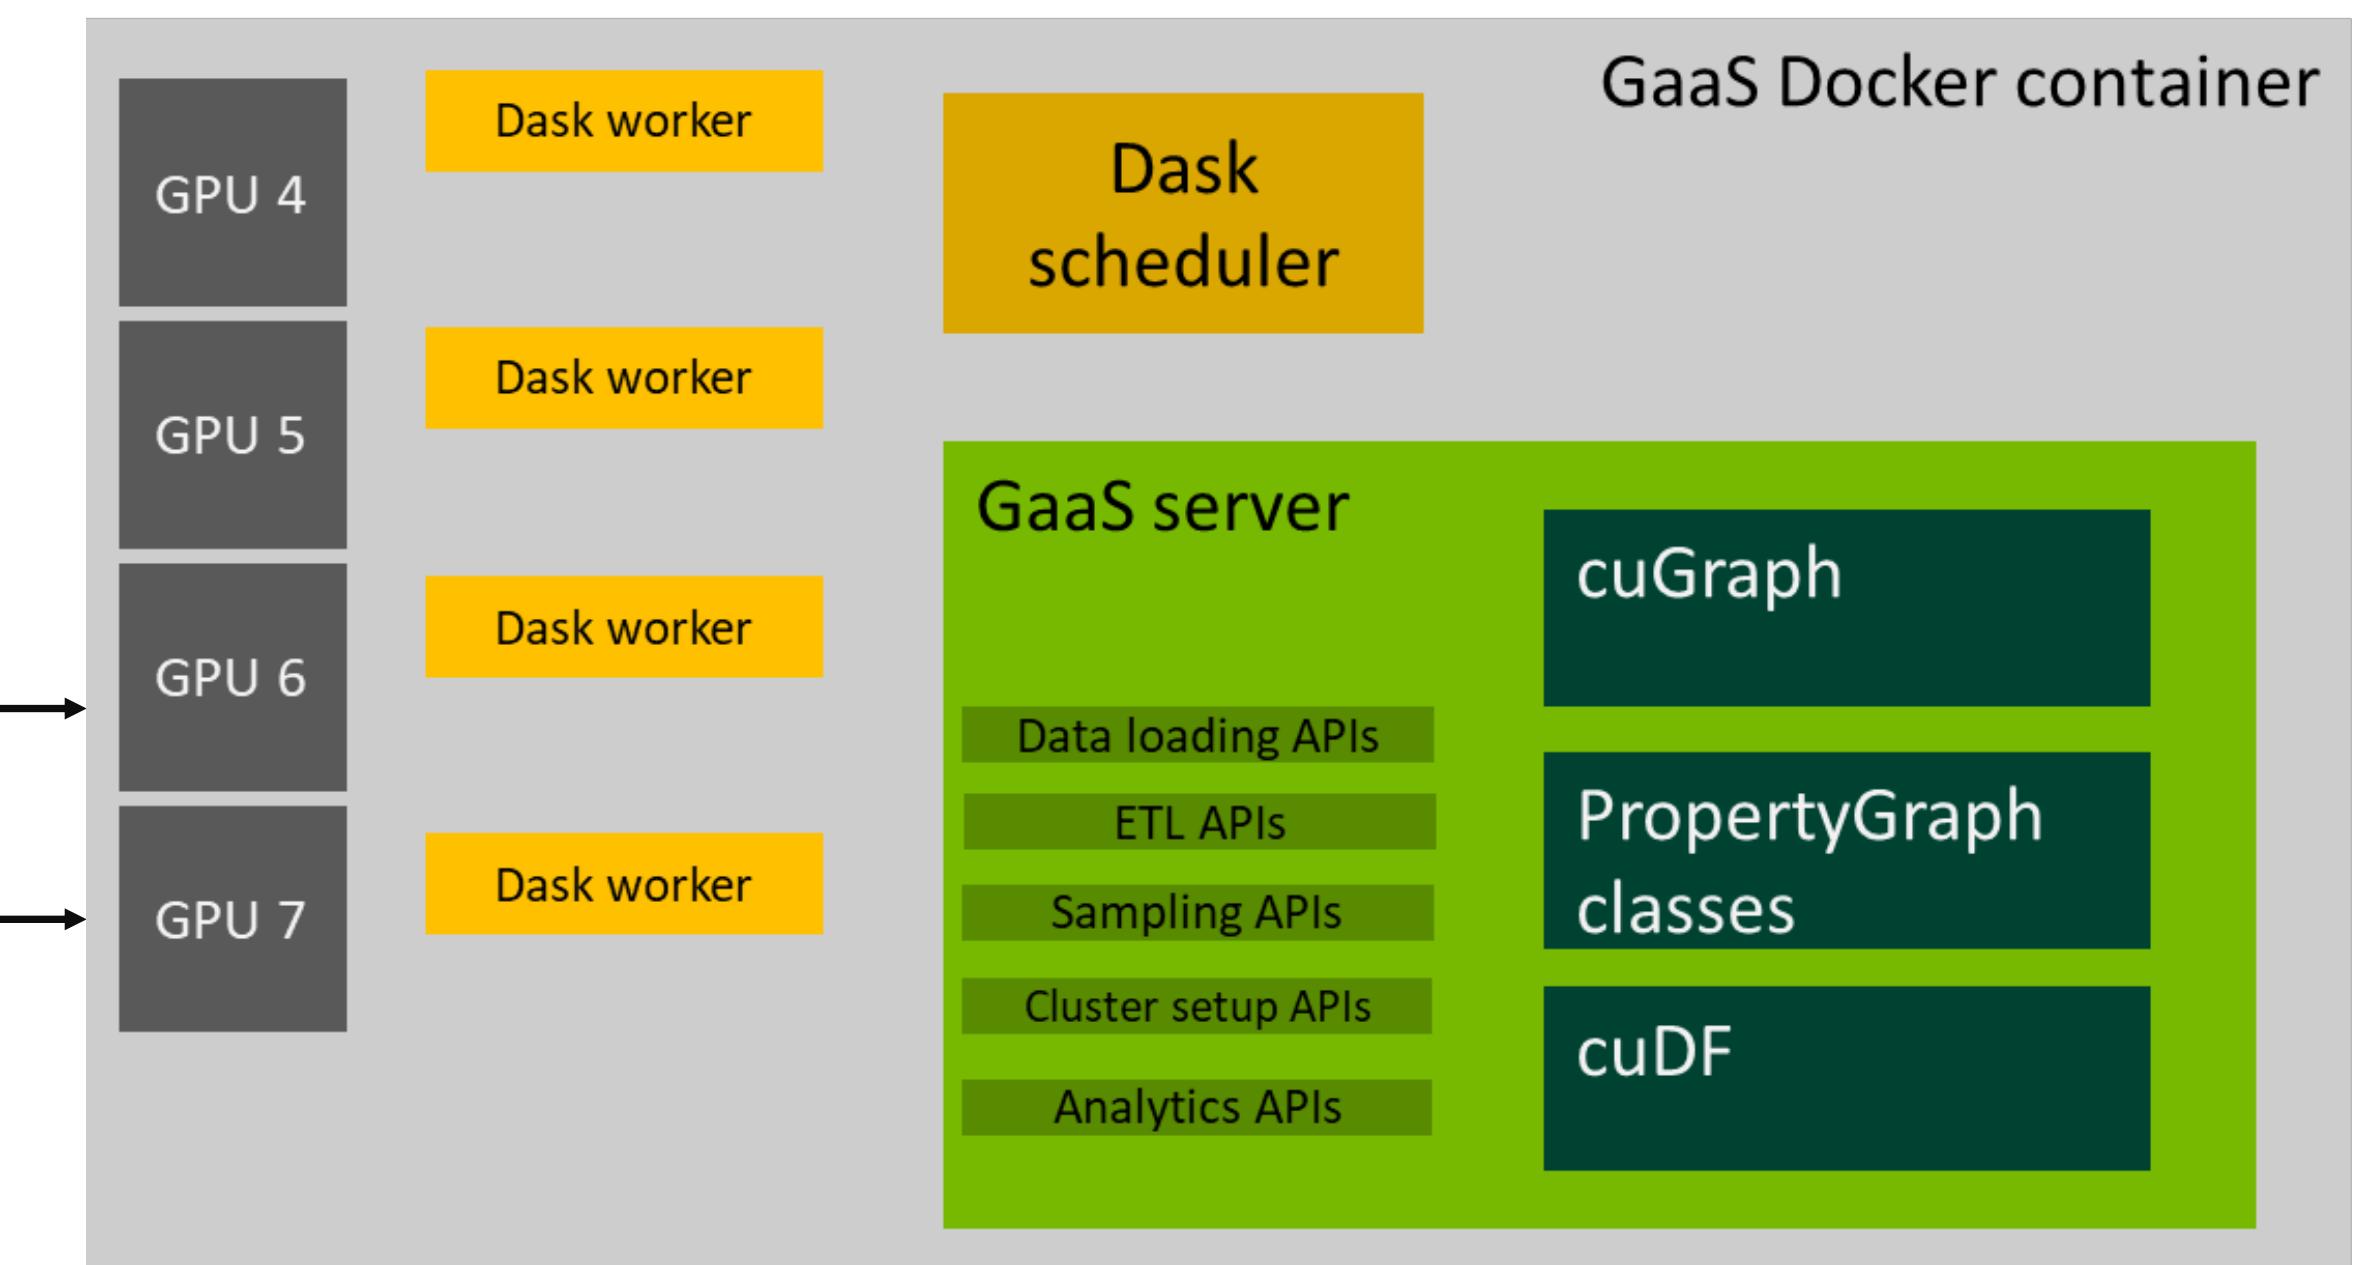
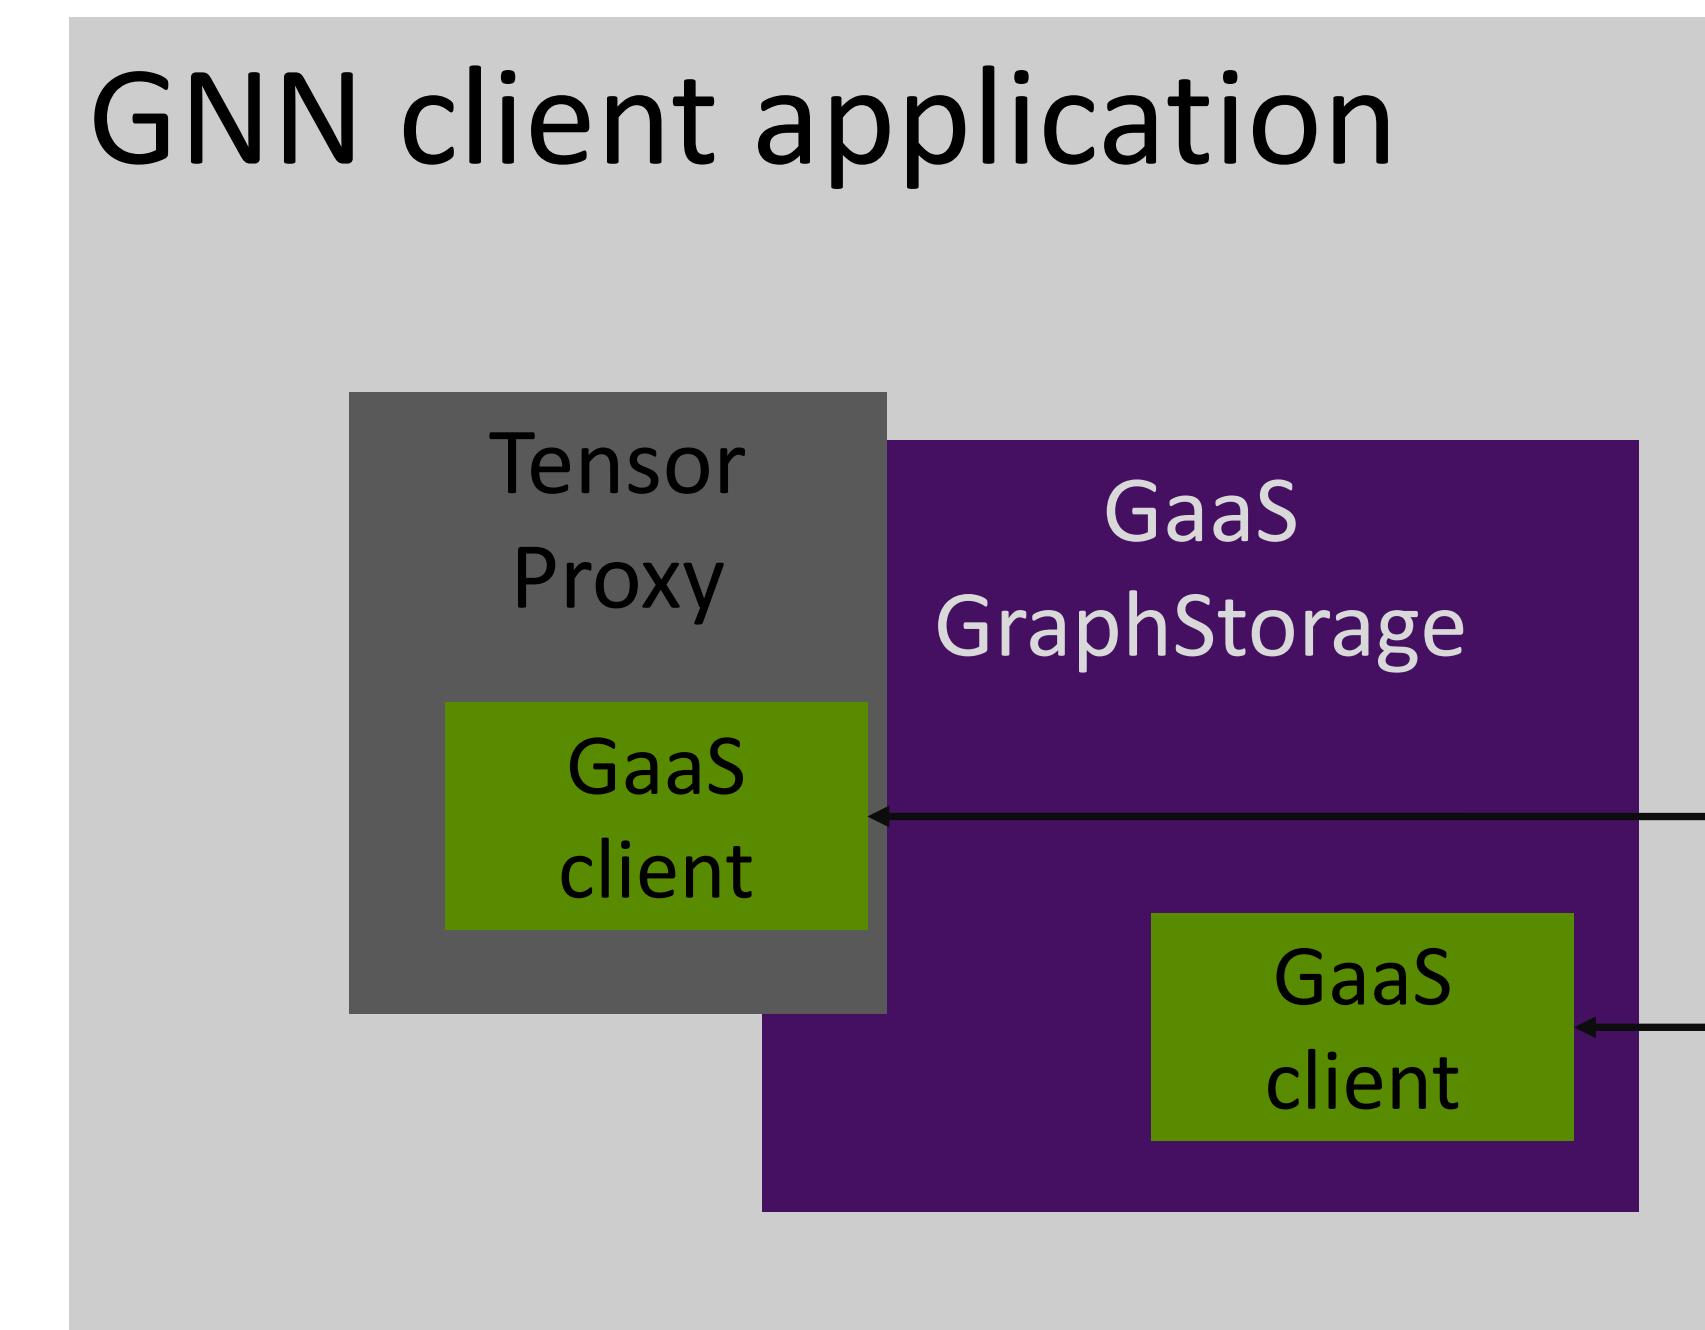
Proxy objects

- Support both cuGraph “local” and “GaaS” modes without code changes
 - Objects for either mode have the same API (duck-typed)
 - “local” mode object accesses cuGraph directly
 - “GaaS” mode object accesses the server via the client API and manages remote object lifecycle using RAII-style implementation
- Simplifies GNN integration by encapsulating integration details

cuGraph “local”



cuGraph “GaaS”



GAAS - DESIGN PATTERNS

Proxy objects

- Use the client to load a graph with vertex and edge data from CSV files.

```
>>> from gaas_client import GaasClient
>>>
>>> client = GaasClient()
>>> G = GaasGraph(client)
>>>
>>> # Add vertex and edge data to the graph on the server
>>> client.load_csv_as_vertex_data("/data/vertex_data.csv",
...                                ...,
...                                graph_id=graph_id)
>>> client.load_csv_as_edge_data("/data/edge_data.csv",
...                                ...,
...                                graph_id=graph_id)
```

- Retrieve the vertex features for vertices 100, 101, 102
- This requires the application code to maintain the client object and graph ID for future calls.

...or

- Construct a TensorProxy using the client object and graph ID for vertex data.
- This tensor can be passed to other application code as-is to access vertex data.

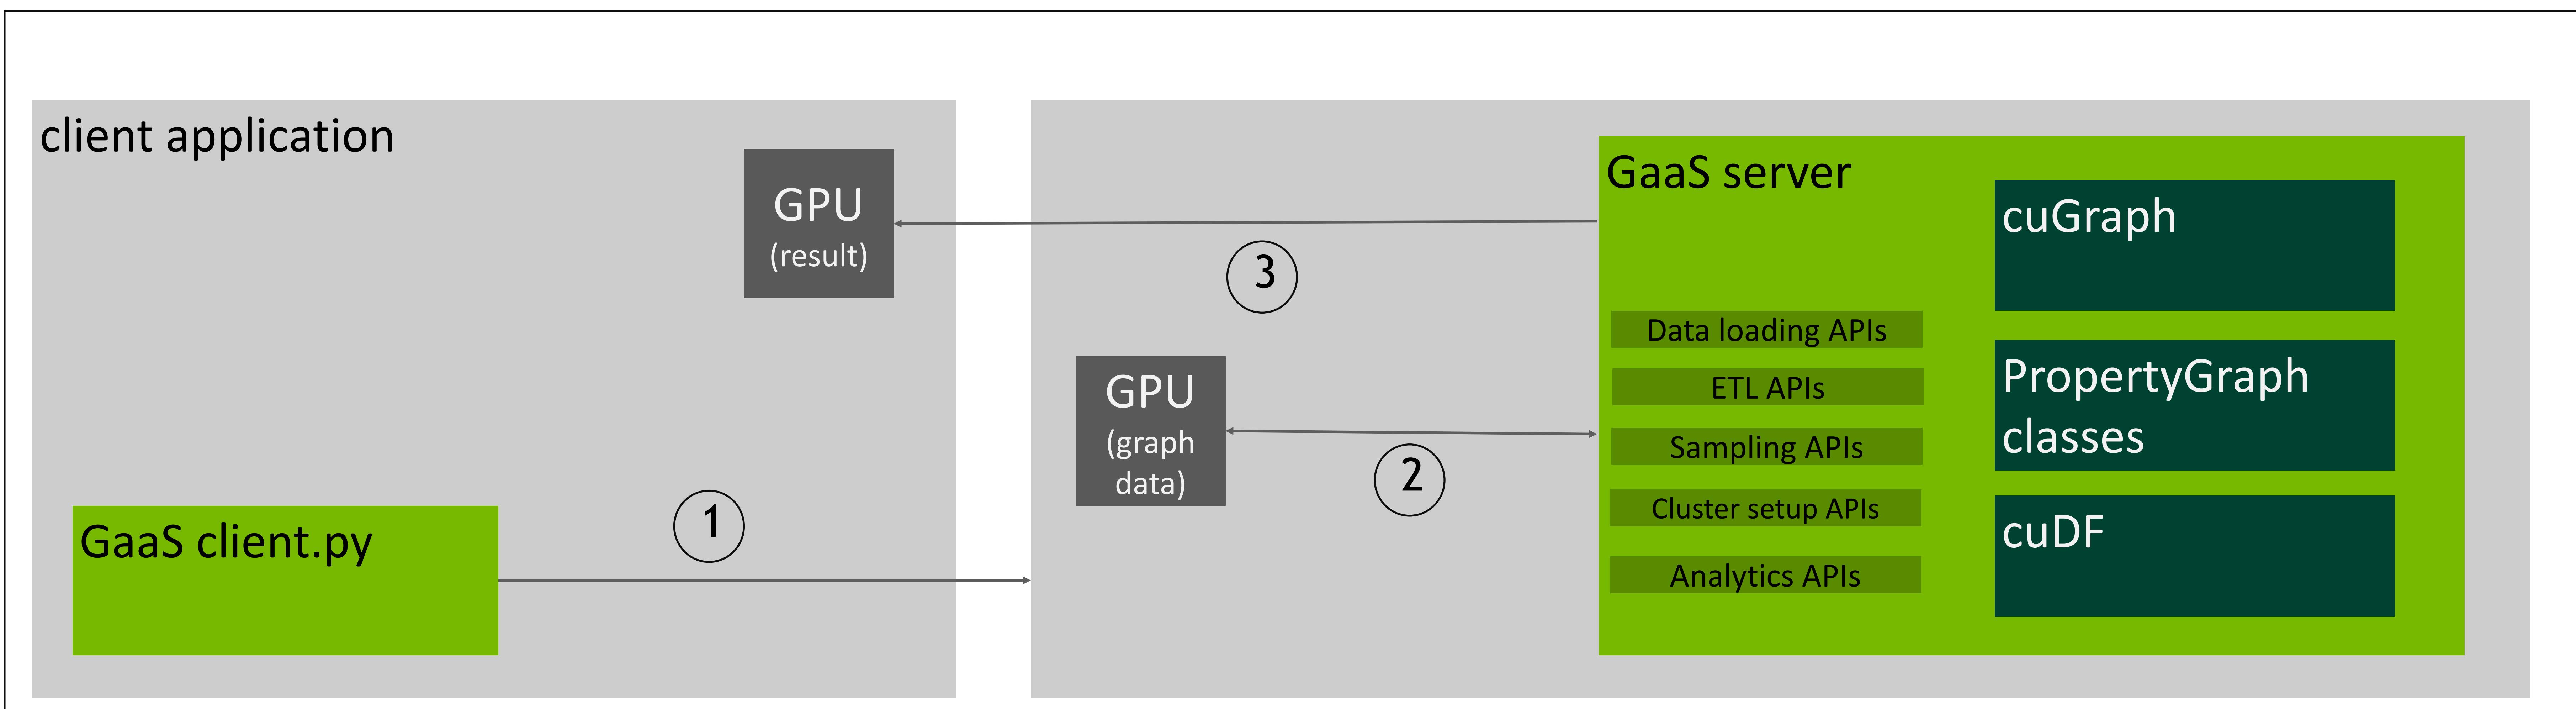
```
>>> client.get_graph_vertex_data(vertex_ids=[100,101,102], graph_id=0)
array([[0, 11223, 8793],
       [1, 11289, 2399],
       [2, 11267, 1500]], dtype='int64')
```

```
>>> vert_features = TensorProxy(client, graph_id, category="vertex")
>>> vert_features[[100,101,102]]
tensor([[    0, 11223, 8793],
        [    1, 11289, 2399],
        [    2, 11267, 1500]])
```

GAAS - UPCOMING ENHANCEMENTS

Performance improvements for clients with GPUs

- Allow GaaS to populate client GPU with result directly
 1. Client makes request
 2. GPU graph analytic performed on graph in server GPU memory
 3. Result is written to client GPU for use by client-side tasks (training, etc.)
- Eliminates the server device-to-host copy, serialization, and transfer steps normally needed when returning results back to the client over RPC.



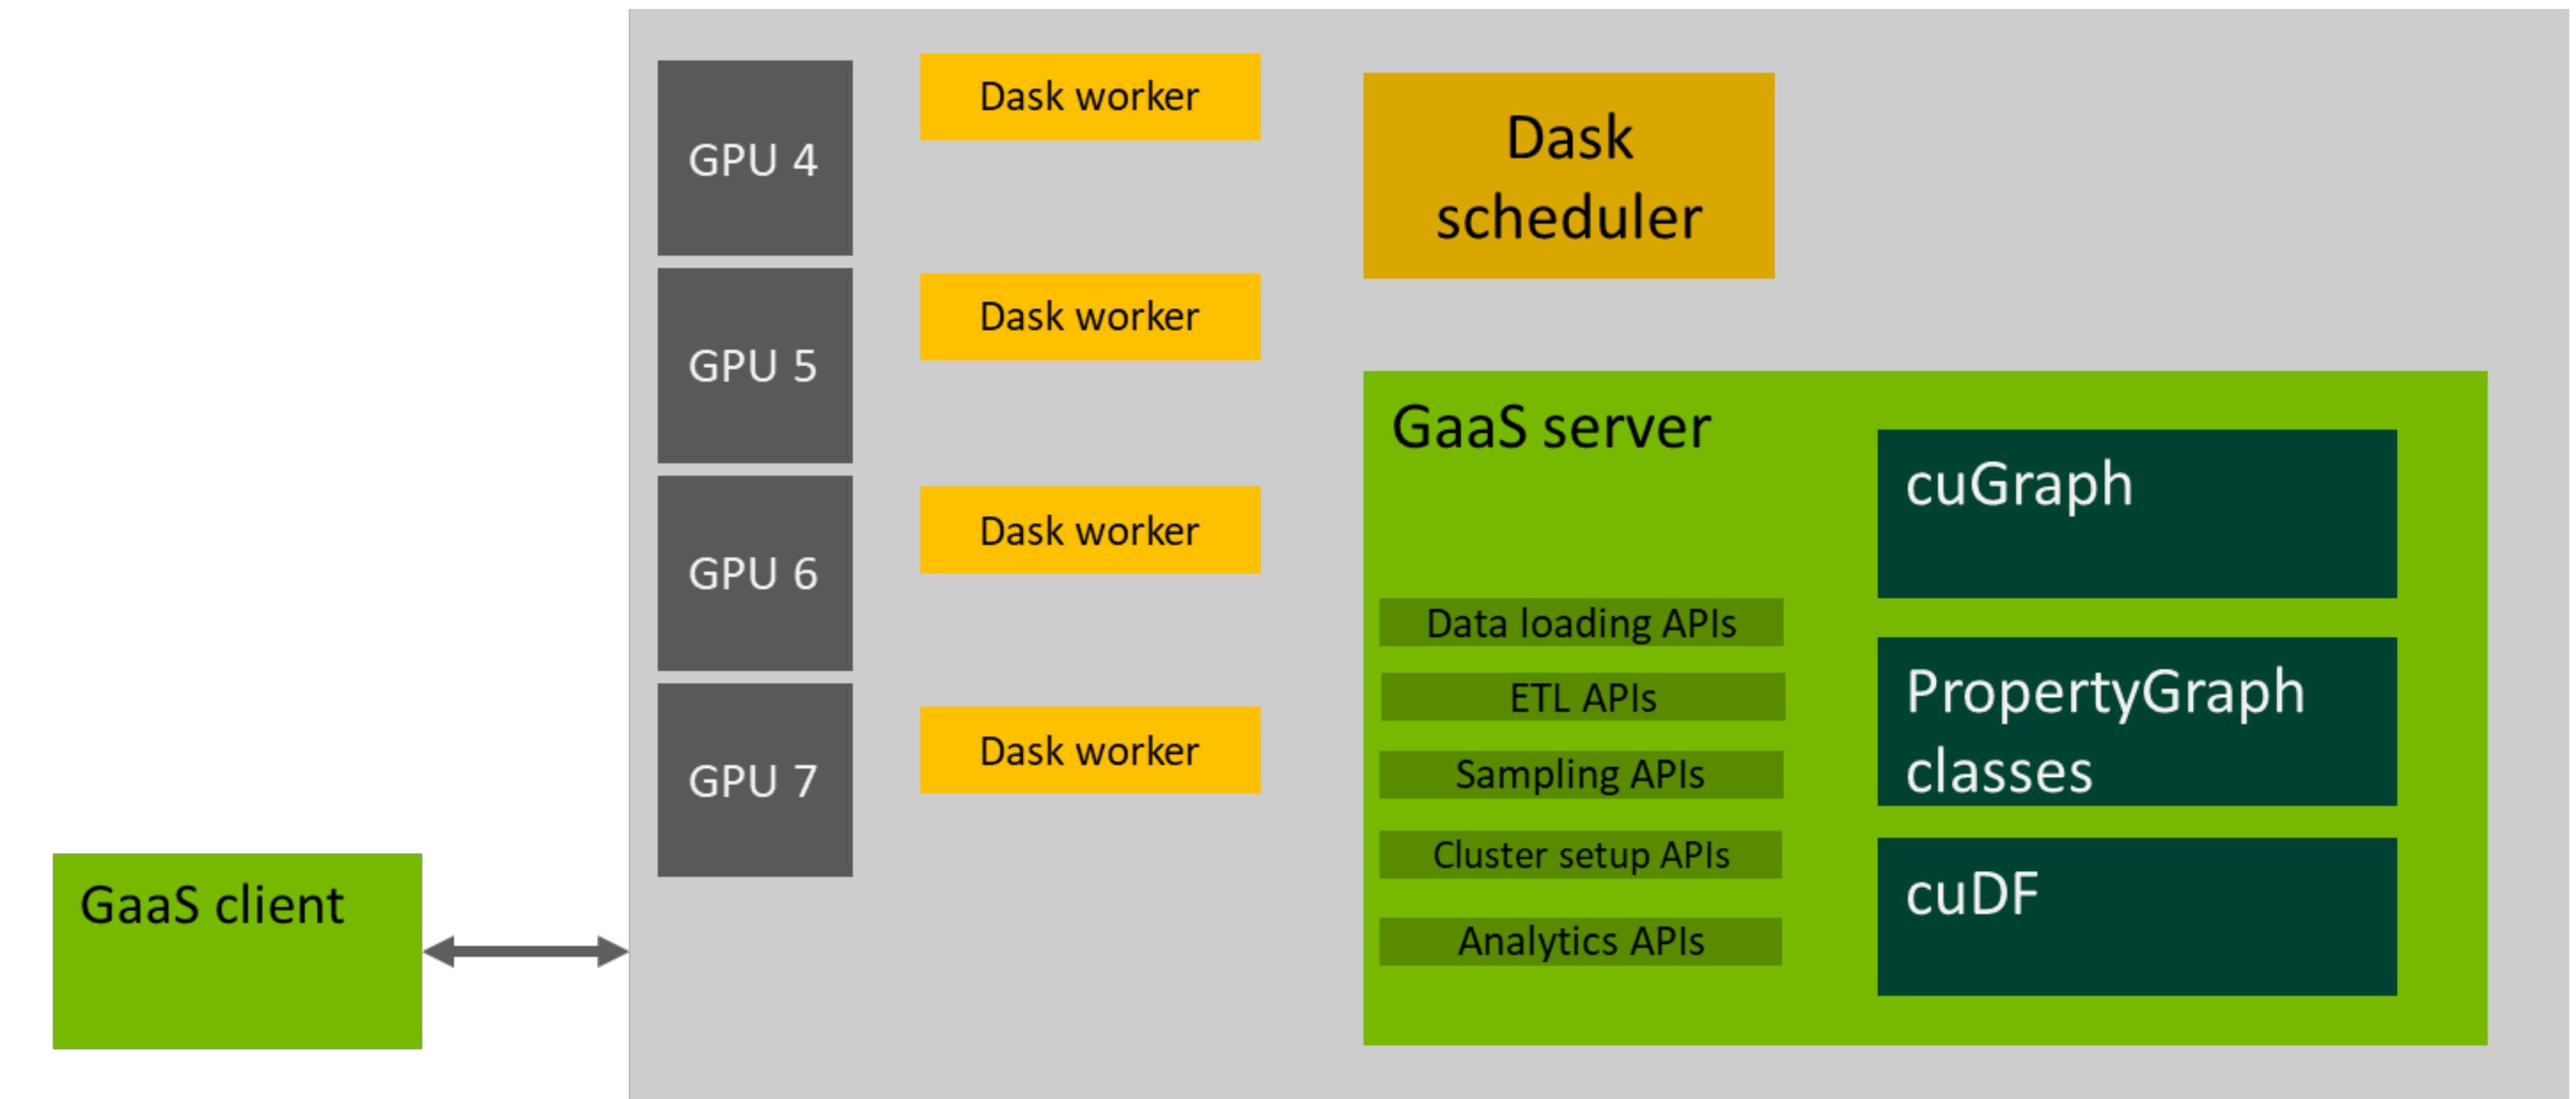
GAAS - UPCOMING ENHANCEMENTS

- Conda packages

- Server - requires GPU and CUDA, installs cuGraph, cuDF, dask, RPC package, others
- Client - no GPU requirement, installs only client, numpy, and RPC package
 - If client has GPUs, UCX-Py can also be installed for device-device copies of results from server to client

- Docker image

- Assign GPUs using “docker run –gpus ...”
- Volume mount a shared dir to “/dask” to share auto-generated dask scheduler JSON configuration for MNMG
- Volume mount a shared dir to “/logs” to access logs from all workers, GaaS, etc.
- Volume mount a shared dir to “/config” to specify other non-default config options (enable/disable NVLink, InfiniBand, force specific protocols, specify device names, etc.)





GNN Support

GNN WITH CUGRAPH AND DGL

FILL THE GAP IN BETWEEN THE SAMPLING AND TRAINING



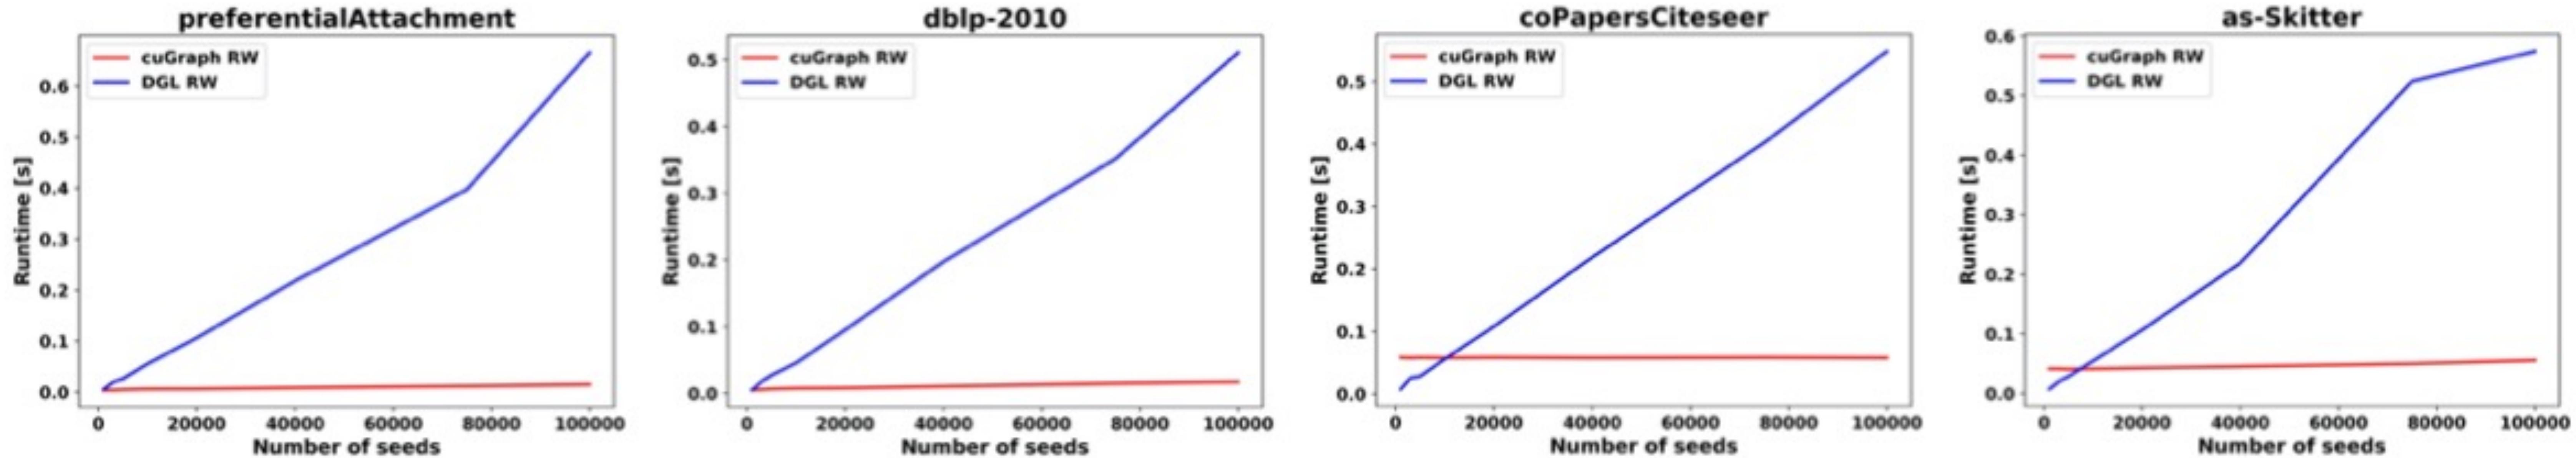
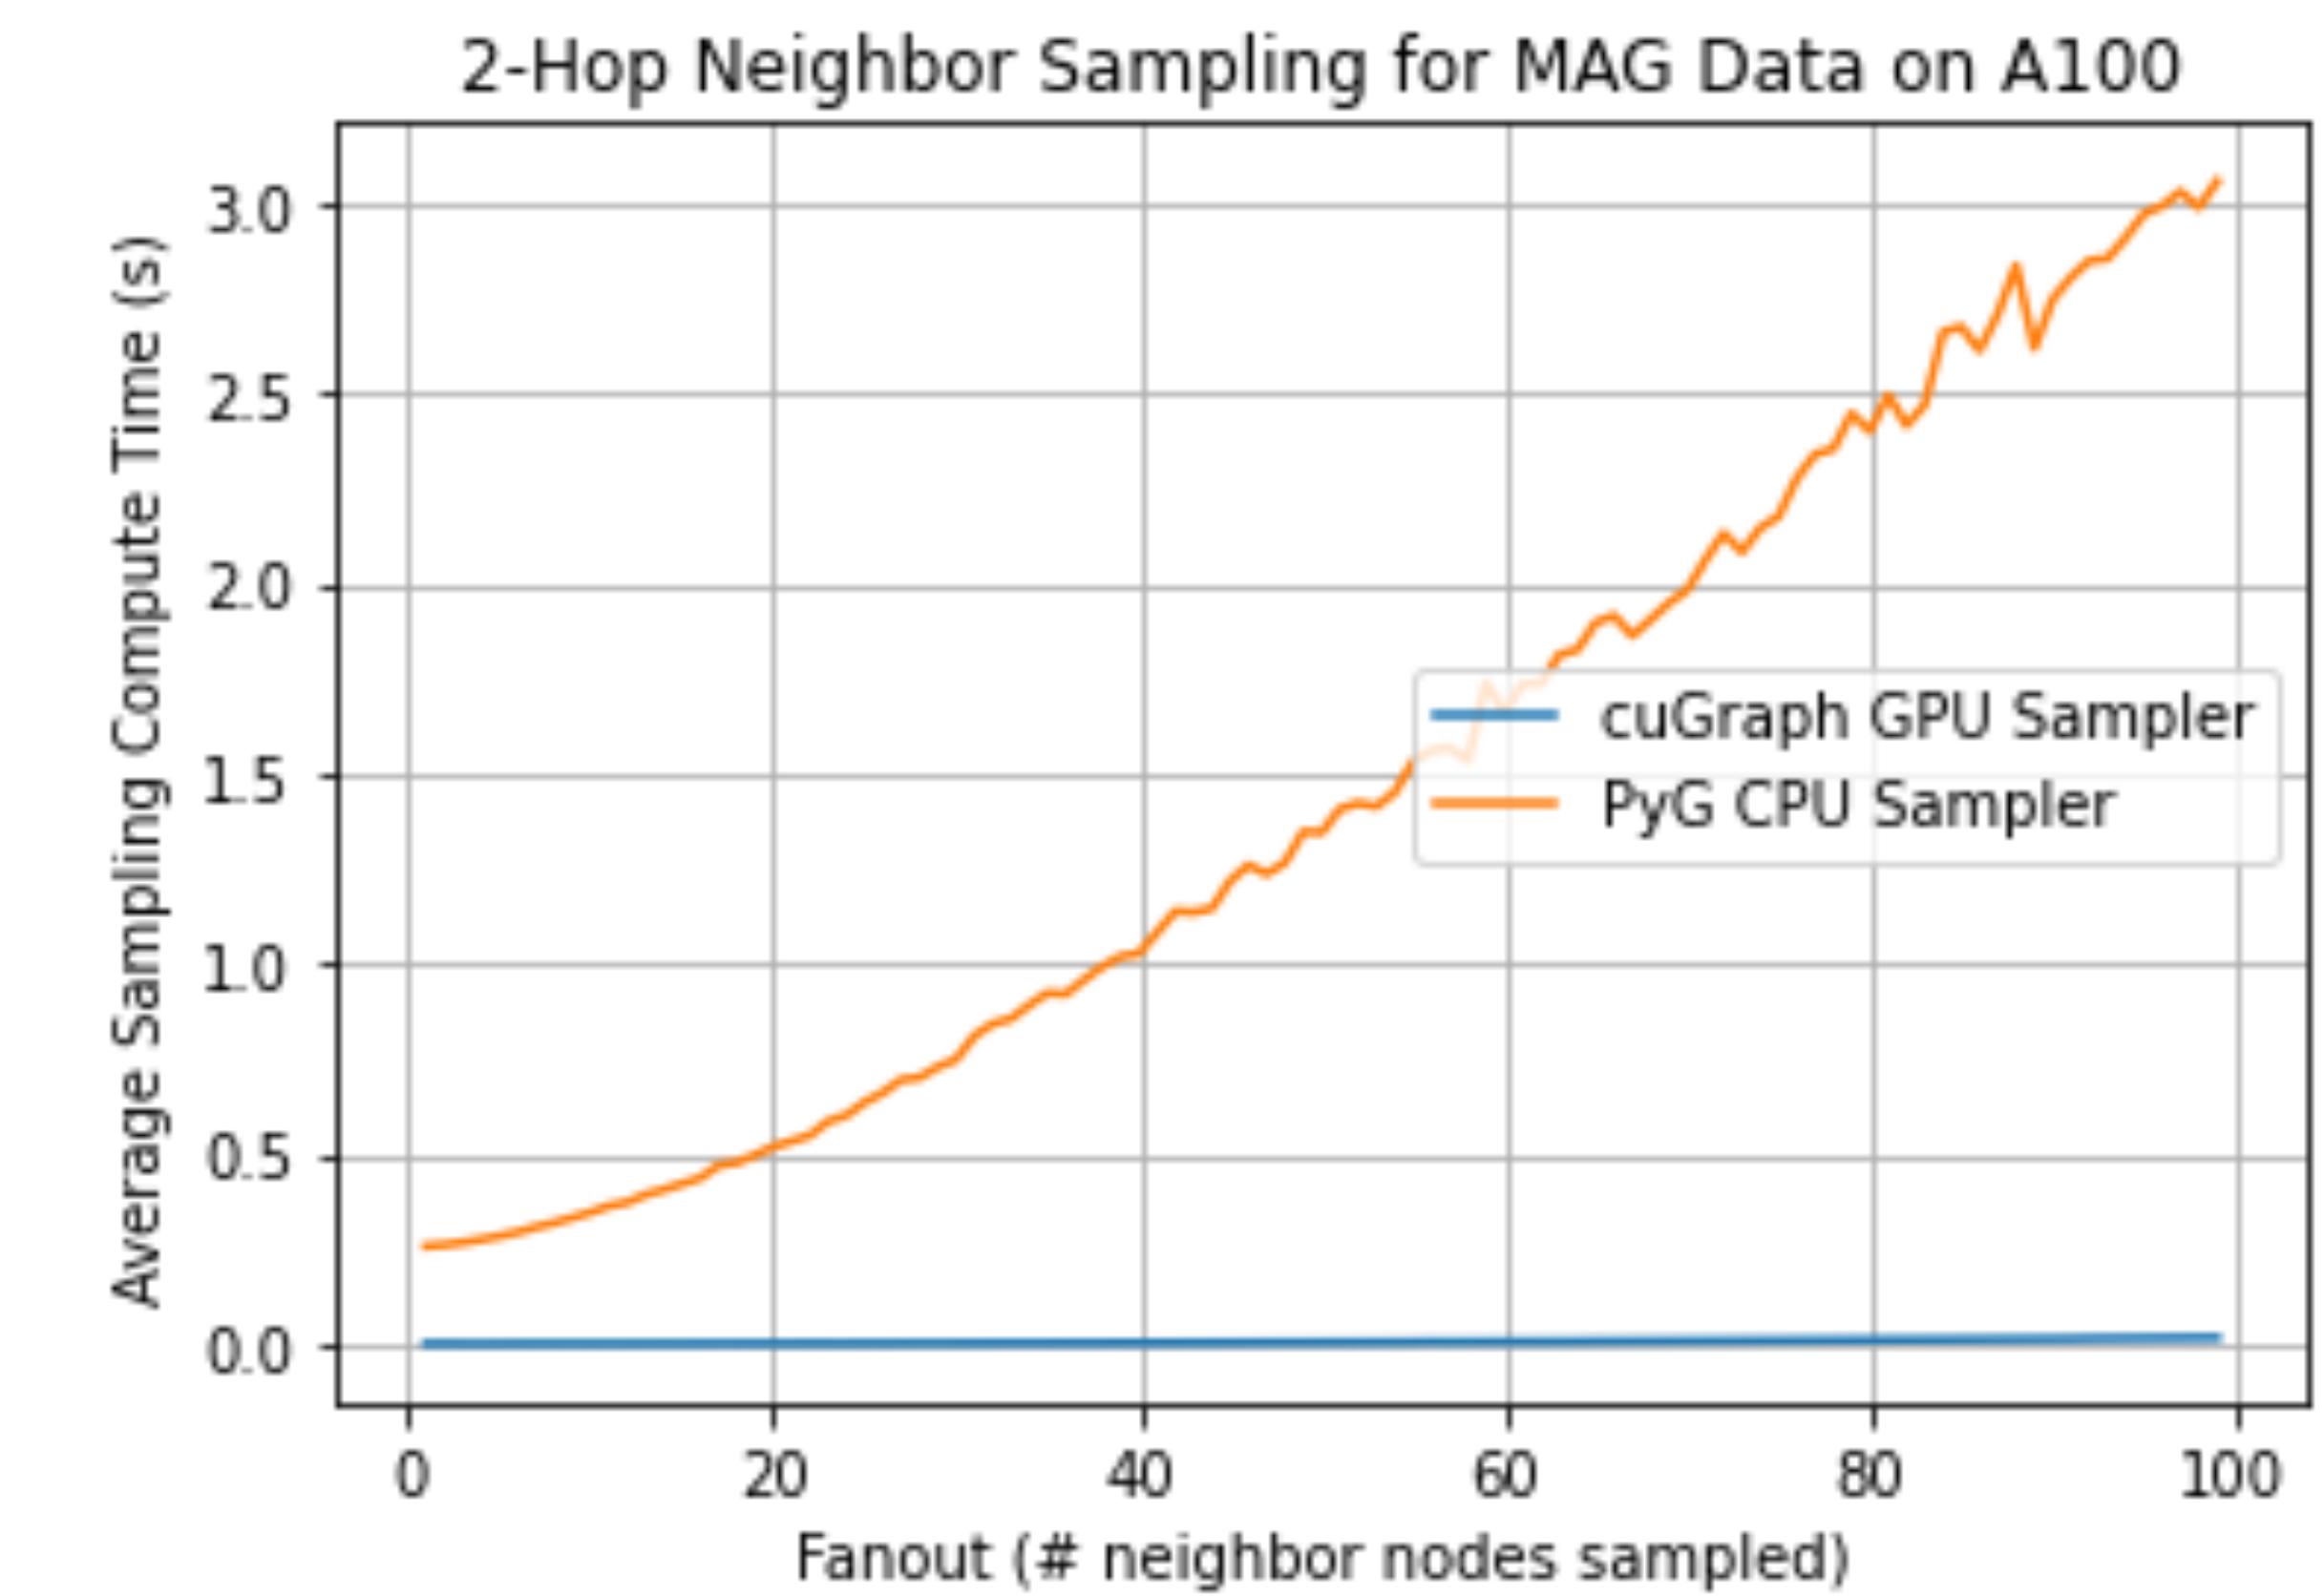
- Motivations:
Existing frameworks do samplings on CPUs and train on GPUs. it is time-consuming.
- Avoid moving data number of epochs X batch between GPUs and CPUs.
- Sampling and training has closer relationship than data processing.

Operations	Data processing	Sampling	Training
Device of existing method	CPU or GPU	CPU	GPU
Devices of our method	GPU	GPU	GPU

GPU ACCELERATED RANDOM WALK

Common Sampling Method

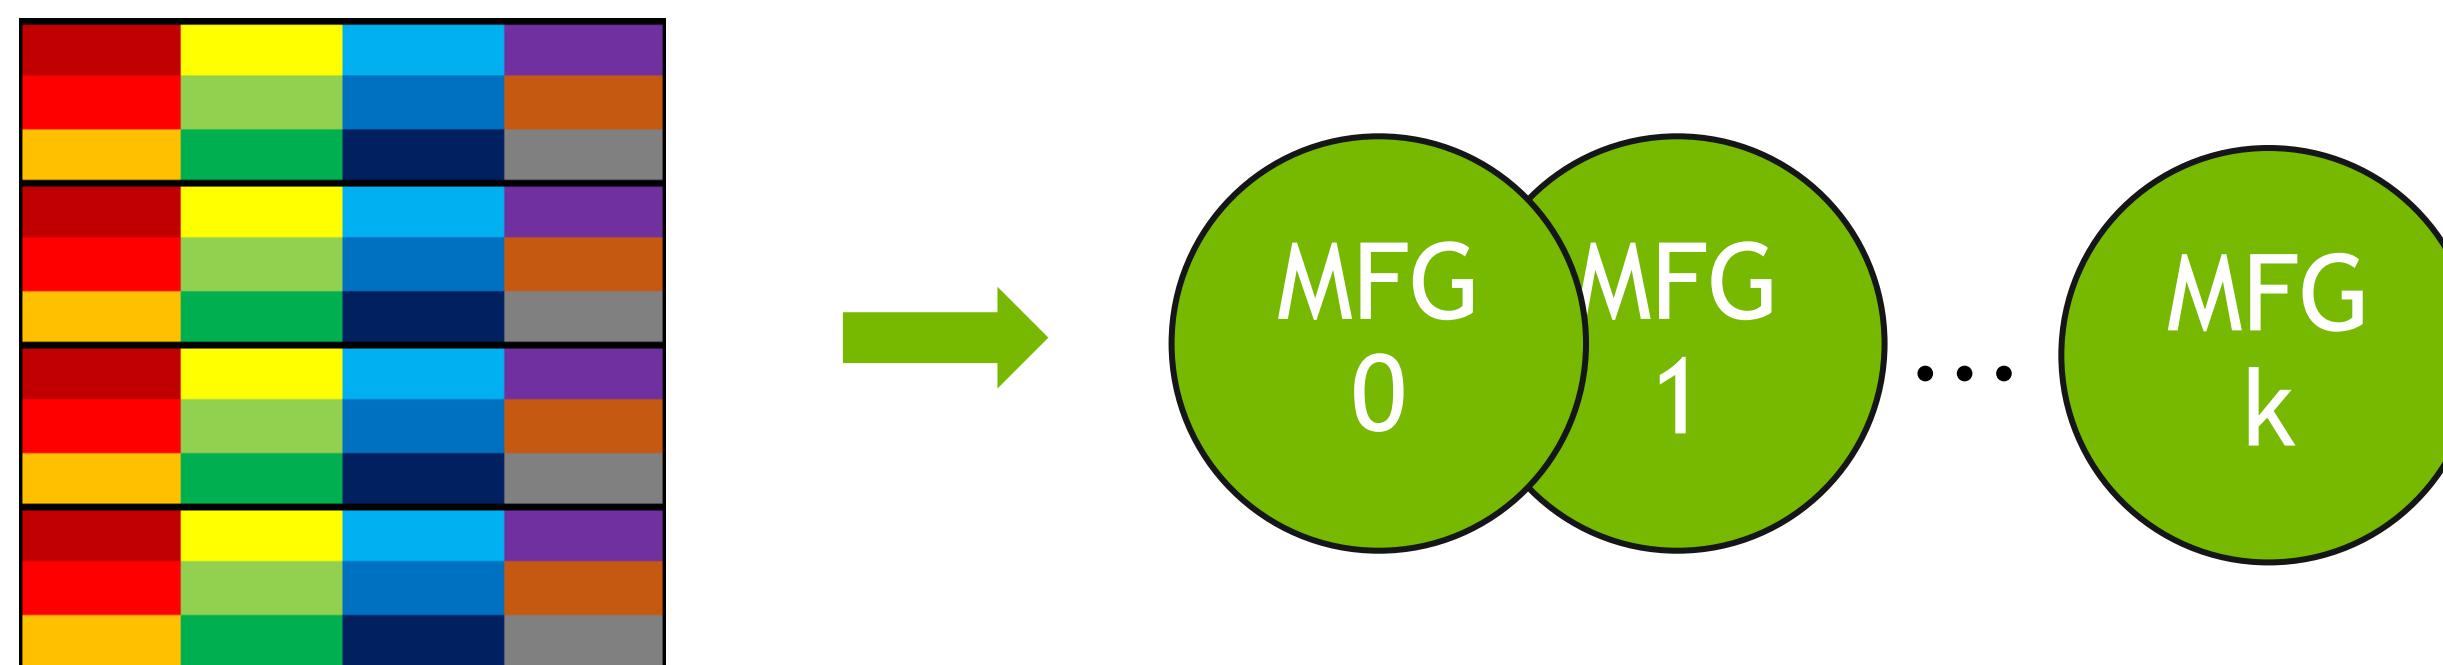
- Graph Sampling can consume 50 - 80% of training time
- cuGraph has been working on expanding the list of graph sampling algorithms and working to have the algorithms support multiple seeds
 - Egonet
 - Random Walk
 - Node2Vec
 - Neighborhood



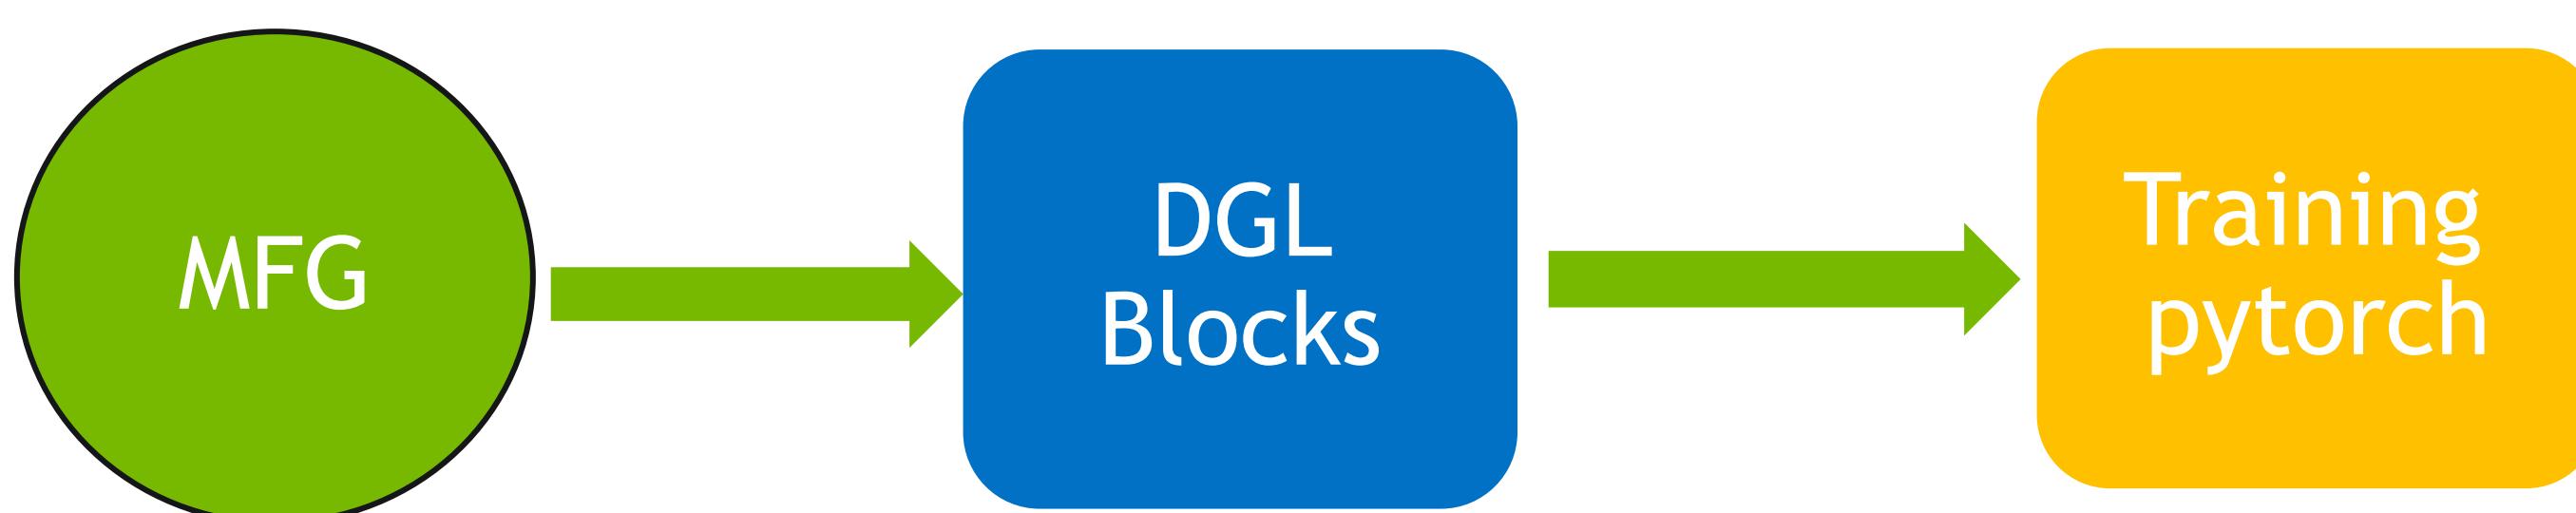
cuGraph RandomWalk vs DGL RandomWalk with various number of seeds

ROLE OF CUGRAPH-OPS

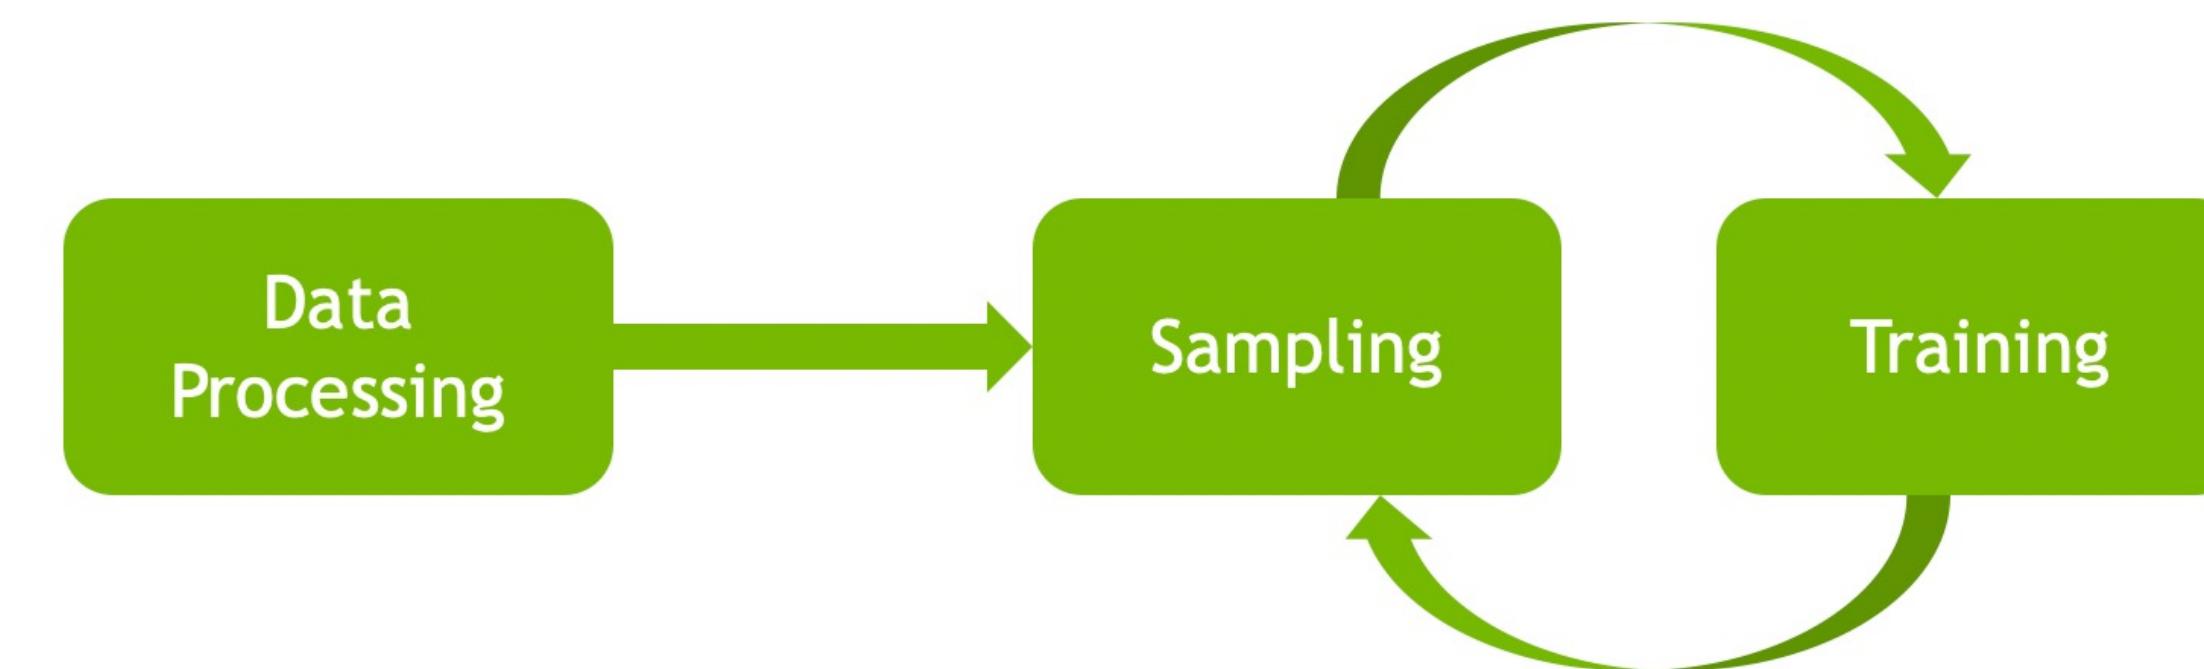
- Neighborhood sampling calls to create Message Flow Graphs



Current workflow



■ Note: Currently we support DGL, will support PyG in future.



Cugraph-ops is a library that is composed of highly optimized and performant primitives associated with GNNs and related graph operations, such as training, sampling and inference.

Currently, we have random walk, neighborhood sampling, and in the future we will support aggregation functions as well.

UPCOMING CUGRAPH FEATURES

DRAFT

- Migrate all Existing algorithms to support MNMG
 - Everything* should smoothly scale from one to thousands of GPUs
 - 32 Graph Algorithms
 - 9 have been upgraded to support Multiple GPUs across Multiple Nodes (MNMG)
- More Algorithms
 - More multi-seed versions,
 - all-pair shortest path,
 - propagation algorithms,
 - flow algorithms, etc..
 - See cuGraph GitHub site for roadmap list of algorithms
- Better Graph Type Support
 - Property Graphs, Bipartite Graphs, Multigraphs, HyperGraphs
- More and Better support for GNNs
 - More sampling algorithms
 - All sampling algorithm supporting Multi-Node Multi-GPU (MNMG) environments for scaling to massive graphs
 - Heterogeneous Graph Support
 - Improved Property Graph
 - More model optimization support in cugraph-ops library
- Data Masking
 - Provide a filter over the data for processing just a portion without extracting or reloading the graph
 - Time Windowing or View by Attribute(s).
- Performance and Scaling improvements
 - This is a constant on-going activity